

Storm Topologies

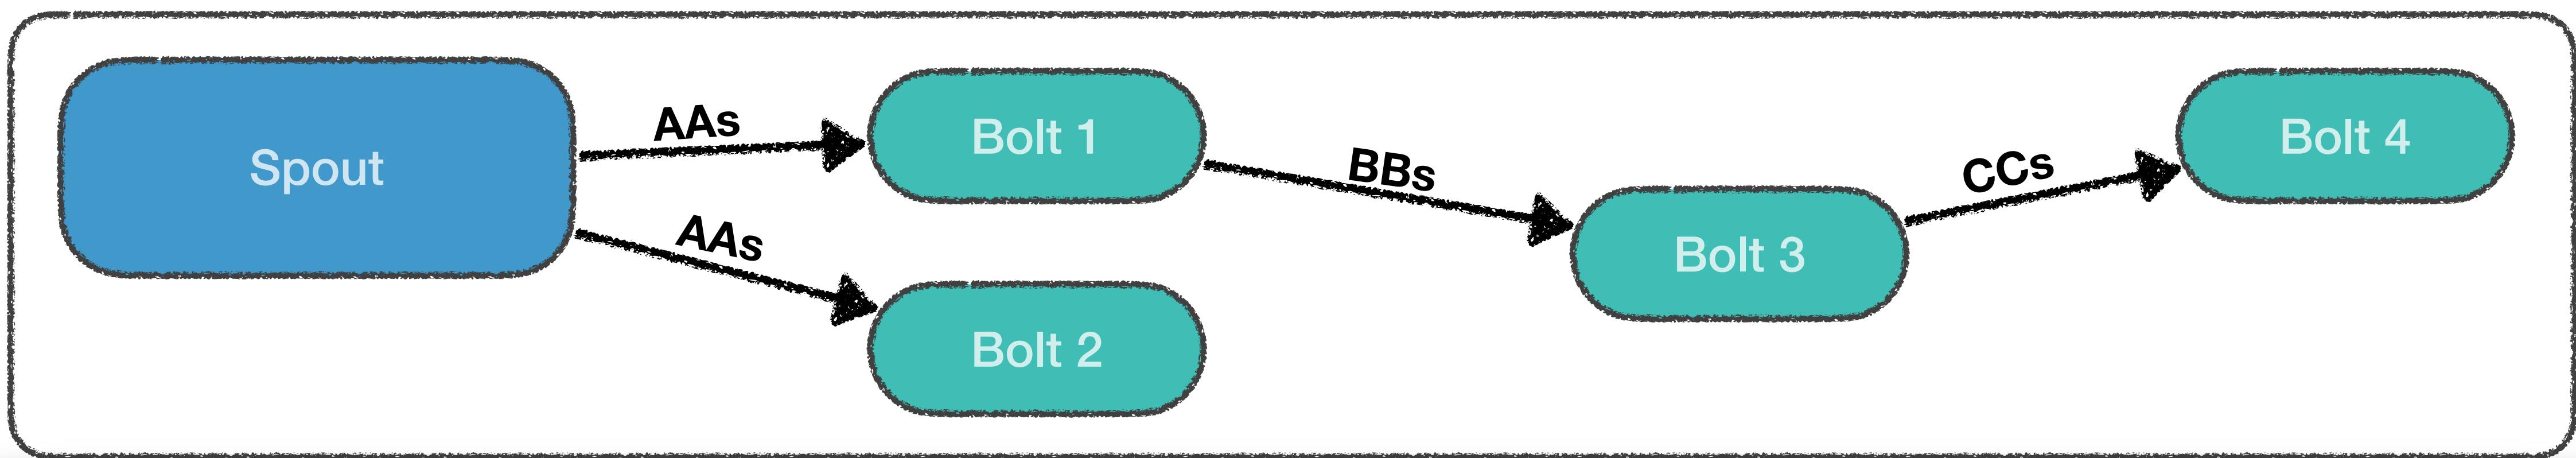
Storm Topologies

- What's a topology;
- Spouts;
- Bolts;
- Topologies
- Reliability;
- Parallelism;
- Stream Grouping.

What's a topology

- A topology encapsulates all the logic for a realtime application;
- Unlike Hadoop's MapReduce jobs that eventually finish, Storm topologies can run forever;
- A topology is a graph of spouts and bolts that are connected with stream groupings;

What's a topology



Topology Direct Acyclic Graph

Spouts

Spouts read data from an external source (Kafka, Rabbit MQ, Kestrel, etc) and emit them into the topology:

- Reliable **spout** can replay tuples if they were failed to be processed;
- Unreliable **spout** forgets about tuples as soon as they are emitted;

Spouts

```
class InstructionTranslator extends RecordTranslator[String, String] {

    override def getFieldsFor(s: String): Fields = new Fields(INSTRUCTION.name)

    override def streams(): util.List[String] = RecordTranslator.DEFAULT_STREAM

    override def apply(consumerRecord: ConsumerRecord[String, String]): util.List[AnyRef] = {

        val message      = consumerRecord.value()
        val partition    = consumerRecord.partition().toString
        val instruction  = getInstruction(message)
        val result       = instruction.map(it => new Values(it))

        result.toOption.getOrElse(null)
    }

    def getInstruction(message: String): Try[Instruction] = Try {...}

    def deserialize(event: String): Feed = {...}
```

Spouts

```
class InstructionTranslator extends RecordTranslator[String] {  
    override def getFieldsFor(s: String): Fields = new Fields(INSTRUCTION.name)  
  
    override def streams(): util.List[String] = RecordTranslator.DEFAULT_STREAM  
  
    override def apply(consumerRecord: ConsumerRecord[String, String]): util.List[AnyRef] = {  
  
        val message      = consumerRecord.value()  
        val partition    = consumerRecord.partition().toString  
        val instruction  = getInstruction(message)  
        val result       = instruction.map(it => new Values(it))  
  
        result.toOption.getOrElse(null)  
    }  
  
    def getInstruction(message: String): Try[Instruction] = Try {...}  
  
    def deserialize(event: String): Feed = {...}
```

Tuple fields to be emitted

Spouts

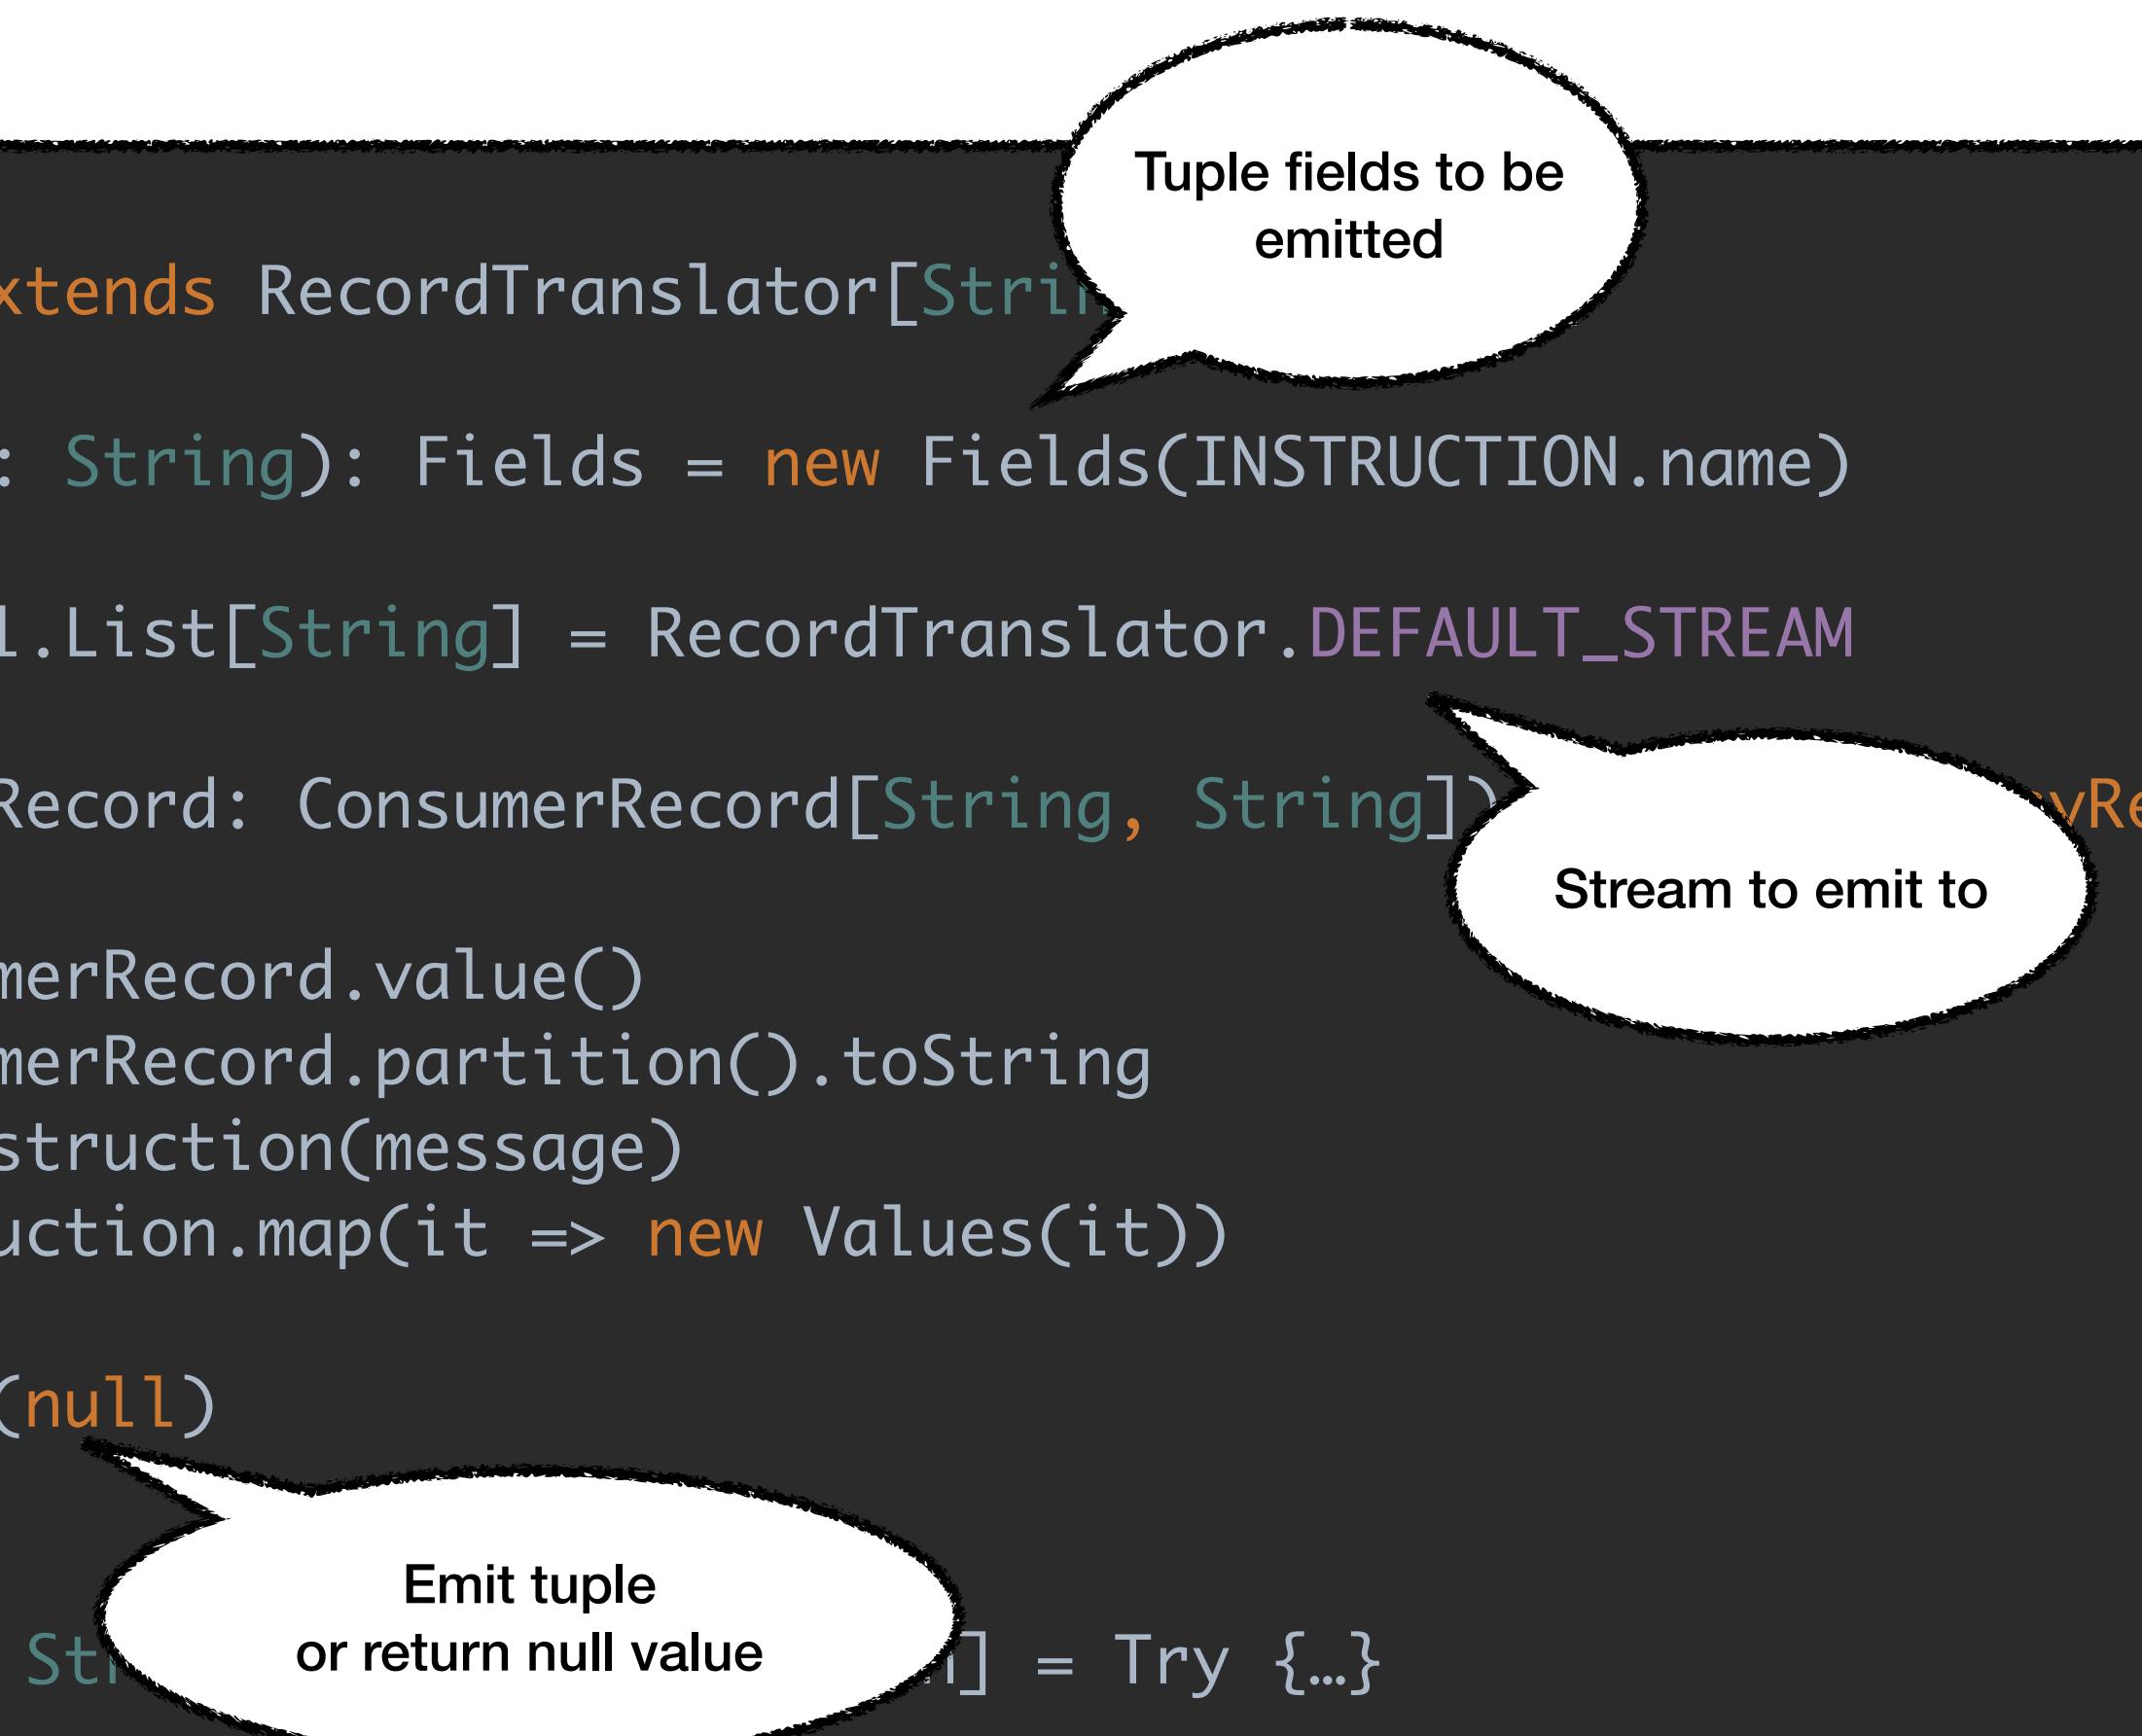
```
class InstructionTranslator extends RecordTranslator[String] {  
    override def getFieldsFor(s: String): Fields = new Fields(INSTRUCTION.name)  
  
    override def streams(): util.List[String] = RecordTranslator.DEFAULT_STREAM  
  
    override def apply(consumerRecord: ConsumerRecord[String, String]): Try[Values] = {  
        val message      = consumerRecord.value()  
        val partition   = consumerRecord.partition().toString  
        val instruction = getInstruction(message)  
        val result       = instruction.map(it => new Values(it))  
  
        result.toOption.getOrElse(null)  
    }  
  
    def getInstruction(message: String): Try[Instruction] = Try {...}  
  
    def deserialize(event: String): Feed = {...}}
```

Tuple fields to be emitted

Stream to emit to

Spouts

```
class InstructionTranslator extends RecordTranslator[String] {  
    override def getFieldsFor(s: String): Fields = new Fields(INSTRUCTION.name)  
  
    override def streams(): util.List[String] = RecordTranslator.DEFAULT_STREAM  
  
    override def apply(consumerRecord: ConsumerRecord[String, String]): Try[Values] = {  
        val message      = consumerRecord.value()  
        val partition   = consumerRecord.partition().toString  
        val instruction = getInstruction(message)  
        val result       = instruction.map(it => new Values(it))  
  
        result.toOption.getOrElse(null)  
    }  
  
    def getInstruction(message: String): Try[Values] = Try {...}  
  
    def deserialize(event: String): Feed = {...}
```



A diagram illustrating the logic of the `apply` method in the `InstructionTranslator` class. Three callout bubbles point to different parts of the code:

- A bubble above the `instruction.map` line contains the text "Tuple fields to be emitted".
- A bubble to the right of the `Try[Values]` return type contains the text "Stream to emit to".
- A bubble below the `result.toOption.getOrElse` line contains the text "Emit tuple or return null value".

Bolts

All the processing in topologies is done in **bolts**:

- **Bolts** can do filtering, functions, aggregations, joins, persisting to a data store, collect data in-memory, connect to an external service, etc.
- They can also emit to more than one stream.

Bolts

```
trait ProcessorBolt extends BaseRichBolt {

    Var output: OutputCollector = _
    Val processor: Processor[Instruction]

    override def prepare(map: util.Map[_, _], context: TopologyContext, output: OutputCollector): Unit = {
        output = output
    }

    override def execute(tuple: Tuple): Unit = {
        val instruction = tuple.getValueByField(INSTRUCTION.name).asInstanceOf[Instruction]
        val results = processor.process(instruction)

        results.map(result => output.emit(tuple, new Values(result)))
        output.ack(tuple)
    }

    override def declareOutputFields(outputFieldsDeclarer: OutputFieldsDeclarer): Unit = {
        outputFieldsDeclarer.declare(new Fields(MARKET.name))
    }
}
```

Bolts

```
trait ProcessorBolt extends BoltInterface[BoltContext, Tuple, OutputCollector] with BoltFn[INSTRUCTION, MARKET] with BoltFn[Instruction, Values] with BoltFn[Tuple, Values] {  
    var output: OutputCollector  
    val processor: Processor[INSTRUCTION, MARKET, Instruction, Values]  
  
    override def prepare(map: util.Map[_, _], context: TopologyContext, output: OutputCollector): Unit = {  
        this.output = output  
    }  
  
    override def execute(tuple: Tuple): Unit = {  
  
        val instruction = tuple.getValueByField(INSTRUCTION.name).asInstanceOf[Instruction]  
        val results = processor.process(instruction)  
  
        results.map(result => output.emit(tuple, new Values(result)))  
        output.ack(tuple)  
    }  
  
    override def declareOutputFields(outputFieldsDeclarer: OutputFieldsDeclarer): Unit = {  
        outputFieldsDeclarer.declare(new Fields(MARKET.name))  
    }  
}
```



Dependencies are initialized here

Bolts

```
trait ProcessorBolt extends BaseProcessorBolt {
    var output: OutputCollector
    val processor: Processor[Instruction]
    override def prepare(map: util.Map[_, _], context: TopologyContext, output: OutputCollector): Unit = {
        output = output
    }

    override def execute(tuple: Tuple): Unit = {
        val instruction = tuple.getValueByFieldIndex(0).asInstanceOf[Instruction]
        val results = processor.process(instruction)
        results.map(result => output.emit(tuple, new Values(result)))
        output.ack(tuple)
    }

    override def declareOutputFields(outputFieldsDeclarer: OutputFieldsDeclarer): Unit = {
        outputFieldsDeclarer.declare(new Fields(MARKET.name))
    }
}
```

Dependencies are initialized here

Anchoring to the input tuple

Bolts

```
trait ProcessorBolt extends Bolt {
    var output: OutputCollector
    val processor: Processor[Instruction]
    override def prepare(map: util.Map[_, _], context: TopologyContext, output: OutputCollector): Unit = {
        output = output
    }
    override def execute(tuple: Tuple): Unit = {
        val instruction = tuple.getValueByField(MARKET.name).asInstanceOf[Instruction]
        val results = processor.process(instruction)
        results.map(result => output.emit(tuple, new Values(result)))
        output.ack(tuple)
    }
    override def declareOutputFields(outputFieldsDeclarer: OutputFieldsDeclarer): Unit = {
        outputFieldsDeclarer.declare(new Fields(MARKET.name))
    }
}
```

Dependencies are initialized here

Anchoring to the input tuple

Tuple fields to be emitted

Bolts

```
trait OutputBolt extends BaseRichBolt with LazyLogging {

  var output: OutputCollector = _
  val source: OutputSource[Market]

  override def prepare(map: util.Map[_, _], topologyContext: TopologyContext, outputCollector: OutputCollector): Unit = {
    output = outputCollector
  }

  override def execute(tuple: Tuple): Unit = {
    val market = tuple.getValueByField(MARKET.name).asInstanceOf[Market]

    source.write(key, market) match {
      case Success(_) => output.ack(tuple)
      case Failure(_) => output.fail(tuple)
    }
  }

  override def declareOutputFields(declarer: OutputFieldsDeclarer): Unit = {}
}
```

Bolts

```
trait OutputBolt extends BaseRichBolt with LazyLogging {  
  
  Var output: OutputCollector = _  
  val source: OutputSource[Market]  
  
  override def prepare(map: util.Map[_, _], topologyContext: TopologyContext, outputCollector: OutputCollector): Unit = {  
    output = outputCollector  
  }  
  
  override def execute(tuple: Tuple): Unit = {  
    val market = tuple.Acking the input tupleField(MARKET.name).asInstanceOf[Market]  
  
    source.write(key, market) match {  
      case Success(_) => output.ack(tuple)  
      case Failure(_) => output.fail(tuple)  
    }  
  }  
  
  override def declareOutputFields(declarer: OutputFieldsDeclarer): Unit = {}
```

Bolts

```
trait OutputBolt extends BaseRichBolt with LazyLogging {  
  
  var output: OutputCollector = _  
  val source: OutputSource[Market]  
  
  override def prepare(map: util.Map[_, _], topologyContext: TopologyContext, outputCollector: OutputCollector): Unit = {  
    output = outputCollector  
  }  
  
  override def execute(tuple: Tuple): Unit = {  
    val market = tuple.Acking the input tupleField(MARKET).Field(MARKET) name.asInstanceOf[Market]  
    source.write(key, market) match {  
      case Success(_) => output.ack(tuple)  
      case Failure(_) => output.fail(tuple)  
    }  
  }  
  
  override def declareOutputFields(declarer: OutputFieldsDeclarer): Unit = {}
```

Bolts

```
trait OutputBolt extends BaseRichBolt with LazyLogging {  
  
  var output: OutputCollector = _  
  val source: OutputSource[Market]  
  
  override def prepare(map: util.Map[_, _], topologyContext: TopologyContext, outputCollector: OutputCollector): Unit = {  
    output = outputCollector  
  }  
  
  override def execute(tuple: Tuple): Unit = {  
    val market = tuple.Acking the input tupleField(MARKET).Field(MARKET) name.asInstanceOf[Market]  
    source.write(key, market) match {  
      case Success(_) => output.ack(tuple)  
      case Failure(_) => output.fail(tuple)  
    }  
  }  
  
  override def declareOutputFields(declarer: OutputFieldsDeclarer): Unit = {}
```

You can also fail the
input tuple

There's no fields to
declare

Topologies

As said before **topologies** encapsulate all the logic for the application, it is in the **topology builder** where we define stream groupings between spout and bolts.

Topologies

```
trait SimpleTopology extends SpoutModule with ProcessorModule with OutputModule {

    val SPOUT      : String = "SimpleSpout"
    val PROCESSOR : String = "SimpleBolt"
    val OUTPUT     : String = "SecondSimpleBolt"

    def name() = "SimpleTopology"

    def create(): StormTopology = {

        val builder = new TopologyBuilder()

        builder.setSpout(SPOUT, getSpout(), 1)

        builder.setBolt(PROCESSOR, getProcessor(), 1).shuffleGrouping(SPOUT)

        builder.setBolt(OUTPUT, getOutput(), 1).shuffleGrouping(PROCESSOR)

        builder.createTopology()
    }

    def config(): Config = {
        val cg = new Config
        cg.setNumWorkers(1)
        cg
    }
}
```

Topologies

```
trait SimpleTopology extends SpoutModule with ProcessorModule with OutputModule {

    val SPOUT      : String = "SimpleSpout"
    val PROCESSOR : String = "SimpleBolt"
    val OUTPUT     : String = "SecondSimpleBolt"

    def name() = "SimpleTopology"

    def create(): StormTopology = {
        val builder = new TopologyBuilder()
        builder.setSpout(SPOUT, getSpout(), 1)

        builder.setBolt(PROCESSOR, getProcessor(), 1).shuffleGrouping(SPOUT)
        builder.setBolt(OUTPUT, getOutput(), 1).shuffleGrouping(PROCESSOR)
        builder.createTopology()
    }

    def config(): Config = {
        val cg = new Config
        cg.setNumWorkers(1)
        cg
    }
}
```

Topologies

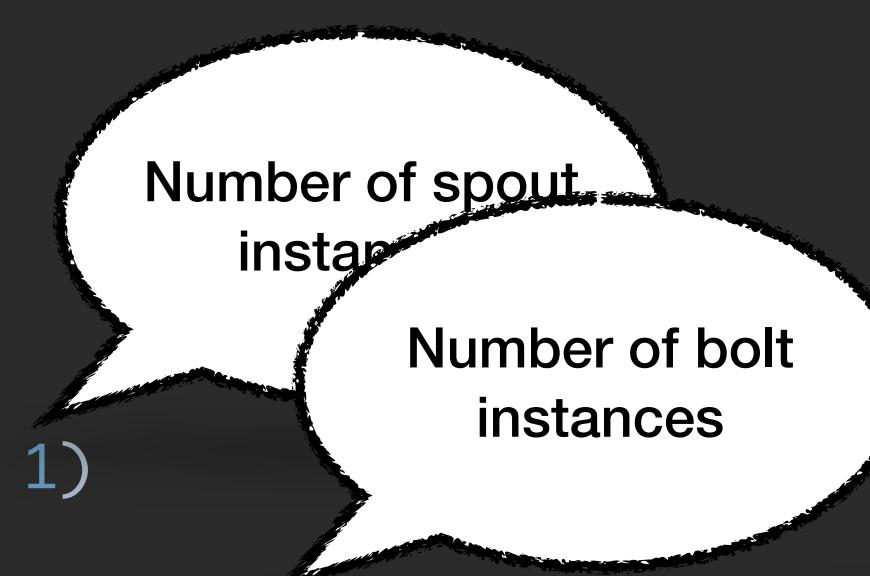
```
trait SimpleTopology extends SpoutModule with ProcessorModule with OutputModule {

    val SPOUT      : String = "SimpleSpout"
    val PROCESSOR : String = "SimpleBolt"
    val OUTPUT     : String = "SecondSimpleBolt"

    def name() = "SimpleTopology"

    def create(): StormTopology = {
        val builder = new TopologyBuilder()
        builder.setSpout(SPOUT, getSpout(), 1)
        builder.setBolt(PROCESSOR, getProcessor(), 1).shuffleGrouping(SPOUT)
        builder.setBolt(OUTPUT, getOutput(), 1).shuffleGrouping(PROCESSOR)
        builder.createTopology()
    }

    def config(): Config = {
        val cg = new Config
        cg.setNumWorkers(1)
        cg
    }
}
```



Topologies

```
trait SimpleTopology extends SpoutModule with ProcessorModule with OutputModule {

    val SPOUT      : String = "SimpleSpout"
    val PROCESSOR : String = "SimpleBolt"
    val OUTPUT     : String = "SecondSimpleBolt"

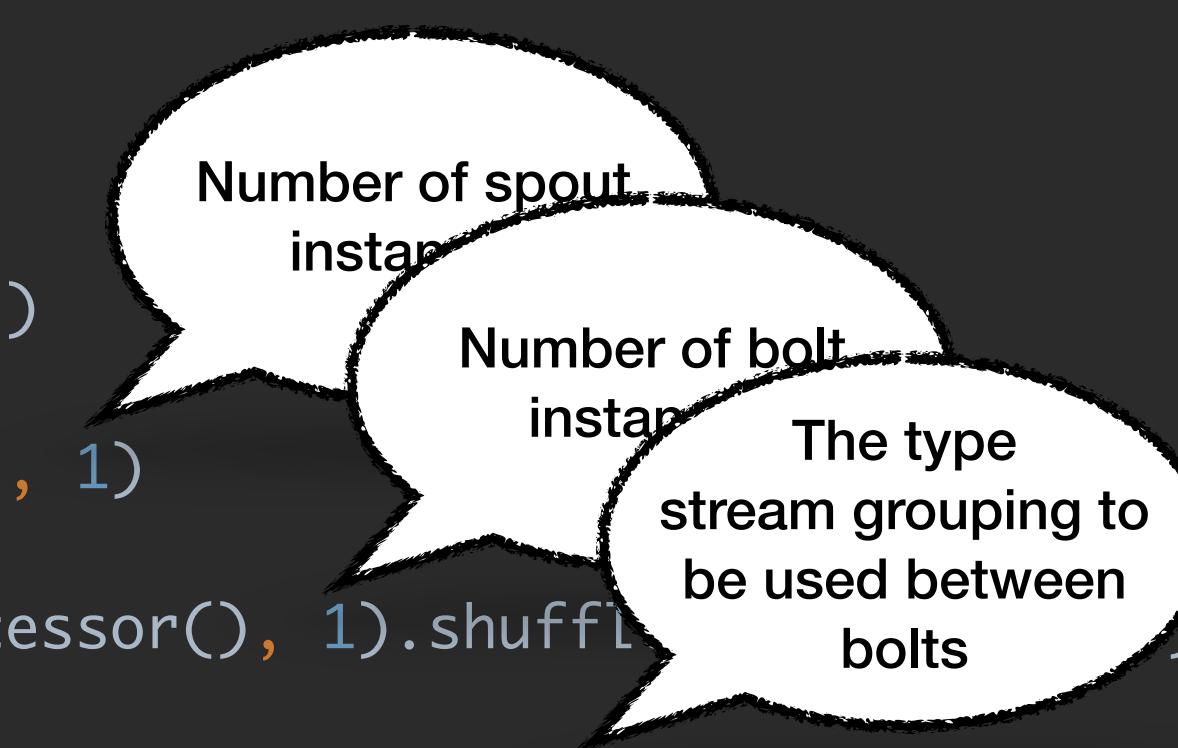
    def name() = "SimpleTopology"

    def create(): StormTopology = {
        val builder = new TopologyBuilder()

        builder.setSpout(SPOUT, getSpout(), 1)
        builder.setBolt(PROCESSOR, getProcessor(), 1).shuffleGrouping()
        builder.setBolt(OUTPUT, getOutput(), 1).shuffleGrouping(PROCESSOR)

        builder.createTopology()
    }

    def config(): Config = {
        val cg = new Config
        cg.setNumWorkers(1)
        cg
    }
}
```



Topologies

```
trait SimpleTopology extends SpoutModule with ProcessorModule with OutputModule {

    val SPOUT      : String = "SimpleSpout"
    val PROCESSOR : String = "SimpleBolt"
    val OUTPUT     : String = "SecondSimpleBolt"

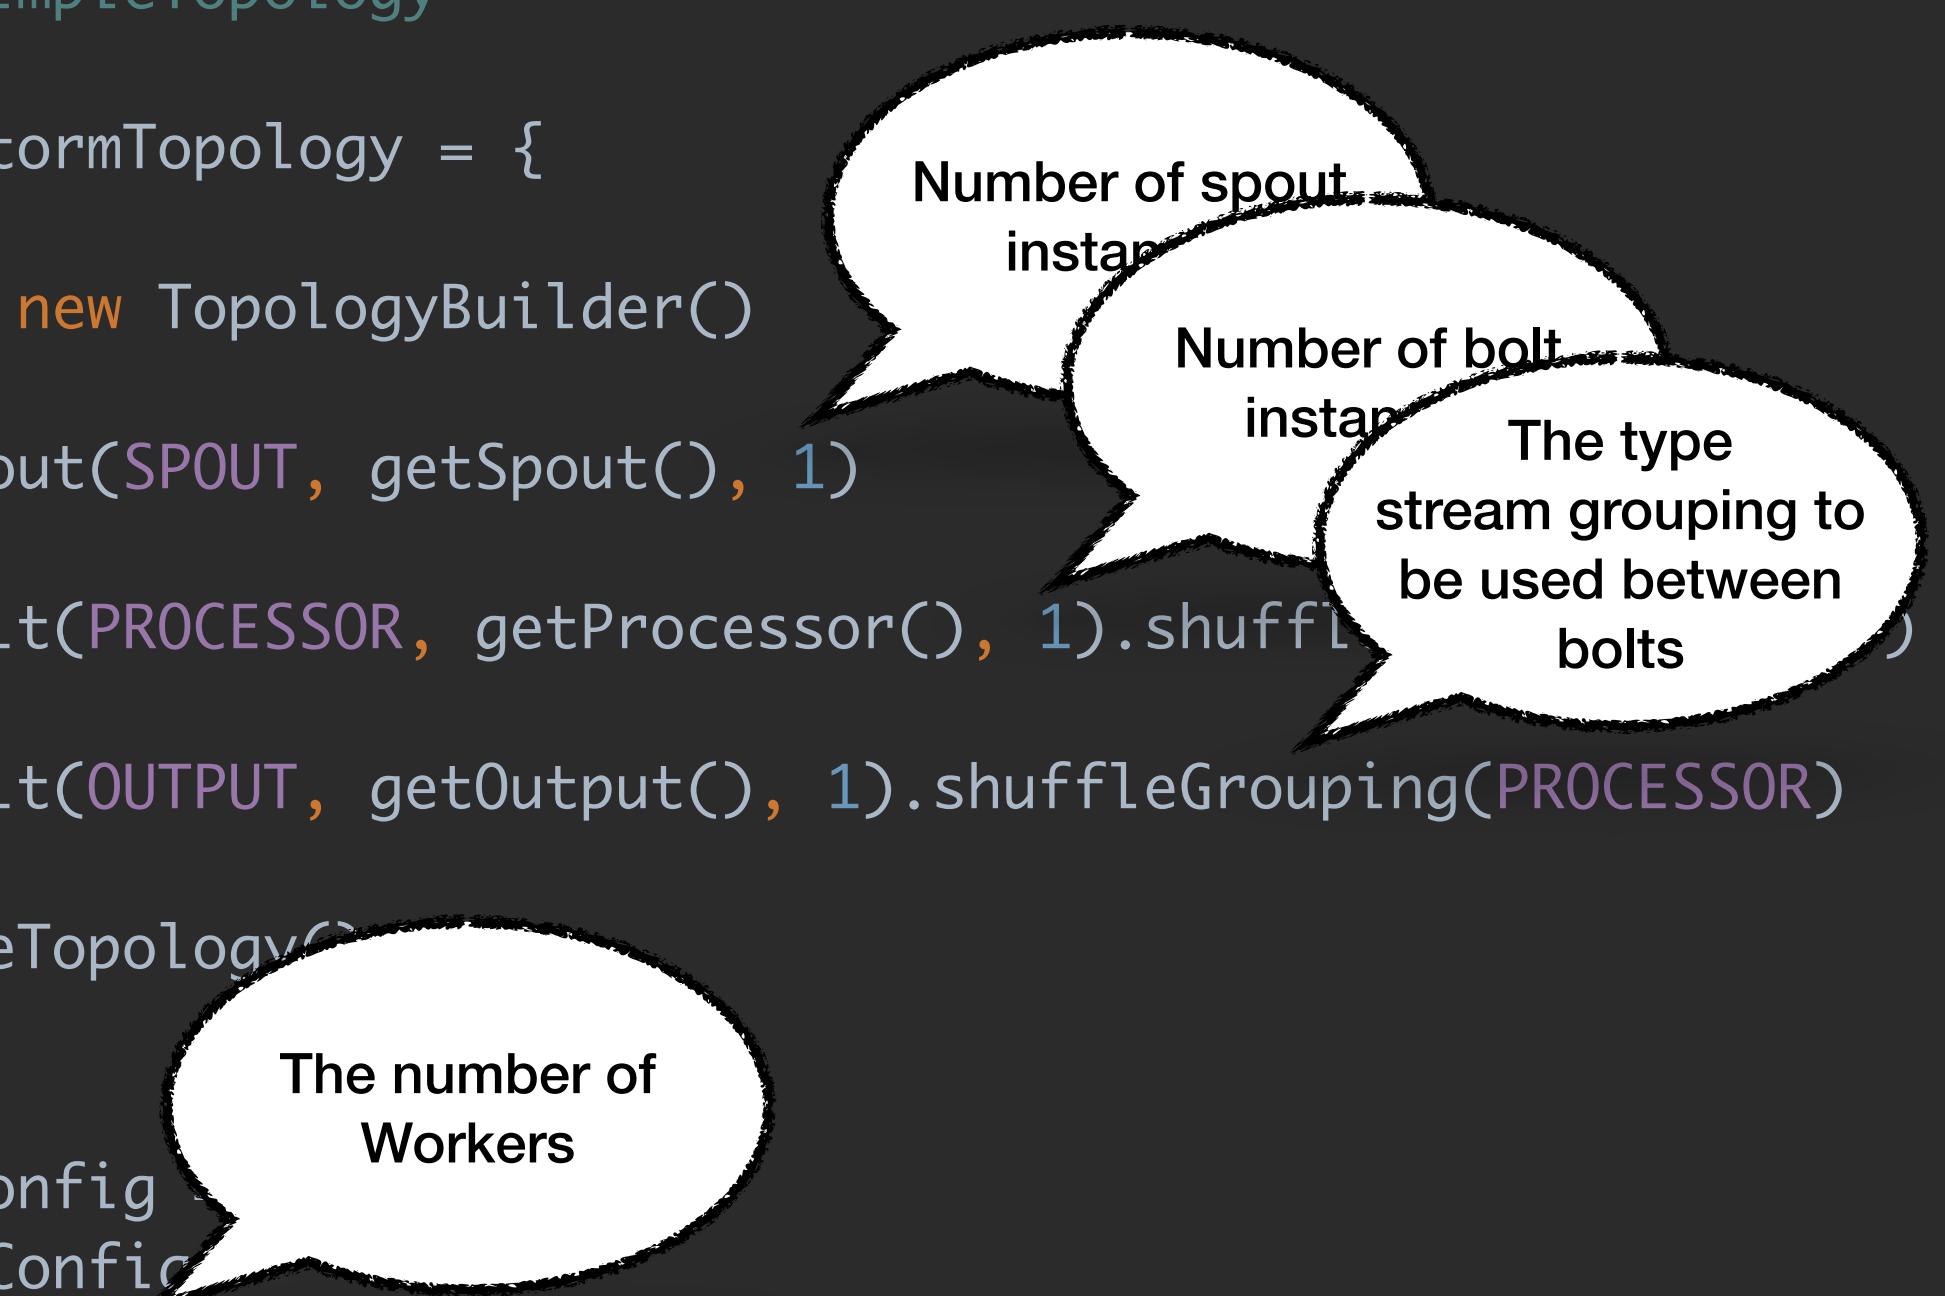
    def name() = "SimpleTopology"

    def create(): StormTopology = {
        val builder = new TopologyBuilder()

        builder.setSpout(SPOUT, getSpout(), 1)
        builder.setBolt(PROCESSOR, getProcessor(), 1).shuffleGrouping()
        builder.setBolt(OUTPUT, getOutput(), 1).shuffleGrouping(PROCESSOR)

        builder.createTopology()
    }

    def config(): Config = {
        val cg = new Config()
        cg.setNumWorkers(1)
        cg
    }
}
```



The code defines a trait `SimpleTopology` that extends `SpoutModule`, `ProcessorModule`, and `OutputModule`. It includes constants for the spout and processor types, and an output bolt. The `name()` method returns "SimpleTopology". The `create()` method builds a topology using a `TopologyBuilder`. It sets up a spout with 1 instance, a processor bolt with 1 instance, and an output bolt with 1 instance. Stream grouping is used between the processor and output bolts. The `config()` method creates a `Config` object and sets the number of workers to 1.

- Number of spout instances
- Number of bolt instances
- The type stream grouping to be used between bolts
- The number of Workers

Topologies

Name	Id	Owner	Status	Uptime	Num workers	Num executors	Num tasks	Replication count	Assigned Mem (MB)	Scheduler Info
RAMP_INSTRUCTIONS_PROCESSOR-football-119913_102	RAMP_INSTRUCTIONS_PROCESSOR-football-119913_102-48-1556614523	storm	ACTIVE	27d 6h 59m 59s	22	1774	1774	3	18304	

Topology stats						
Window	Emitted	Transferred	Complete latency (ms)	Acked	Failed	
10m 0s	72982	82778	6.600	10000		
3h 0m 0s	1380836	1601501	6.173	229306		
1d 0h 0m 0s	15330332	17881216	11.148	2616412		
All time	557816000	649495000	9.753	94589000		

Spouts (All time)												
Search:												
Id	Executors	Tasks	Emitted	Transferred	Complete latency (ms)	Acked	Failed	Error Host	Error Port	Last error	Error Time	
Instructions	50	50	94806000	94806000	9.753	94589000	0					
Replication Instructions	50	50	0	0	0.000	0	0					

Showing 1 to 2 of 2 entries

Bolts (All time)														
Search:														
Id	Executors	Tasks	Emitted	Transferred	Capacity (last 10m)	Execute latency (ms)	Executed	Process latency (ms)	Acked	Failed	Error Host	Error Port	Last error	Error Time
CircuitBreaker Outputsource	1	1	14000	14000	0.000	17.401	157000	17.197	157000	0				
Dual Mapping	154	154	94582000	94582000	0.002	0.036	94602000	0.028	94597000	0				
ES Logger Outputsource	44	44	0	0	0.000	0.235	91689000	0.227	91691000	0				
FMM Create Outputsource	44	44	0	0	0.165	195.152	1759000	216.619	1763000	0				
FMM Refresh Outputsource	88	88	0	0	0.000	96.442	52000	109.000	52000	0				
FMM Result Outputsource	88	88	0	0	0.147	131.147	1131000	123.818	1128000	0				
FMM Update Outputsource	836	836	0	0	0.187	67.665	88795000	67.729	88784000	0				
Instructions Processor	154	154	368396000	460075000	0.008	3.461	94617000	3.463	94615000	0				
PPB Outputsource	44	44	0	0	0.002	0.201	91683000	0.185	91678000	0				
SPORTEX Outputsource	44	44	0	0	0.000	0.130	91694000	0.115	91693000	0				
WatchDog Tick Outputsource	1	1	4000	4000	0.000	87.615	39000	89.375	40000	0				
WatchDog Update Outputsource	154	154	14000	14000	0.002	0.101	93306000	0.088	93311000	0				

Topologies

Name	Id	Owner	Status	Uptime	Num workers	Num executors	Num tasks	Replication count	Assigned Mem (MB)	Scheduler Info
RAMP_INSTRUCTIONS_PROCESSOR-football-119913_102	RAMP_INSTRUCTIONS_PROCESSOR-football-119913_102-48-1556614523	storm	ACTIVE	27d 6h 59m 59s	22	1774	1774	3	18304	

Topology stats						
Window	Emitted	Transferred	Complete latency (ms)	Acked	Failed	
10m 0s	72982	82778	6.600	10000		
3h 0m 0s	1380836	1601501	6.173	229306		
1d 0h 0m 0s	15330332	17881216	11.148	2616412		
All time	557816000	649495000	9.753	94589000		

Spouts (All time)												
Id	Executors	Tasks	Emitted	Transferred	Complete latency (ms)	Acked	Failed	Error Host	Error Port	Last error	Error Time	
Instructions	50	50	94806000	94806000	9.753	94589000	0					
Replication Instructions	50	50	0	0	0.000	0	0					

Showing 1 to 2 of 2 entries

Bolts (All time)														
Id	Executors	Tasks	Emitted	Transferred	Capacity (last 10m)	Execute latency (ms)	Executed	Process latency (ms)	Acked	Failed	Error Host	Error Port	Last error	Error Time
CircuitBreaker Outputsource	1	1	14000	14000	0.000	17.401	157000	17.197	157000	0				
Dual Mapping	154	154	94582000	94582000	0.002	0.036	94602000	0.028	94597000	0				
ES Logger Outputsource	44	44	0	0	0.000	0.235	91689000	0.227	91691000	0				
FMM Create Outputsource	44	44	0	0	0.165	195.152	1759000	216.619	1763000	0				
FMM Refresh Outputsource	88	88	0	0	0.000	96.442	52000	109.000	52000	0				
FMM Result Outputsource	88	88	0	0	0.147	131.147	1131000	123.818	1128000	0				
FMM Update Outputsource	836	836	0	0	0.187	67.665	88795000	67.729	88784000	0				
Instructions Processor	154	154	368396000	460075000	0.008	3.461	94617000	3.463	94615000	0				
PPB Outputsource	44	44	0	0	0.002	0.201	91683000	0.185	91678000	0				
SPORTEX Outputsource	44	44	0	0	0.000	0.130	91694000	0.115	91693000	0				
WatchDog Tick Outputsource	1	1	4000	4000	0.000	87.615	39000	89.375	40000	0				
WatchDog Update Outputsource	154	154	14000	14000	0.002	0.101	93306000	0.088	93311000	0				

Topologies

Name	Id	Owner	Status	Uptime	Num workers	Num executors	Num tasks	Replication count	Assigned Mem (MB)	Scheduler Info
RAMP_INSTRUCTIONS_PROCESSOR-football-119913_102	RAMP_INSTRUCTIONS_PROCESSOR-football-119913_102-48-1556614523	storm	ACTIVE	27d 6h 59m 59s	22	1774	1774	3	18304	

Topology stats

Window	Emitted	Transferred	Complete latency (ms)	Acked	Failed
10m 0s	72982	82778	6.600	10000	
3h 0m 0s	1380836	1601501	6.173	229306	
1d 0h 0m 0s	15330332	17881216	11.148	2616412	
All time	557816000	649495000	9.753	94589000	

Spouts (All time)

Search:

Id	▲ Executors	◆ Tasks	◆ Emitted	◆ Transferred	◆ Complete latency (ms)	◆ Acked	◆ Failed	◆ Error Host	◆ Error Port	◆ Last error	◆ Error Time
Instructions	50	50	94806000	94806000	9.753	94589000	0				
Replication Instructions	50	50	0	0	0.000	0	0				

Bolts (All time)														
Search: <input type="text"/>														
Id	Executors	Tasks	Emitted	Transferred	Capacity (last 10m)	Execute latency (ms)	Executed	Process latency (ms)	Acked	Failed	Error Host	Error Port	Last error	Error Time
CircuitBreaker Outputsource	1	1	14000	14000	0.000	17.401	157000	17.197	157000	0				
Dual Mapping	154	154	94582000	94582000	0.002	0.036	94602000	0.028	94597000	0				
ES Logger Outputsource	44	44	0	0	0.000	0.235	91689000	0.227	91691000	0				
FMM Create Outputsource	44	44	0	0	0.165	195.152	1759000	216.619	1763000	0				
FMM Refresh Outputsource	88	88	0	0	0.000	96.442	52000	109.000	52000	0				
FMM Result Outputsource	88	88	0	0	0.147	131.147	1131000	123.818	1128000	0				
FMM Update Outputsource	836	836	0	0	0.187	67.665	88795000	67.729	88784000	0				
Instructions Processor	154	154	368396000	460075000	0.008	3.461	94617000	3.463	94615000	0				
PPB Outputsource	44	44	0	0	0.002	0.201	91683000	0.185	91678000	0				
SPORTEX Outputsource	44	44	0	0	0.000	0.130	91694000	0.115	91693000	0				
WatchDog Tick Outputsource	1	1	4000	4000	0.000	87.615	39000	89.375	40000	0				
WatchDog Update Outputsource	154	154	14000	14000	0.002	0.101	93306000	0.088	93311000	0				

Topologies

Name	Id	Owner	Status	Uptime	Num workers	Num executors	Num tasks	Replication count	Assigned Mem (MB)	Scheduler Info
RAMP_INSTRUCTIONS_PROCESSOR-football-119913_102	RAMP_INSTRUCTIONS_PROCESSOR-football-119913_102-48-1556614523	storm	ACTIVE	27d 6h 59m 59s	22	1774	1774	3	18304	

Topology stats

Window	Emitted	Transferred	Complete latency (ms)	Acked	Failed
10m 0s	72982	82778	6.600	10000	
3h 0m 0s	1380836	1601501	6.173	229306	
1d 0h 0m 0s	15330332	17881216	11.148	2616412	
All time	557816000	649495000	9.753	94589000	

Spouts (All time)

Search: <input type="text"/>											
Id	Executors	Tasks	Emitted	Transferred	Complete latency (ms)	Acked	Failed	Error Host	Error Port	Last error	Error Time
Instructions	50	50	94806000	94806000	9.753	94589000	0				
Replication Instructions	50	50	0	0	0.000	0	0				

Showing 1 to 2 of 2 entries

Bolts (All time)

Search: <input type="text"/>														
Id	Executors	Tasks	Emitted	Transferred	Capacity (last 10m)	Execute latency (ms)	Executed	Process latency (ms)	Acked	Failed	Error Host	Error Port	Last error	Error Time
CircuitBreaker Outputsource	1	1	14000	14000	0.000	17.401	157000	17.197	157000	0				
Dual Mapping	154	154	94582000	94582000	0.002	0.036	94602000	0.028	94597000	0				
ES Logger Outputsource	44	44	0	0	0.000	0.235	91689000	0.227	91691000	0				
FMM Create Outputsource	44	44	0	0	0.165	195.152	1759000	216.619	1763000	0				
FMM Refresh Outputsource	88	88	0	0	0.000	96.442	52000	109.000	52000	0				
FMM Result Outputsource	88	88	0	0	0.147	131.147	1131000	123.818	1128000	0				
FMM Update Outputsource	836	836	0	0	0.187	67.665	88795000	67.729	88784000	0				
Instructions Processor	154	154	368396000	460075000	0.008	3.461	94617000	3.463	94615000	0				
PPB Outputsource	44	44	0	0	0.002	0.201	91683000	0.185	91678000	0				
SPORTEX Outputsource	44	44	0	0	0.000	0.130	91694000	0.115	91693000	0				
WatchDog Tick Outputsource	1	1	4000	4000	0.000	87.615	39000	89.375	40000	0				
WatchDog Update Outputsource	154	154	14000	14000	0.002	0.101	93306000	0.088	93311000	0				

Reliability

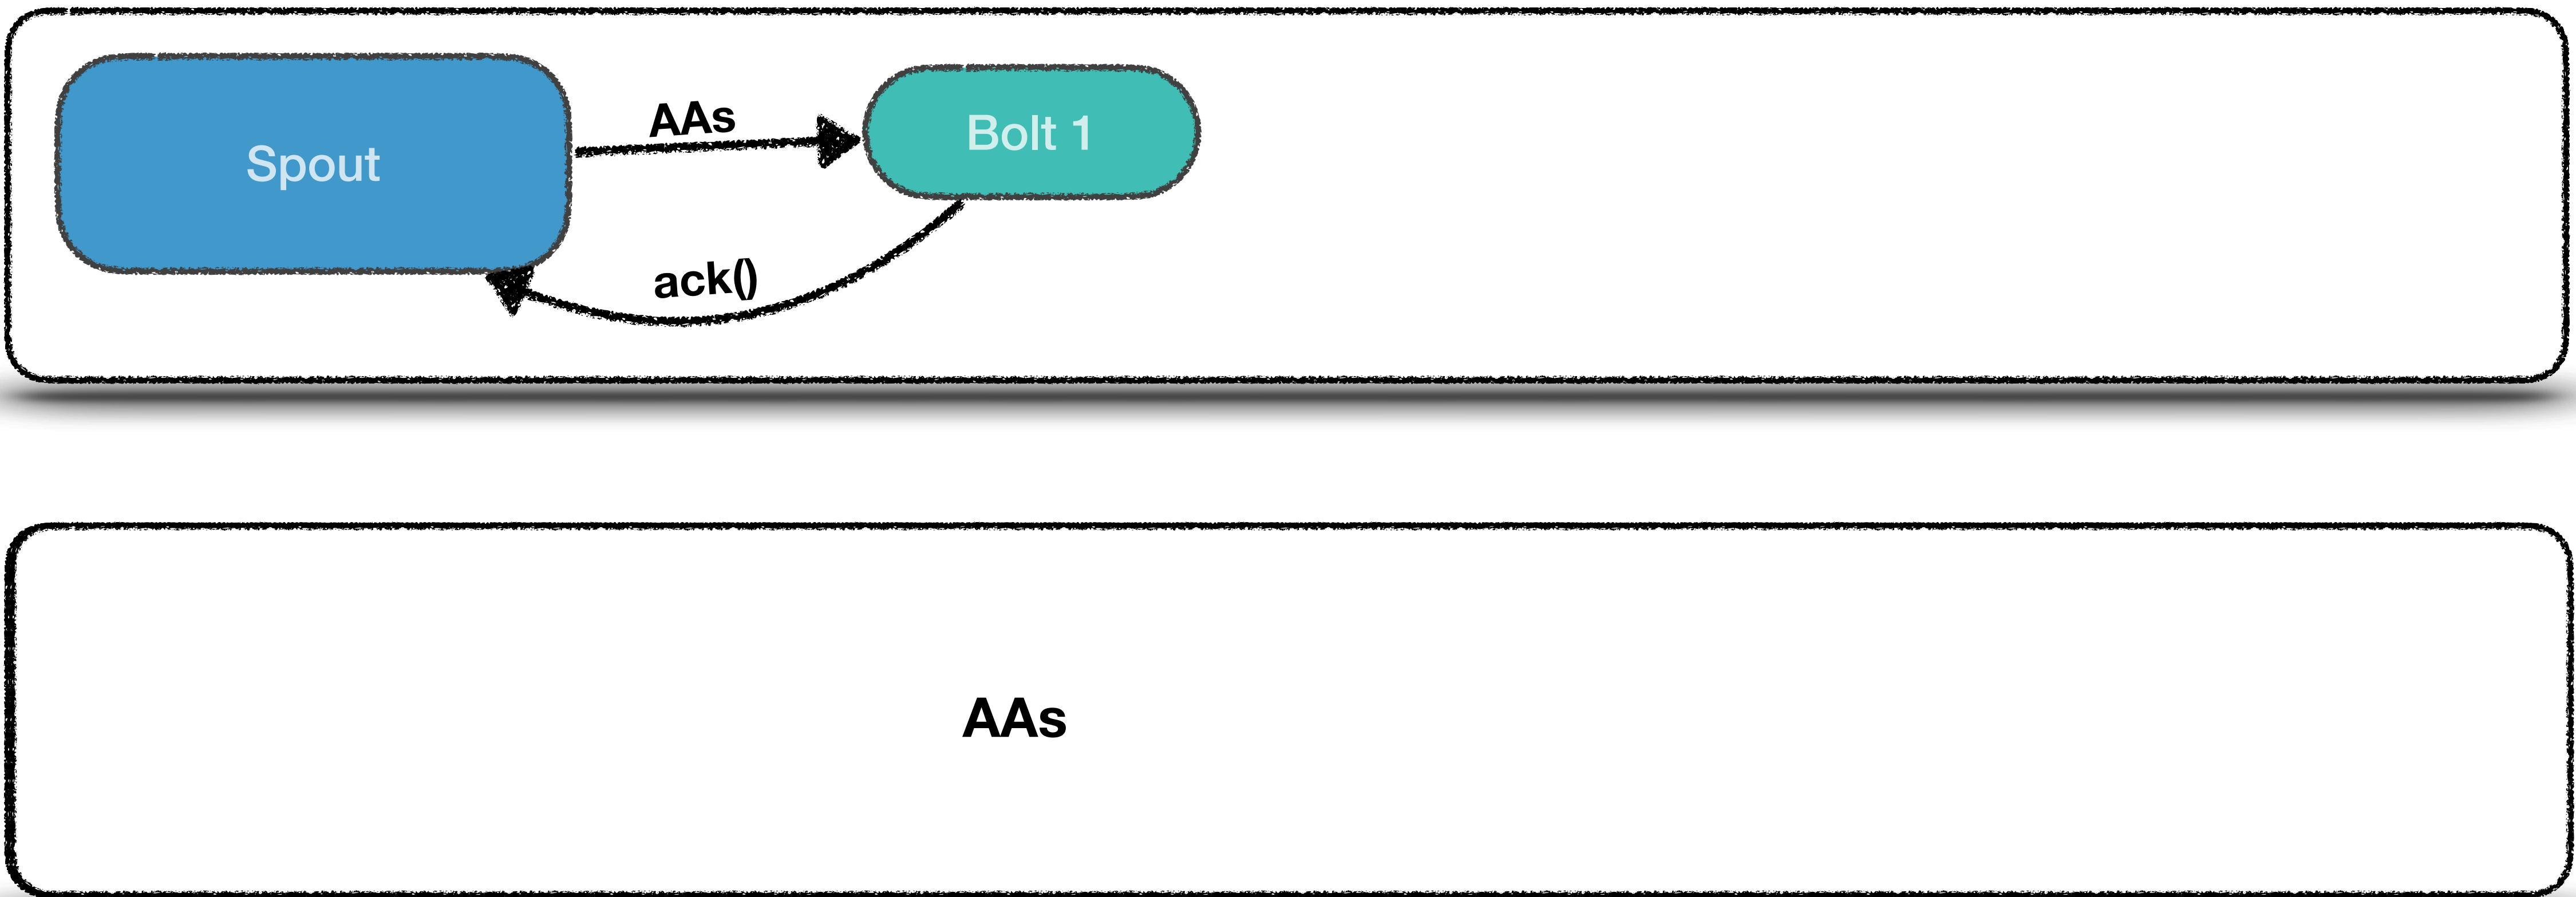
In order to have some level of **reliability**, one needs:

- A fault-tolerant Storm cluster infrastructure;
- Guarantee message processing;
 - A reliable spout
 - An anchored tuple stream
- To ack each processed tuple or notifies of failure;

Reliability

```
trait ISpout extends Serializable {
    def open(var1: Map[String, Any], var2: TopologyContext, var3: SpoutOutputCollector): Unit
    def close(): Unit
    def activate(): Unit
    def deactivate(): Unit
    def nextTuple(): Unit
    def ack(var1: Any): Unit
    def fail(var1: Any): Unit
}
```

Reliability



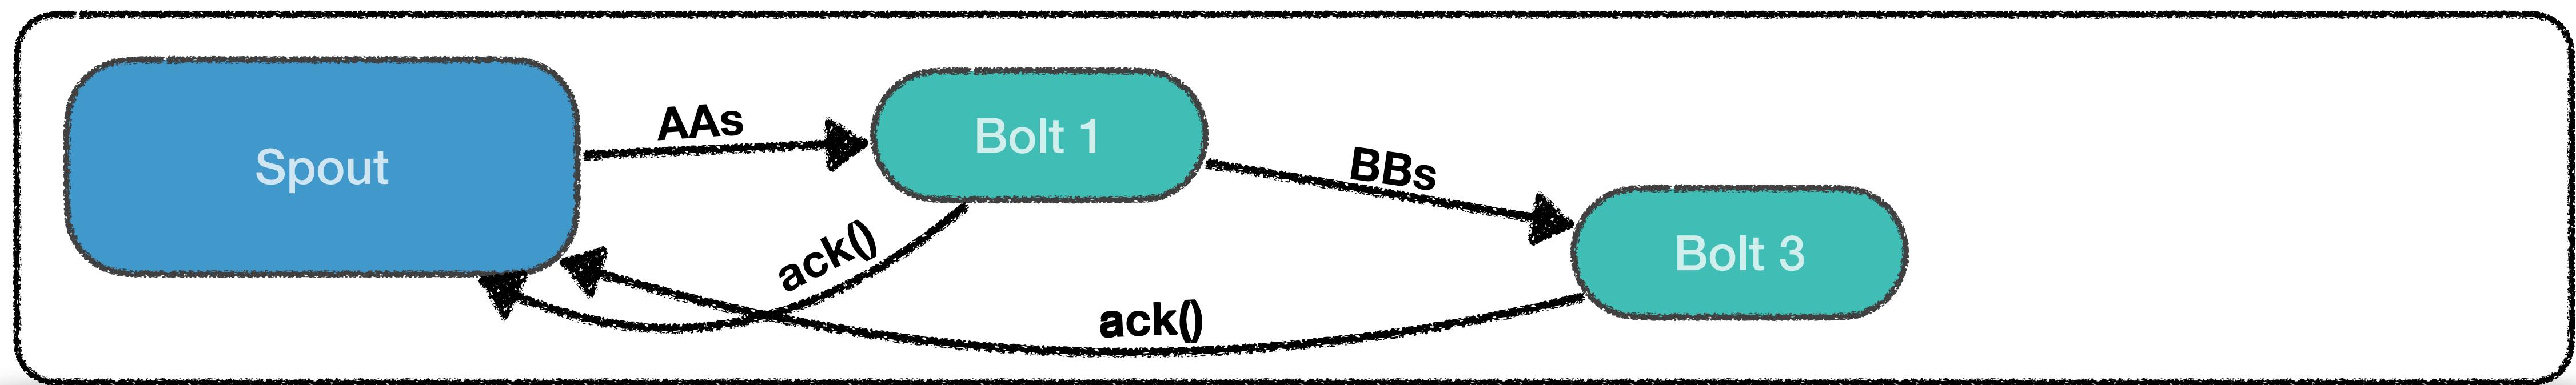
AAs

Acking & Anchoring

Ack



Reliability



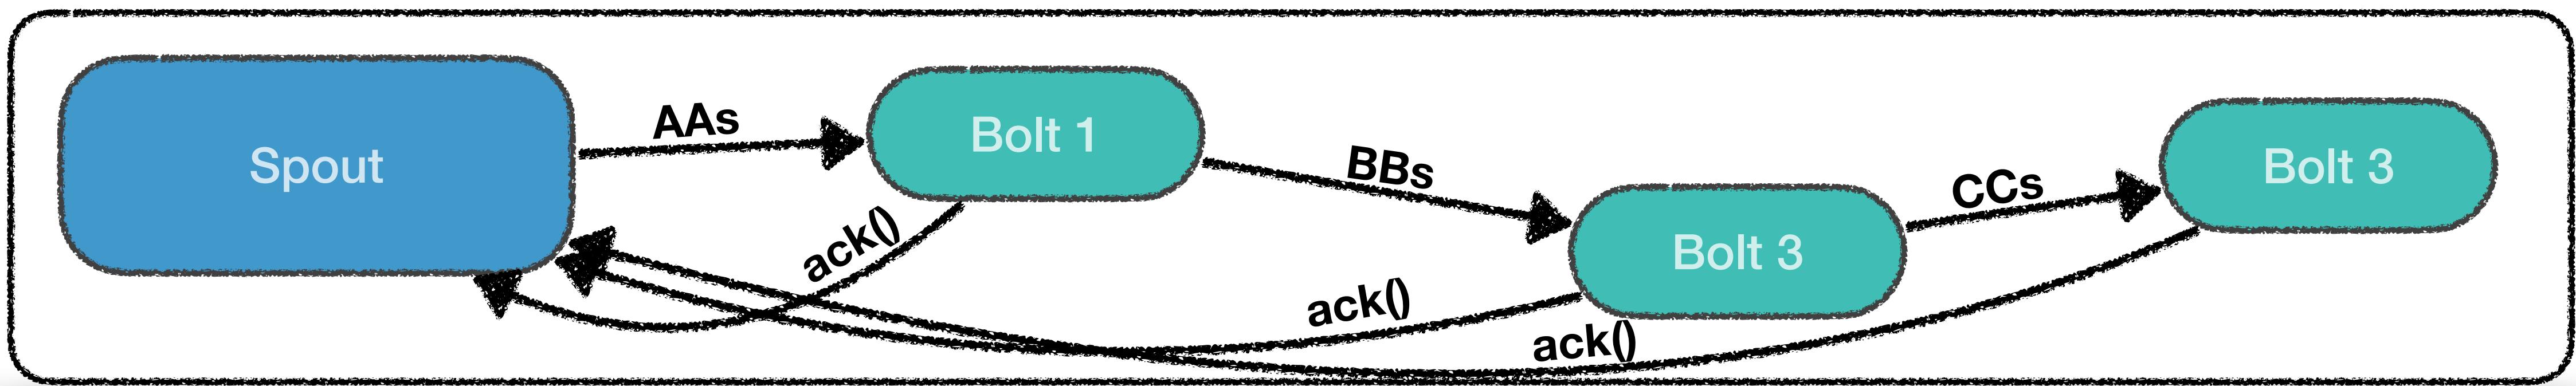
AAs

Acking & Anchoring

Ack



Reliability



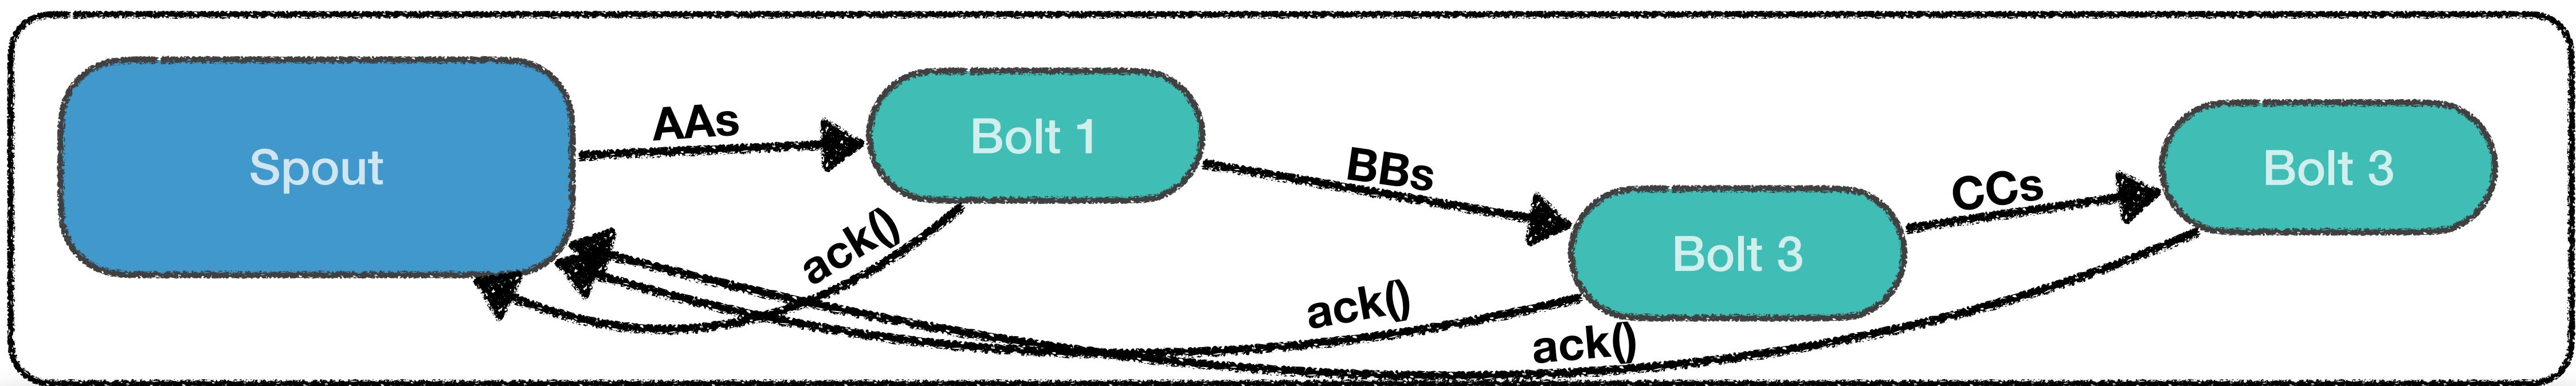
AAs

Acking & Anchoring

Ack



Reliability



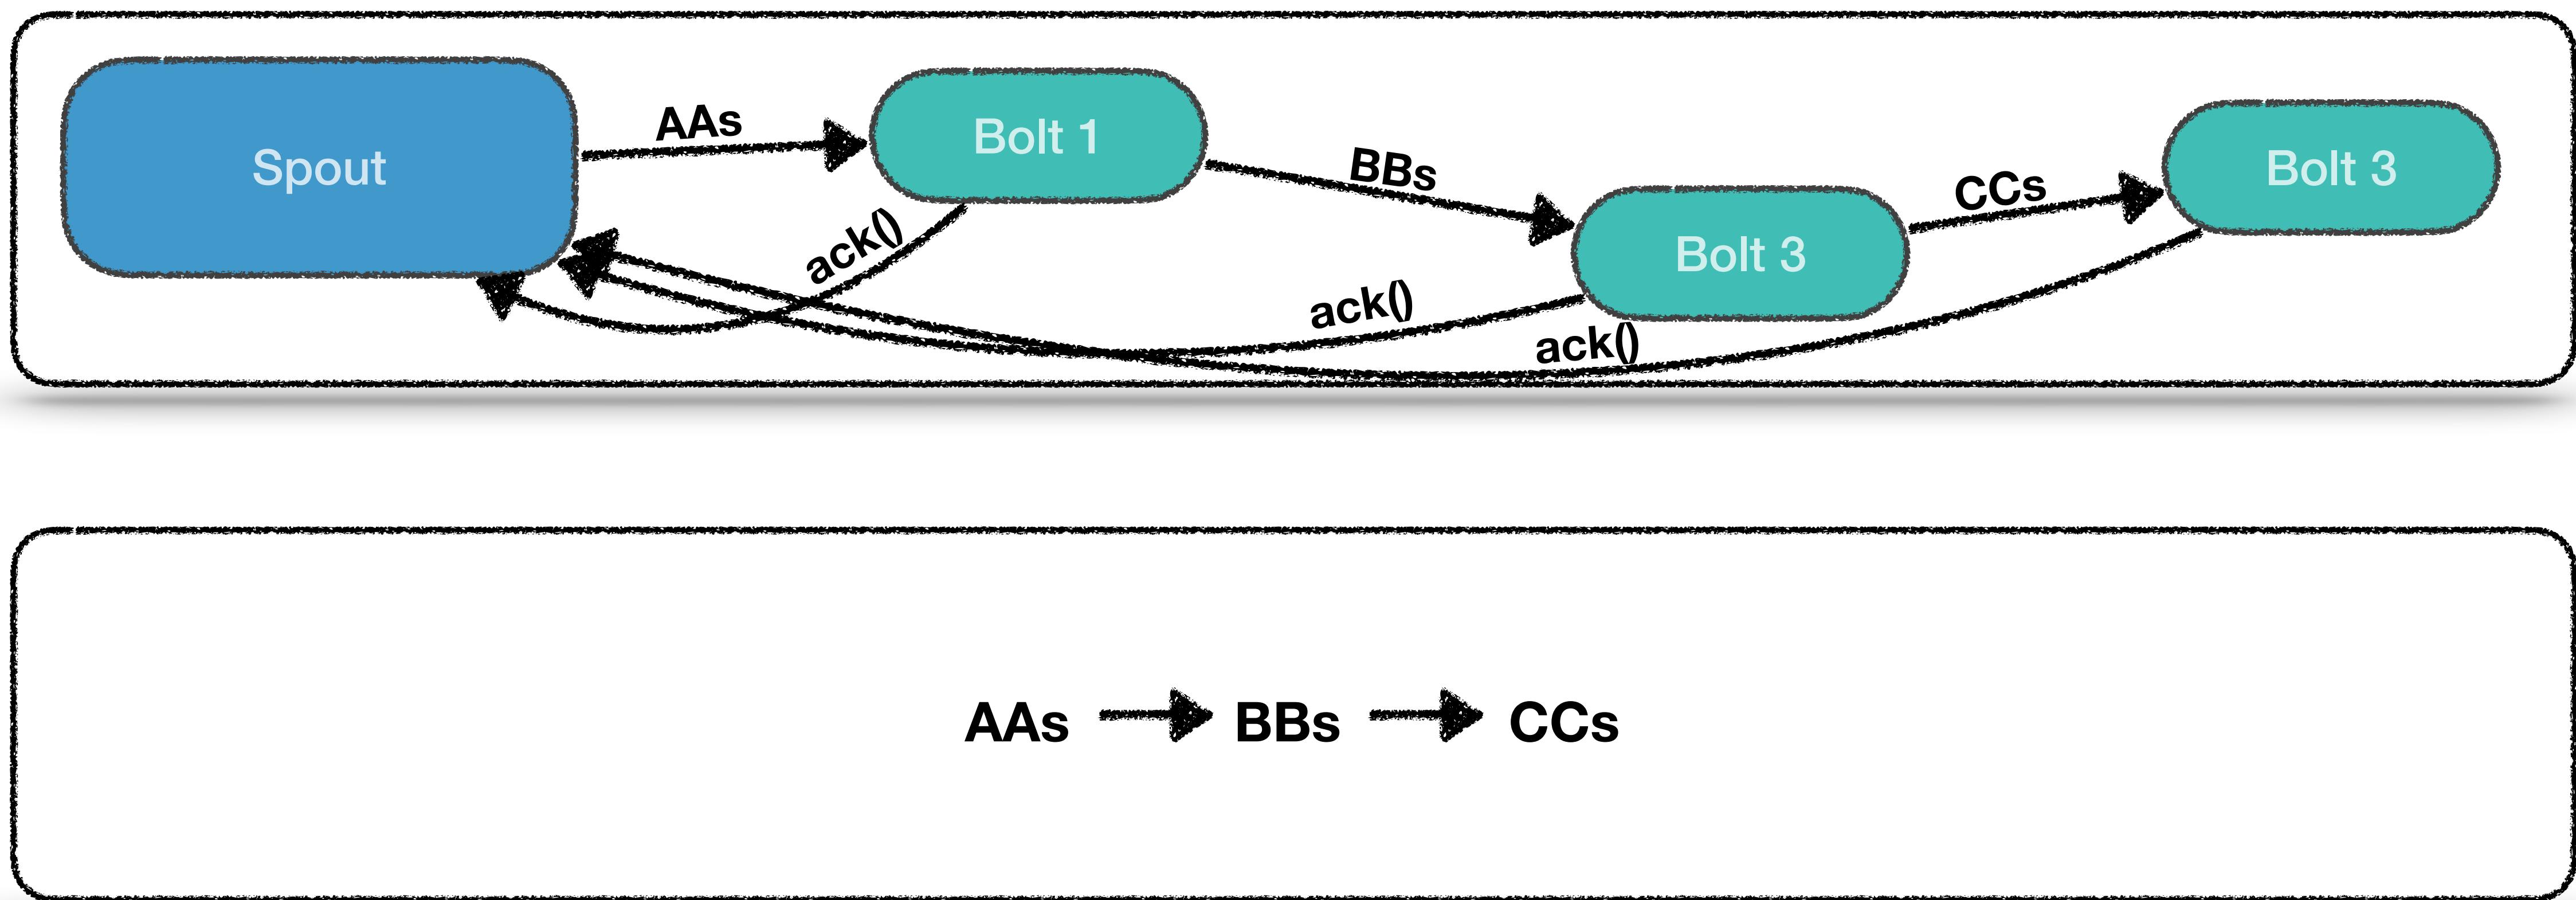
AAs → BBs

Acking & Anchoring

Ack



Reliability

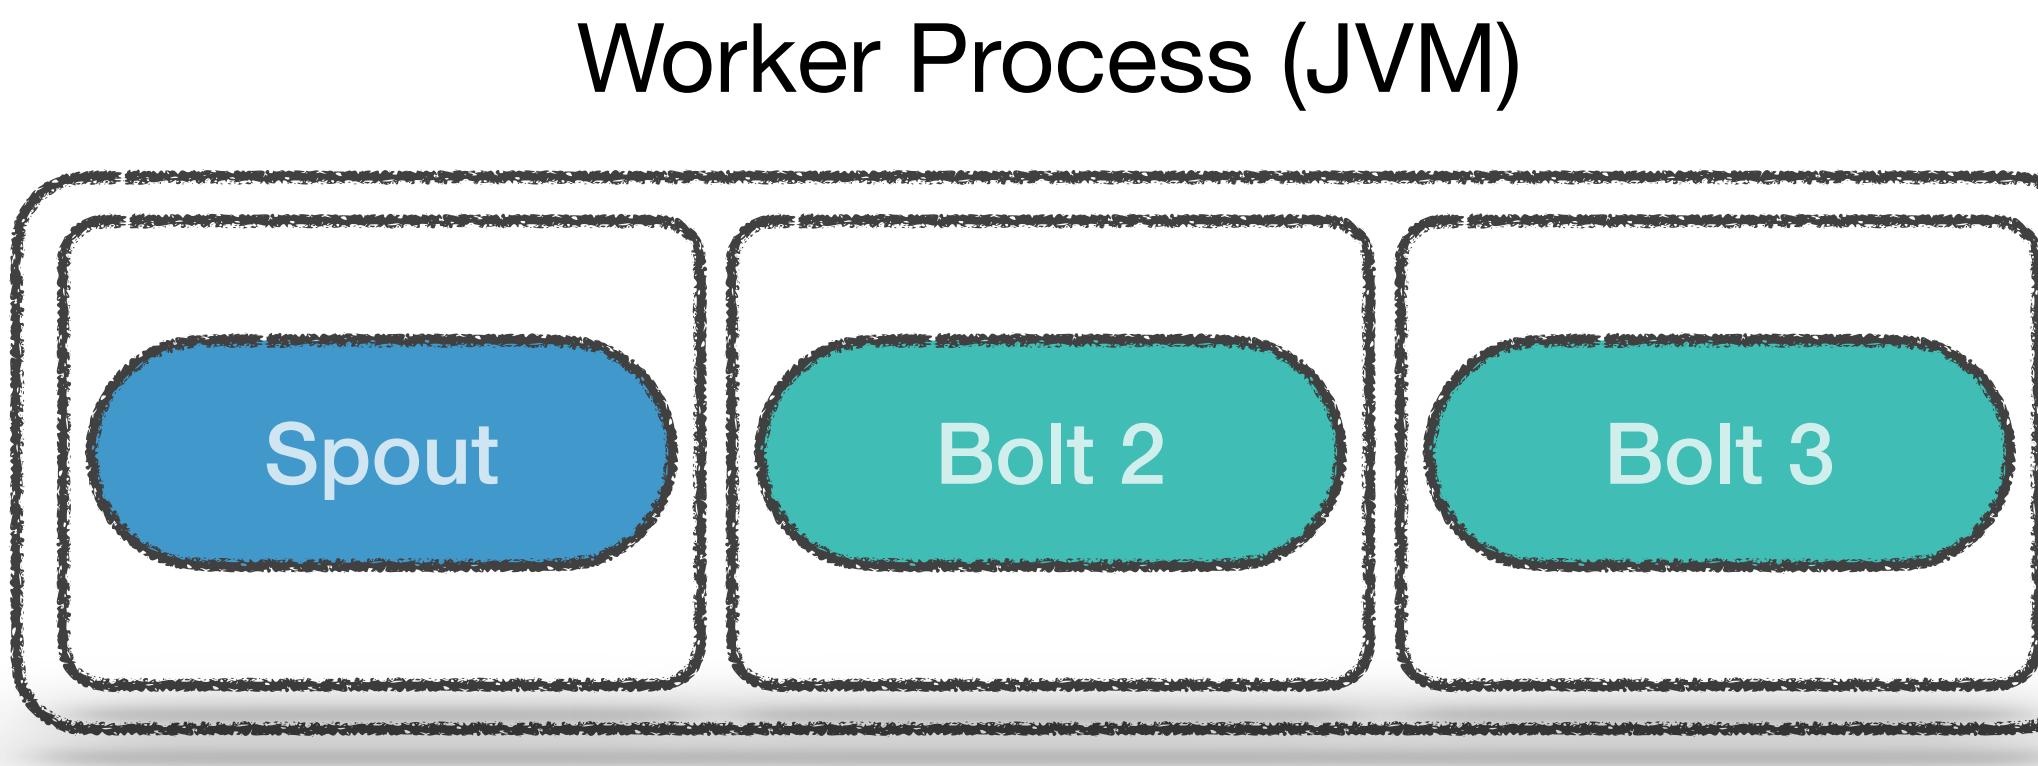


Acking & Anchoring

Ack



Parallelism



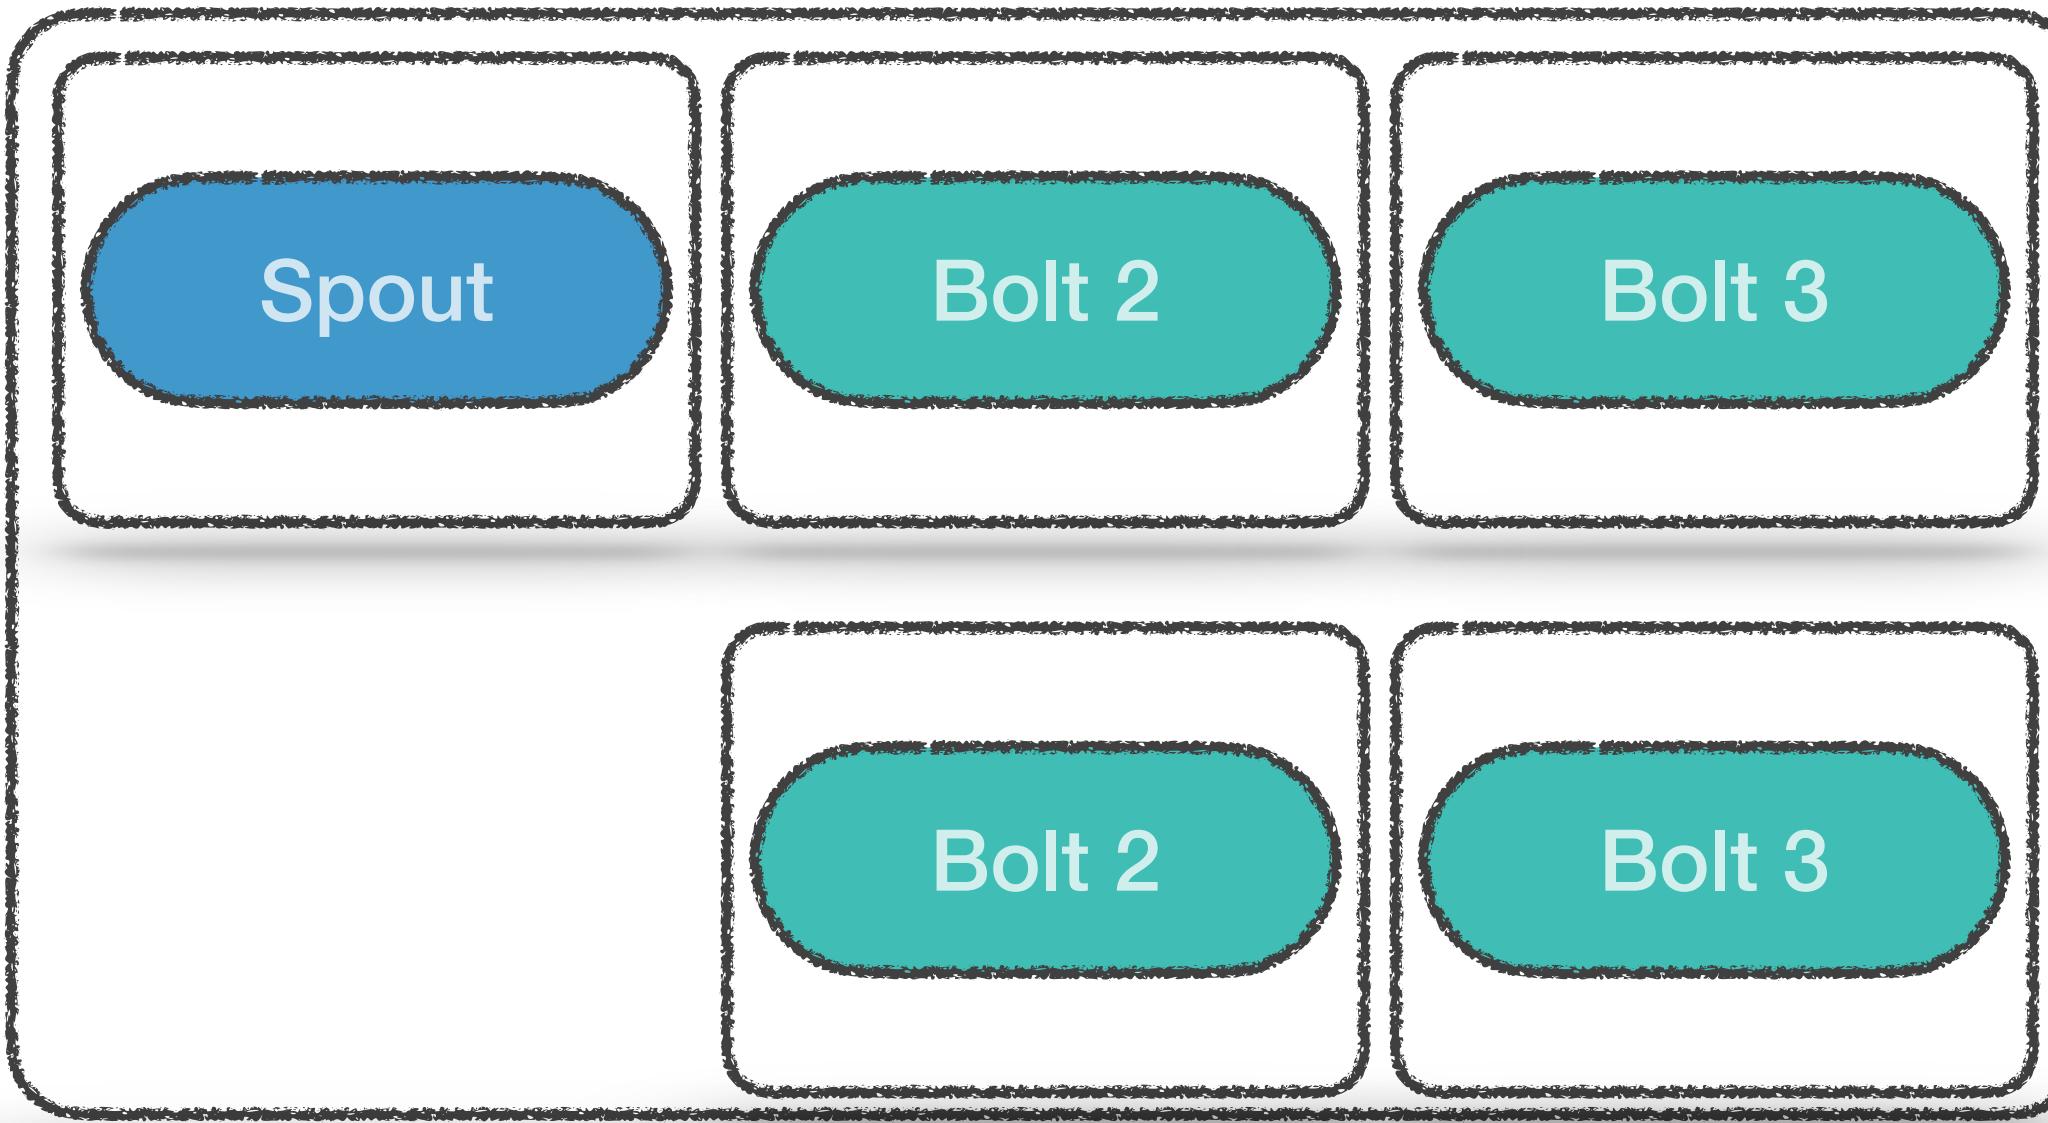
```
.setSpout("Spout Name", spout, 1)  
.setBolt("Bolt 2 Name", bolt2,1)  
.setBolt("Bolt 3 Name", bolt3,1)  
...  
cg.setNumWorkers(1)
```

parallelism hint
==
number of executors

number of worker processes

Parallelism

Worker Process (JVM)



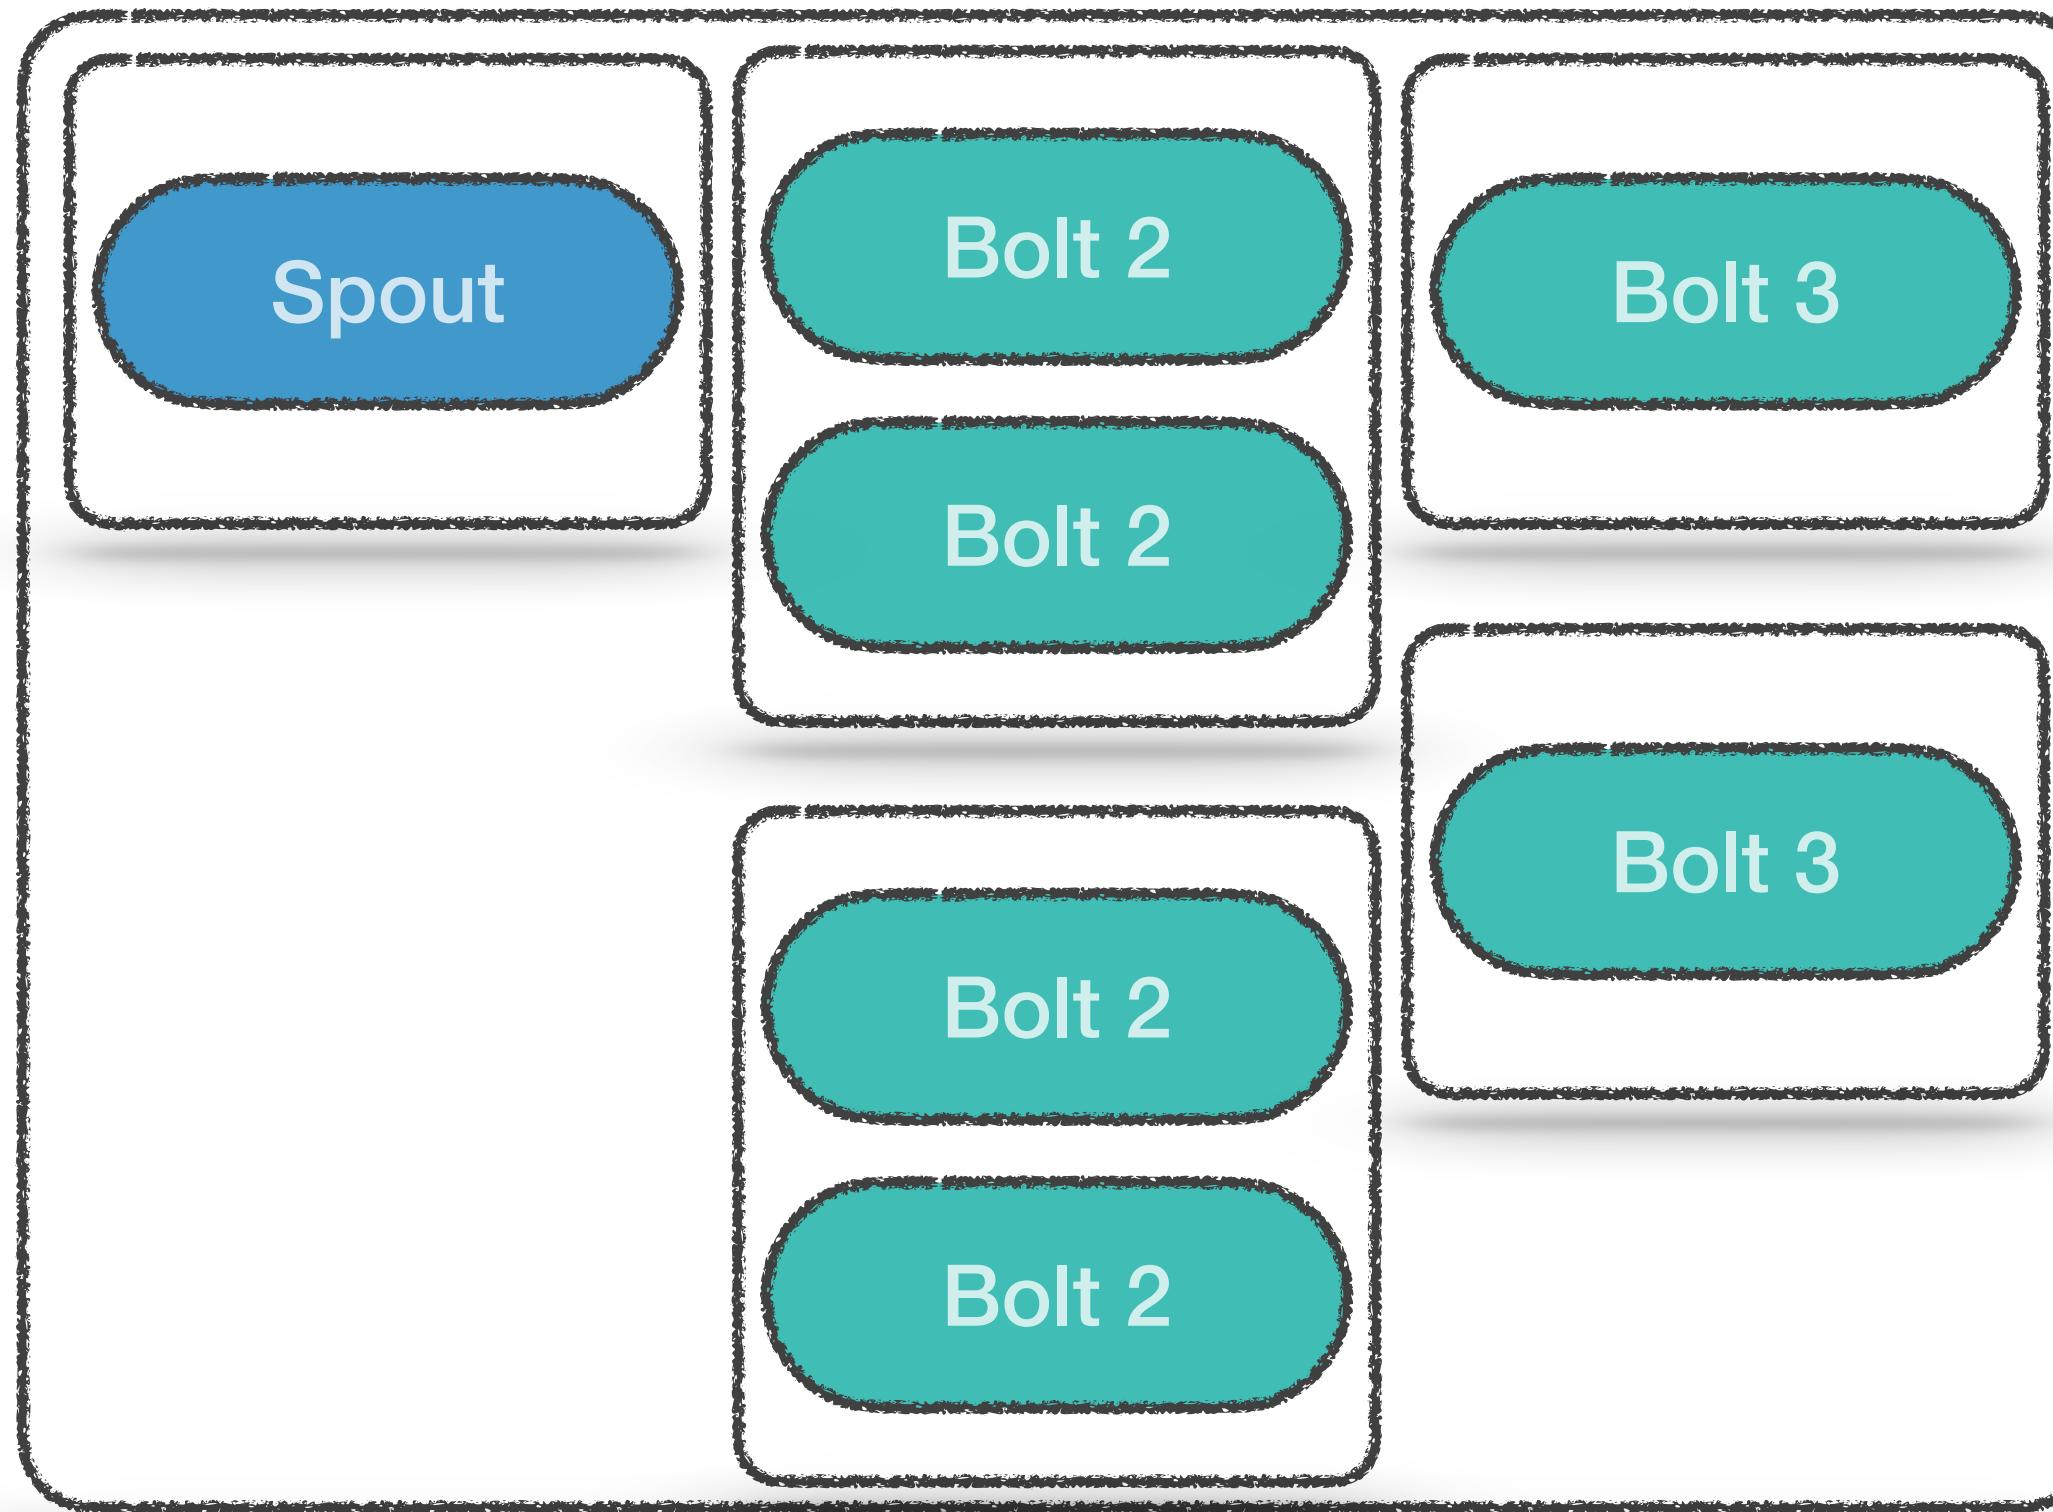
```
.setSpout("Spout Name", spout, 1)  
.setBolt("Bolt 2 Name", bolt2,2)  
.setBolt("Bolt 3 Name", bolt3,2)  
..  
cg.setNumWorkers(1)
```

parallelism hint
==
number of executors

number of worker processes

Parallelism

Worker Process (JVM)



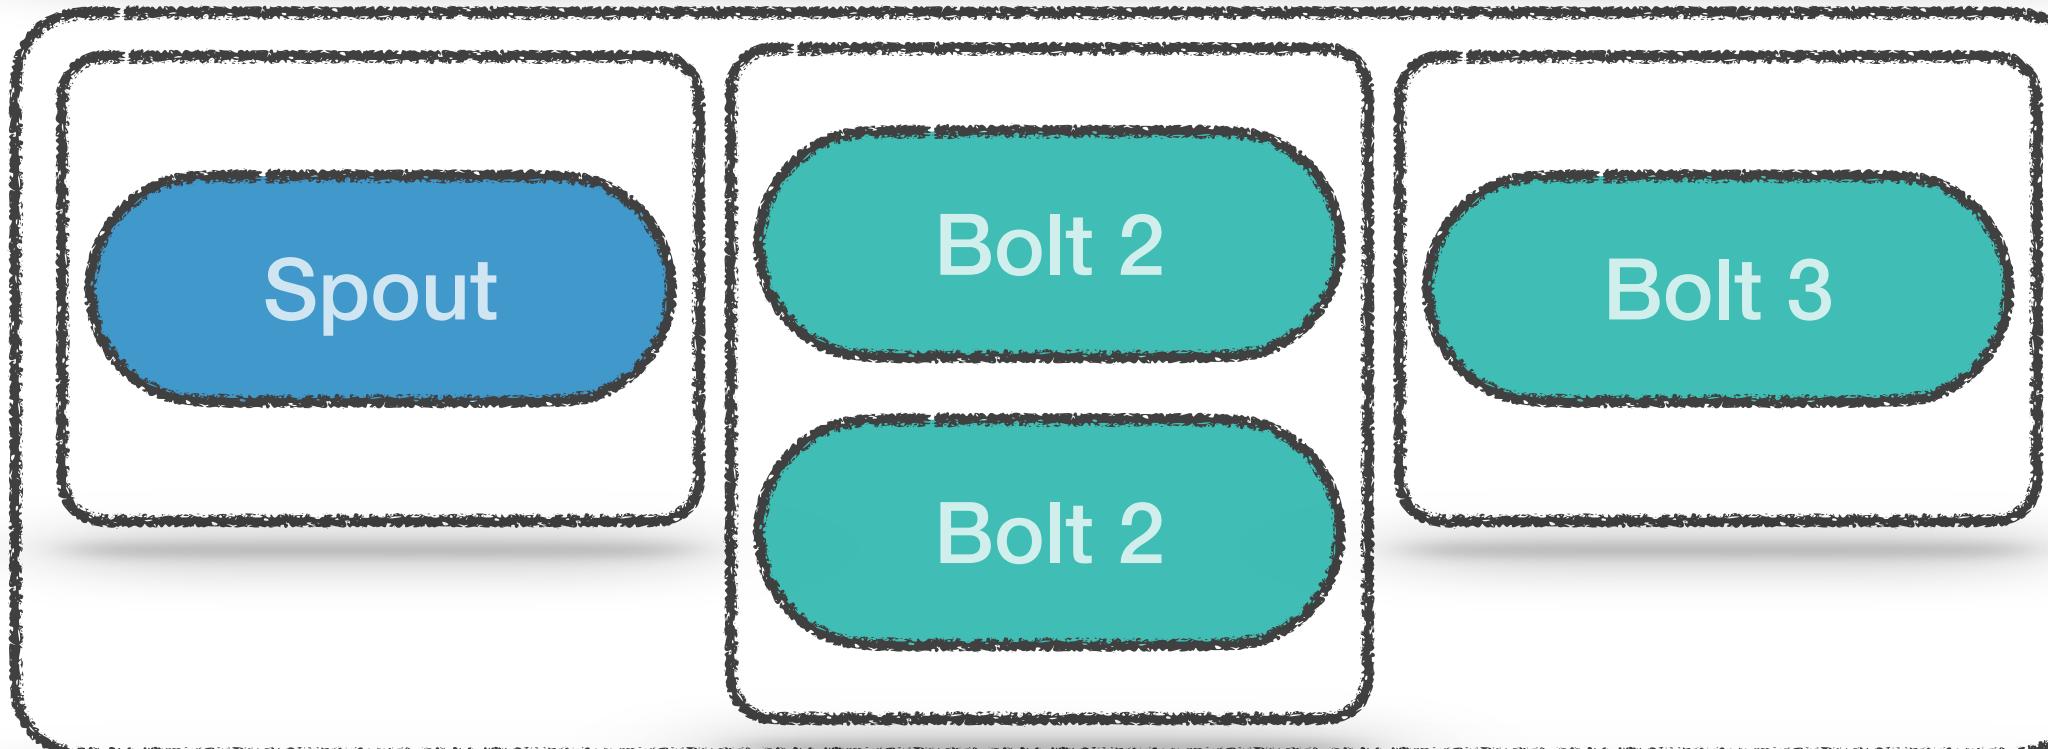
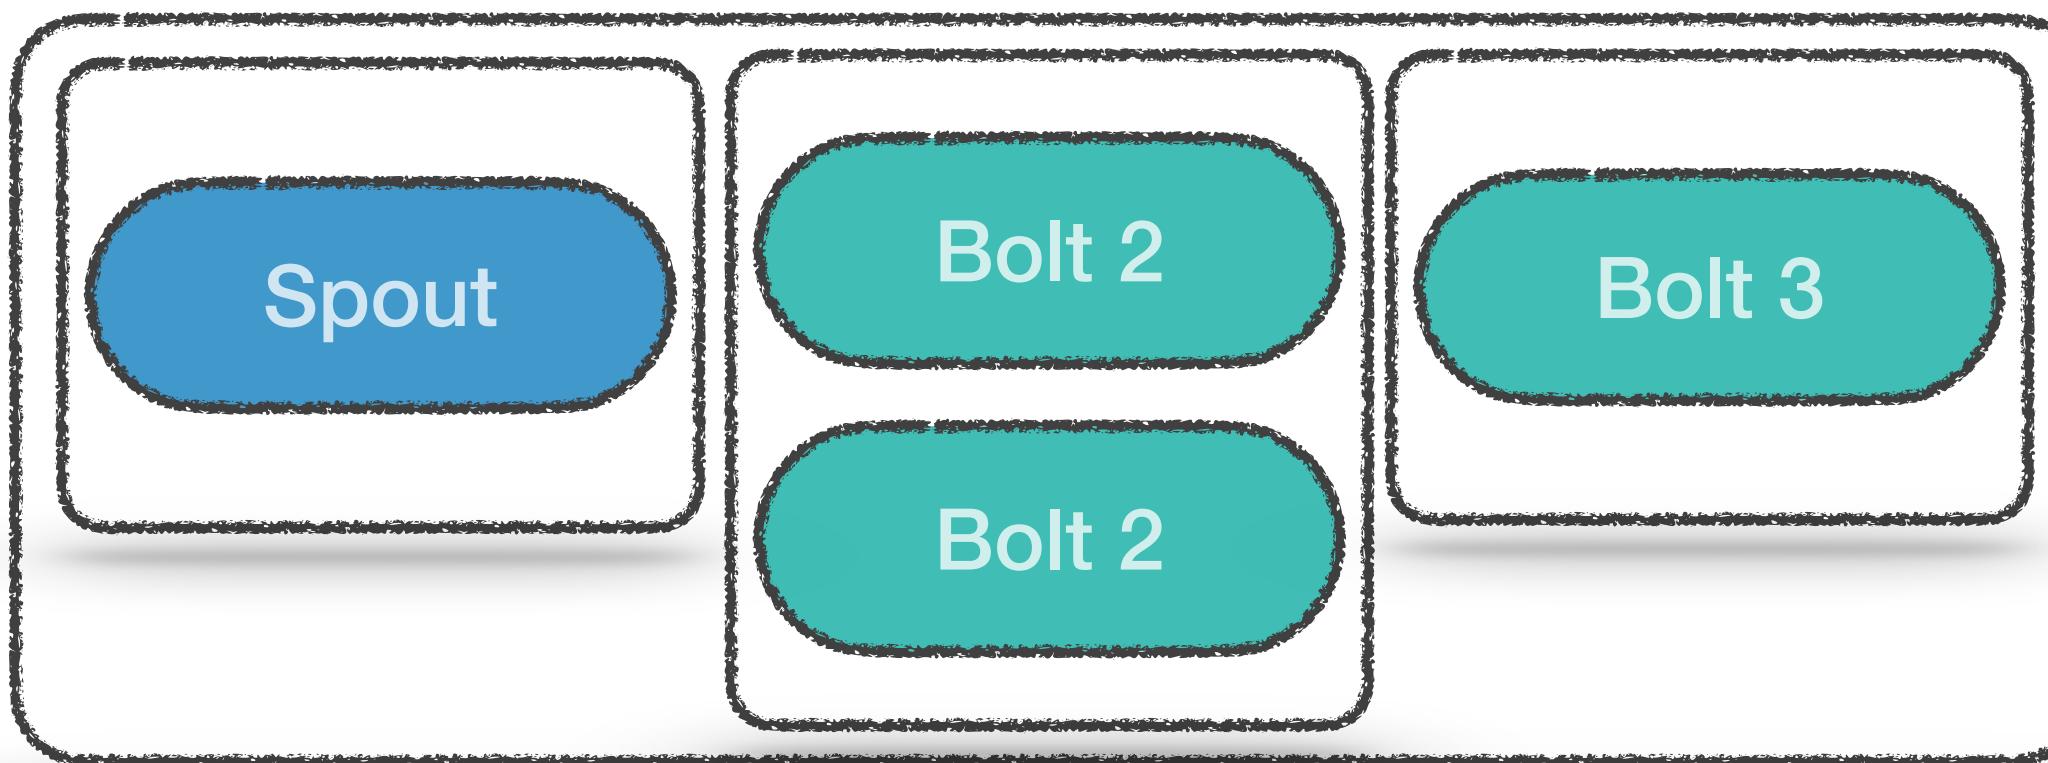
```
.setSpout("Spout Name", spout, 1)  
.setBolt("Bolt 2 Name", bolt2, 2).setNumTasks(4)  
.setBolt("Bolt 3 Name", bolt3, 2)  
"  
cg.setNumWorkers(1)
```

parallelism hint
==
number of executors

number of worker processes

Parallelism

Worker Process (JVM)



```
.setSpout("Spout Name", spout, 1)  
.setBolt("Bolt 2 Name", bolt2, 2).setNumTasks(4)  
.setBolt("Bolt 3 Name", bolt3, 2)  
...  
cg.setNumWorkers(2)
```

parallelism hint
==
number of executors

number of worker processes

Stream grouping

Defines how the tuples are sent between instances of spouts and bolts:

- **Shuffle grouping:** Tuples are randomly distributed across bolt instances;
- **Fields grouping:** The stream is partitioned by the fields specified in the grouping;
- **Partial Key grouping:** It's like Fields grouping, but are load balanced between two downstream bolts;
- **All grouping:** The stream is replicated across all bolt instances;
- **Global grouping:** The entire stream goes to a single one of the bolt's tasks;
- **Direct grouping:** The producer bolt decides which task of the consumer bolt will receive the tuple;

