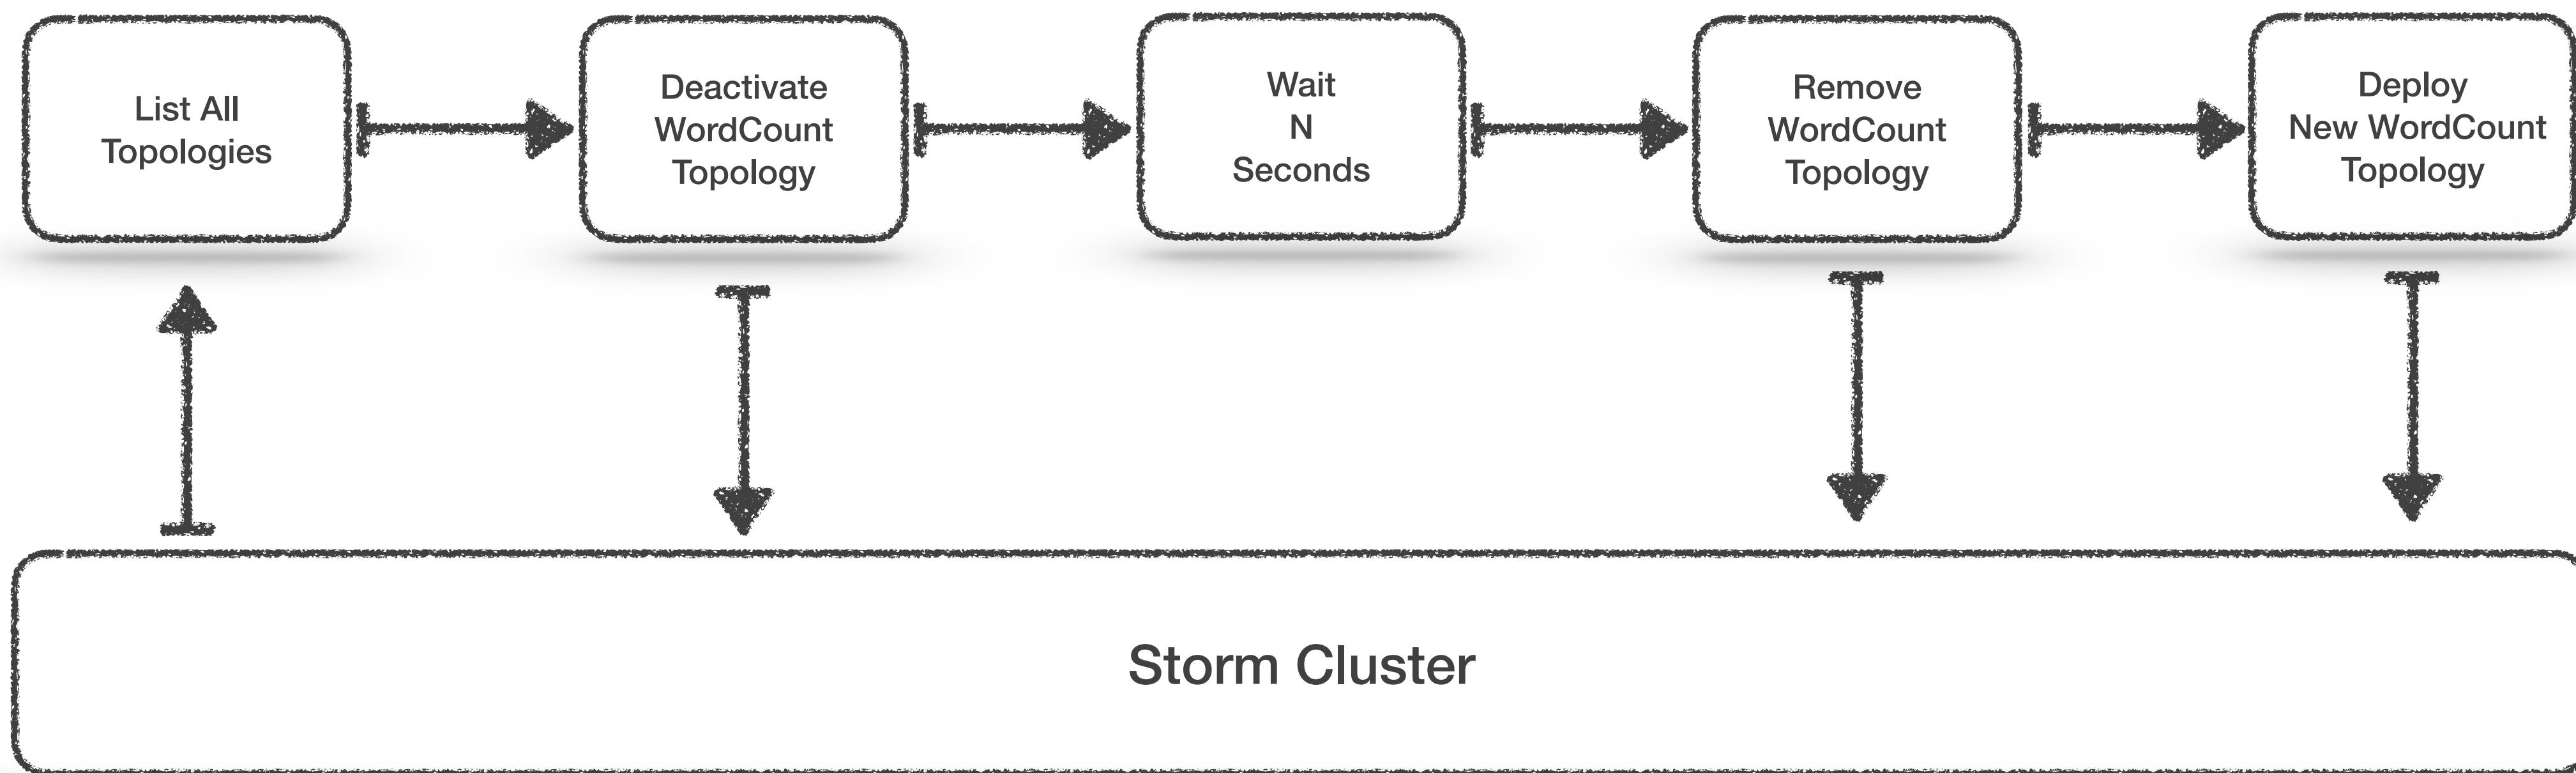


# Topology Insights

# Storm

1. Topology Deployment;
2. Dependency Management;
3. Why Scala;
4. Mock Tests;
5. Metrics.

# Topology Deployment



# Topology Deployment

## Command Line Client

```
→ apache-storm bin/storm list
11:43:00.512 [main] INFO  o.a.s.u.NimbusClient - Found leader nimbus : iopo0782.home:6627
Topology_name      Status    Num_tasks  Num_workers  Uptime_secs  Topology_Id          Owner
-----
WordCountTopology   ACTIVE     28          3            748          WordCountTopology-1-1637321432 veigam
```

```
→ apache-storm bin/storm Kill WordCountTopology
11:49:29.138 [main] INFO  o.a.s.c.Deactivate - Killed topology: WordCountTopology -w 30
```

```
→ apache-storm bin/storm jar wordCount-storm-1.0.0-SNAPSHOT.jar com.ppb.feeds.wordCount.storm.WordCountTopology
Start uploading file 'wordCount-storm-1.0.0-SNAPSHOT.jar' to '/apache-storm-2.3.0/storm-local/nimbus/inbox/
stormjar-2ab0d474-93a5-4eee-91ff-d920cd713503.jar' (18076737 bytes)
[=====] 18076737 / 18076737
```

# Topology Deployment

## Programmatically

```
object WordCountDeployer {

    def main(args: Array[String]): Unit = {

        val wordCount      = new WordCountTopology
        val nimbusClient   = NimbusClient.getConfiguredClient(wordCount.configs()).getClient
        val oldTopology    = getTopologiesName(nimbusClient, wordCount.name())

        killTopology(nimbusClient, oldTopology.head, 30)
        StormSubmitter.submitTopology(wordCount.name(), wordCount.configs(), wordCount.create())
    }

    def getTopologiesName(nimbusClient: Nimbus.Client, prefixName: String): Seq[String] = {...}
    def killTopology(nimbusClient: Nimbus.Client, name: String, waitBeforeKill: Int): Try[String] = Try {...}
}
```

[/Flutter-Global/topology-management-lib](#)

# Dependency Management

**Inversion of Control (IOC)** is a software design principle where the responsibility of controlling the application components and its flow is transferred to external containers or frameworks.

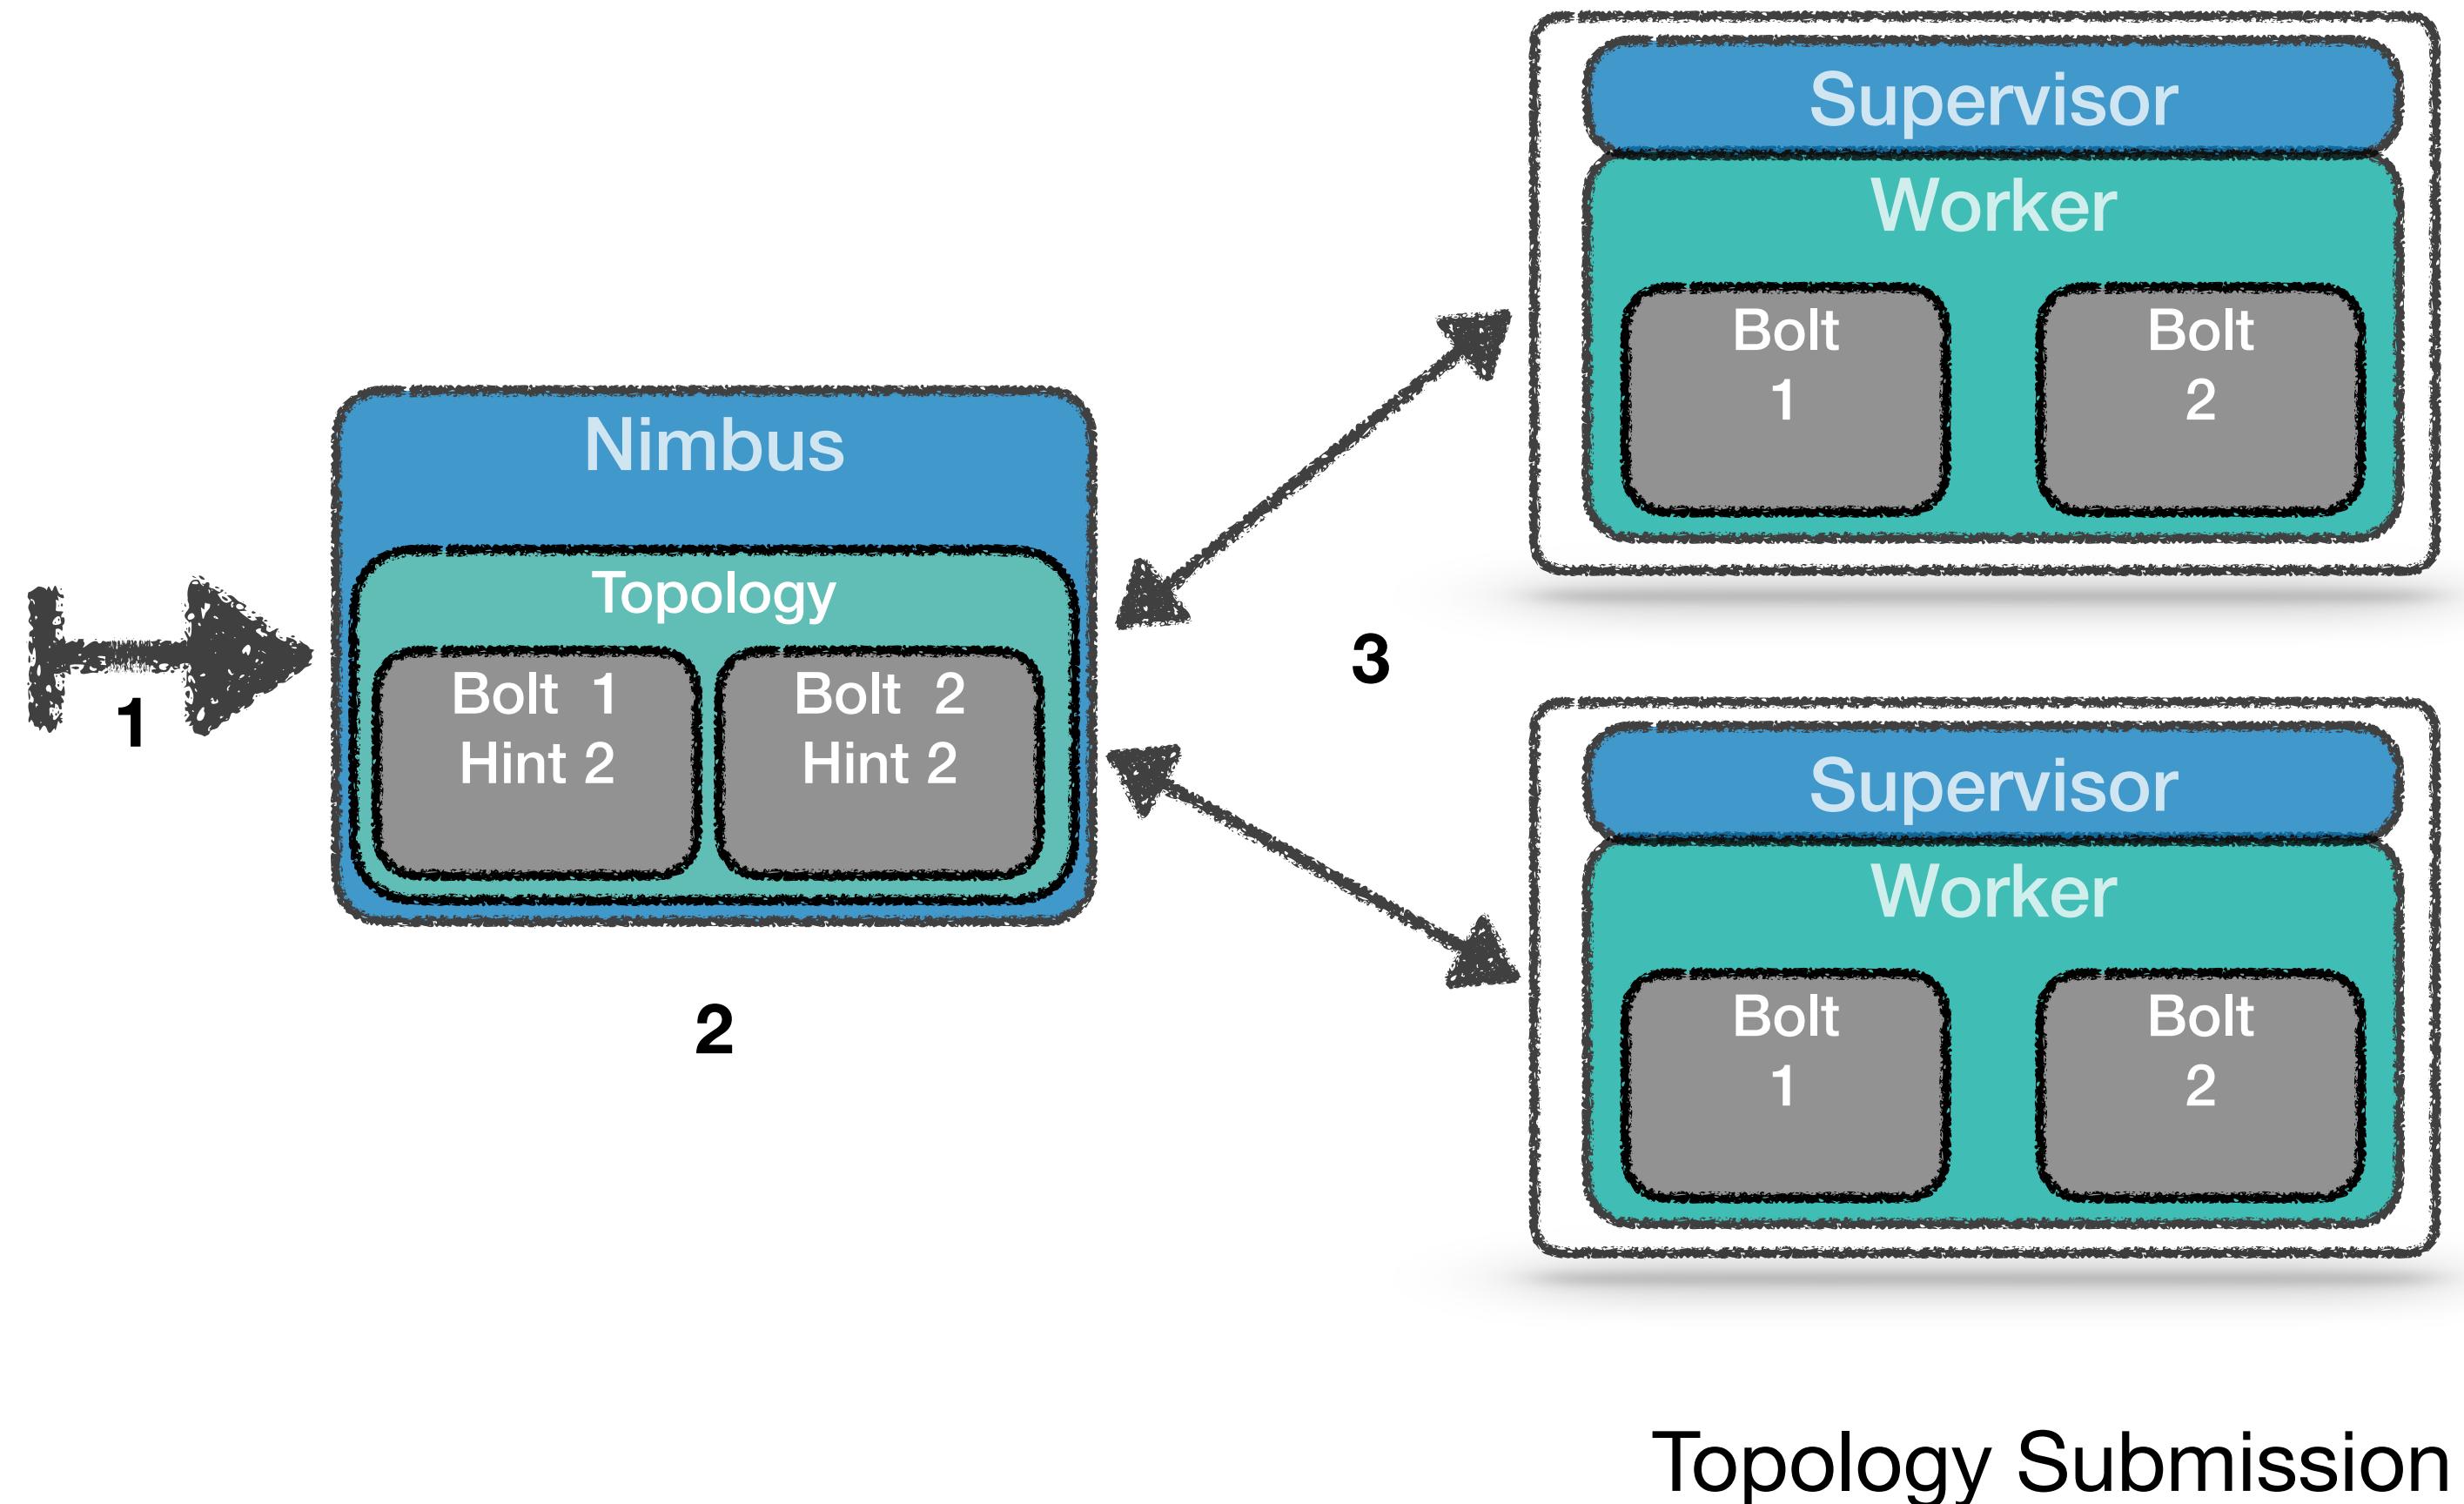
This allows:

- Decoupling the execution of a component from its implementation;
- Better switching between different implementations;
- Greater modularity by breaking an application into isolated components;
- Easier to test.

# Dependency Management

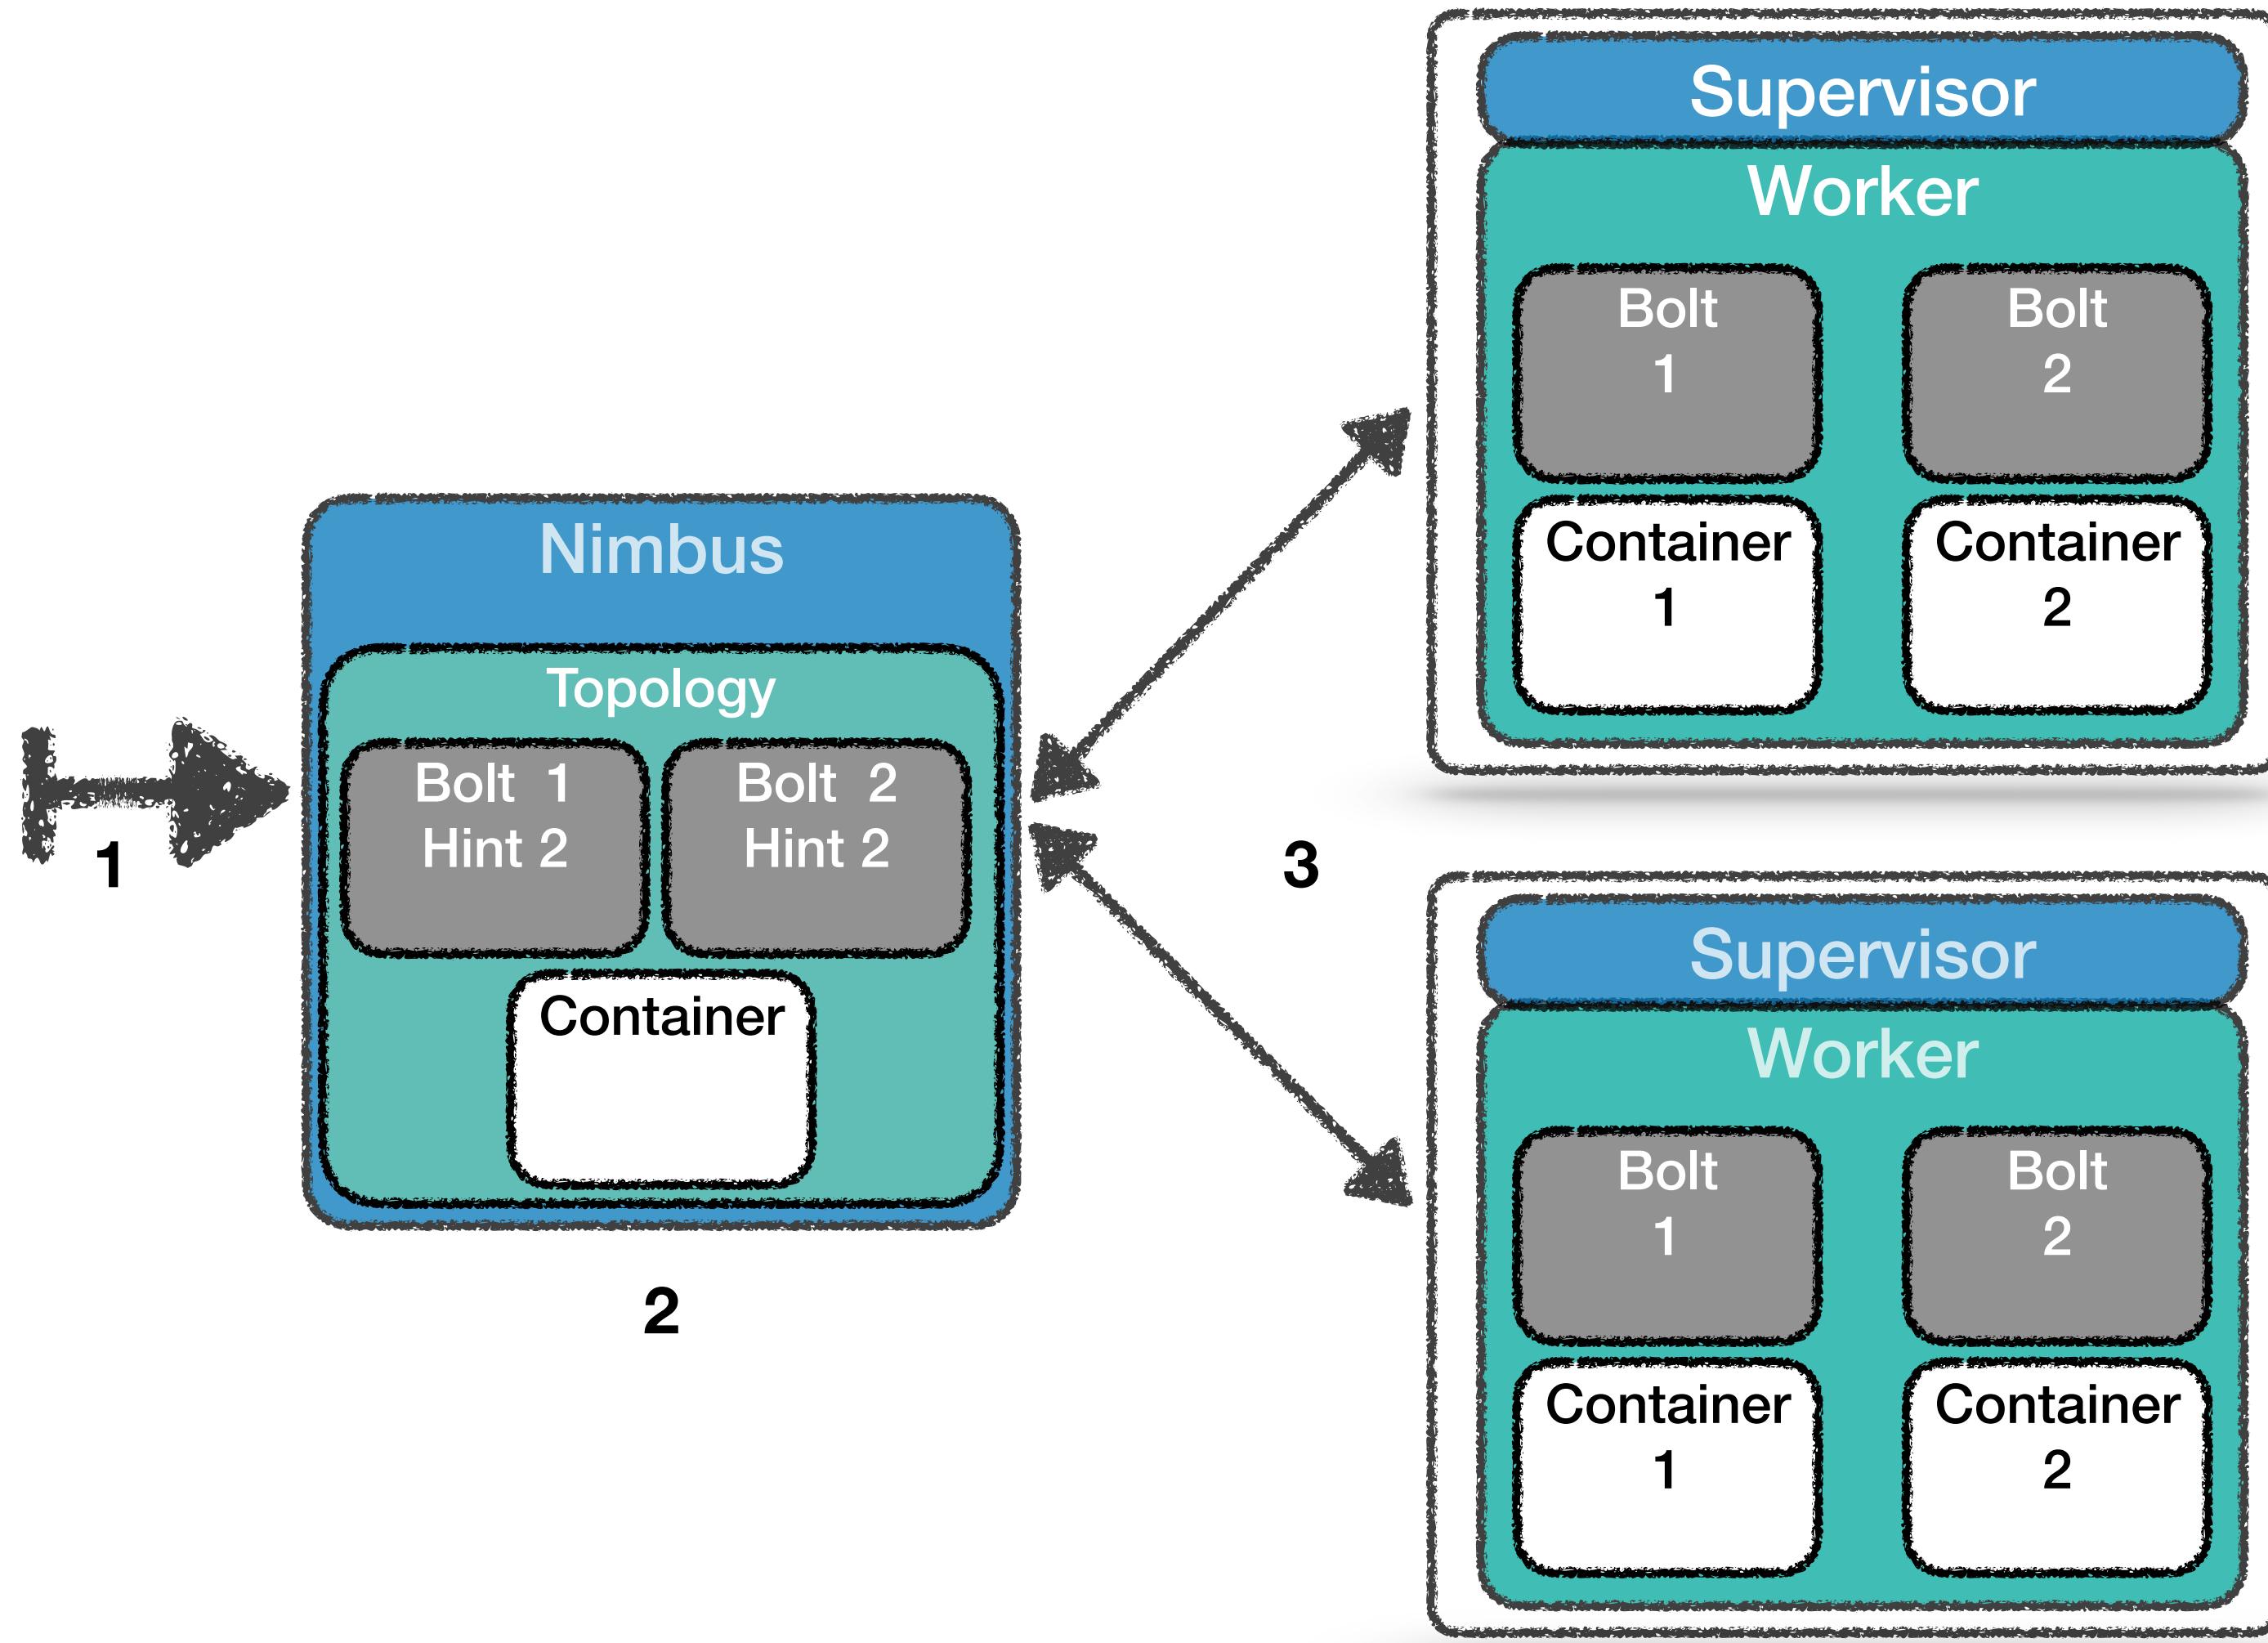
**Dependency injection (DI)** is a design pattern that can be used to apply IOC. In DI, dependencies of an object are provided (injected) by means of an orchestrator instead of having the object construct them itself. Normally, this can be achieved with the help of IOC Container frameworks like Spring or libraries like Guice.

# Dependency Management



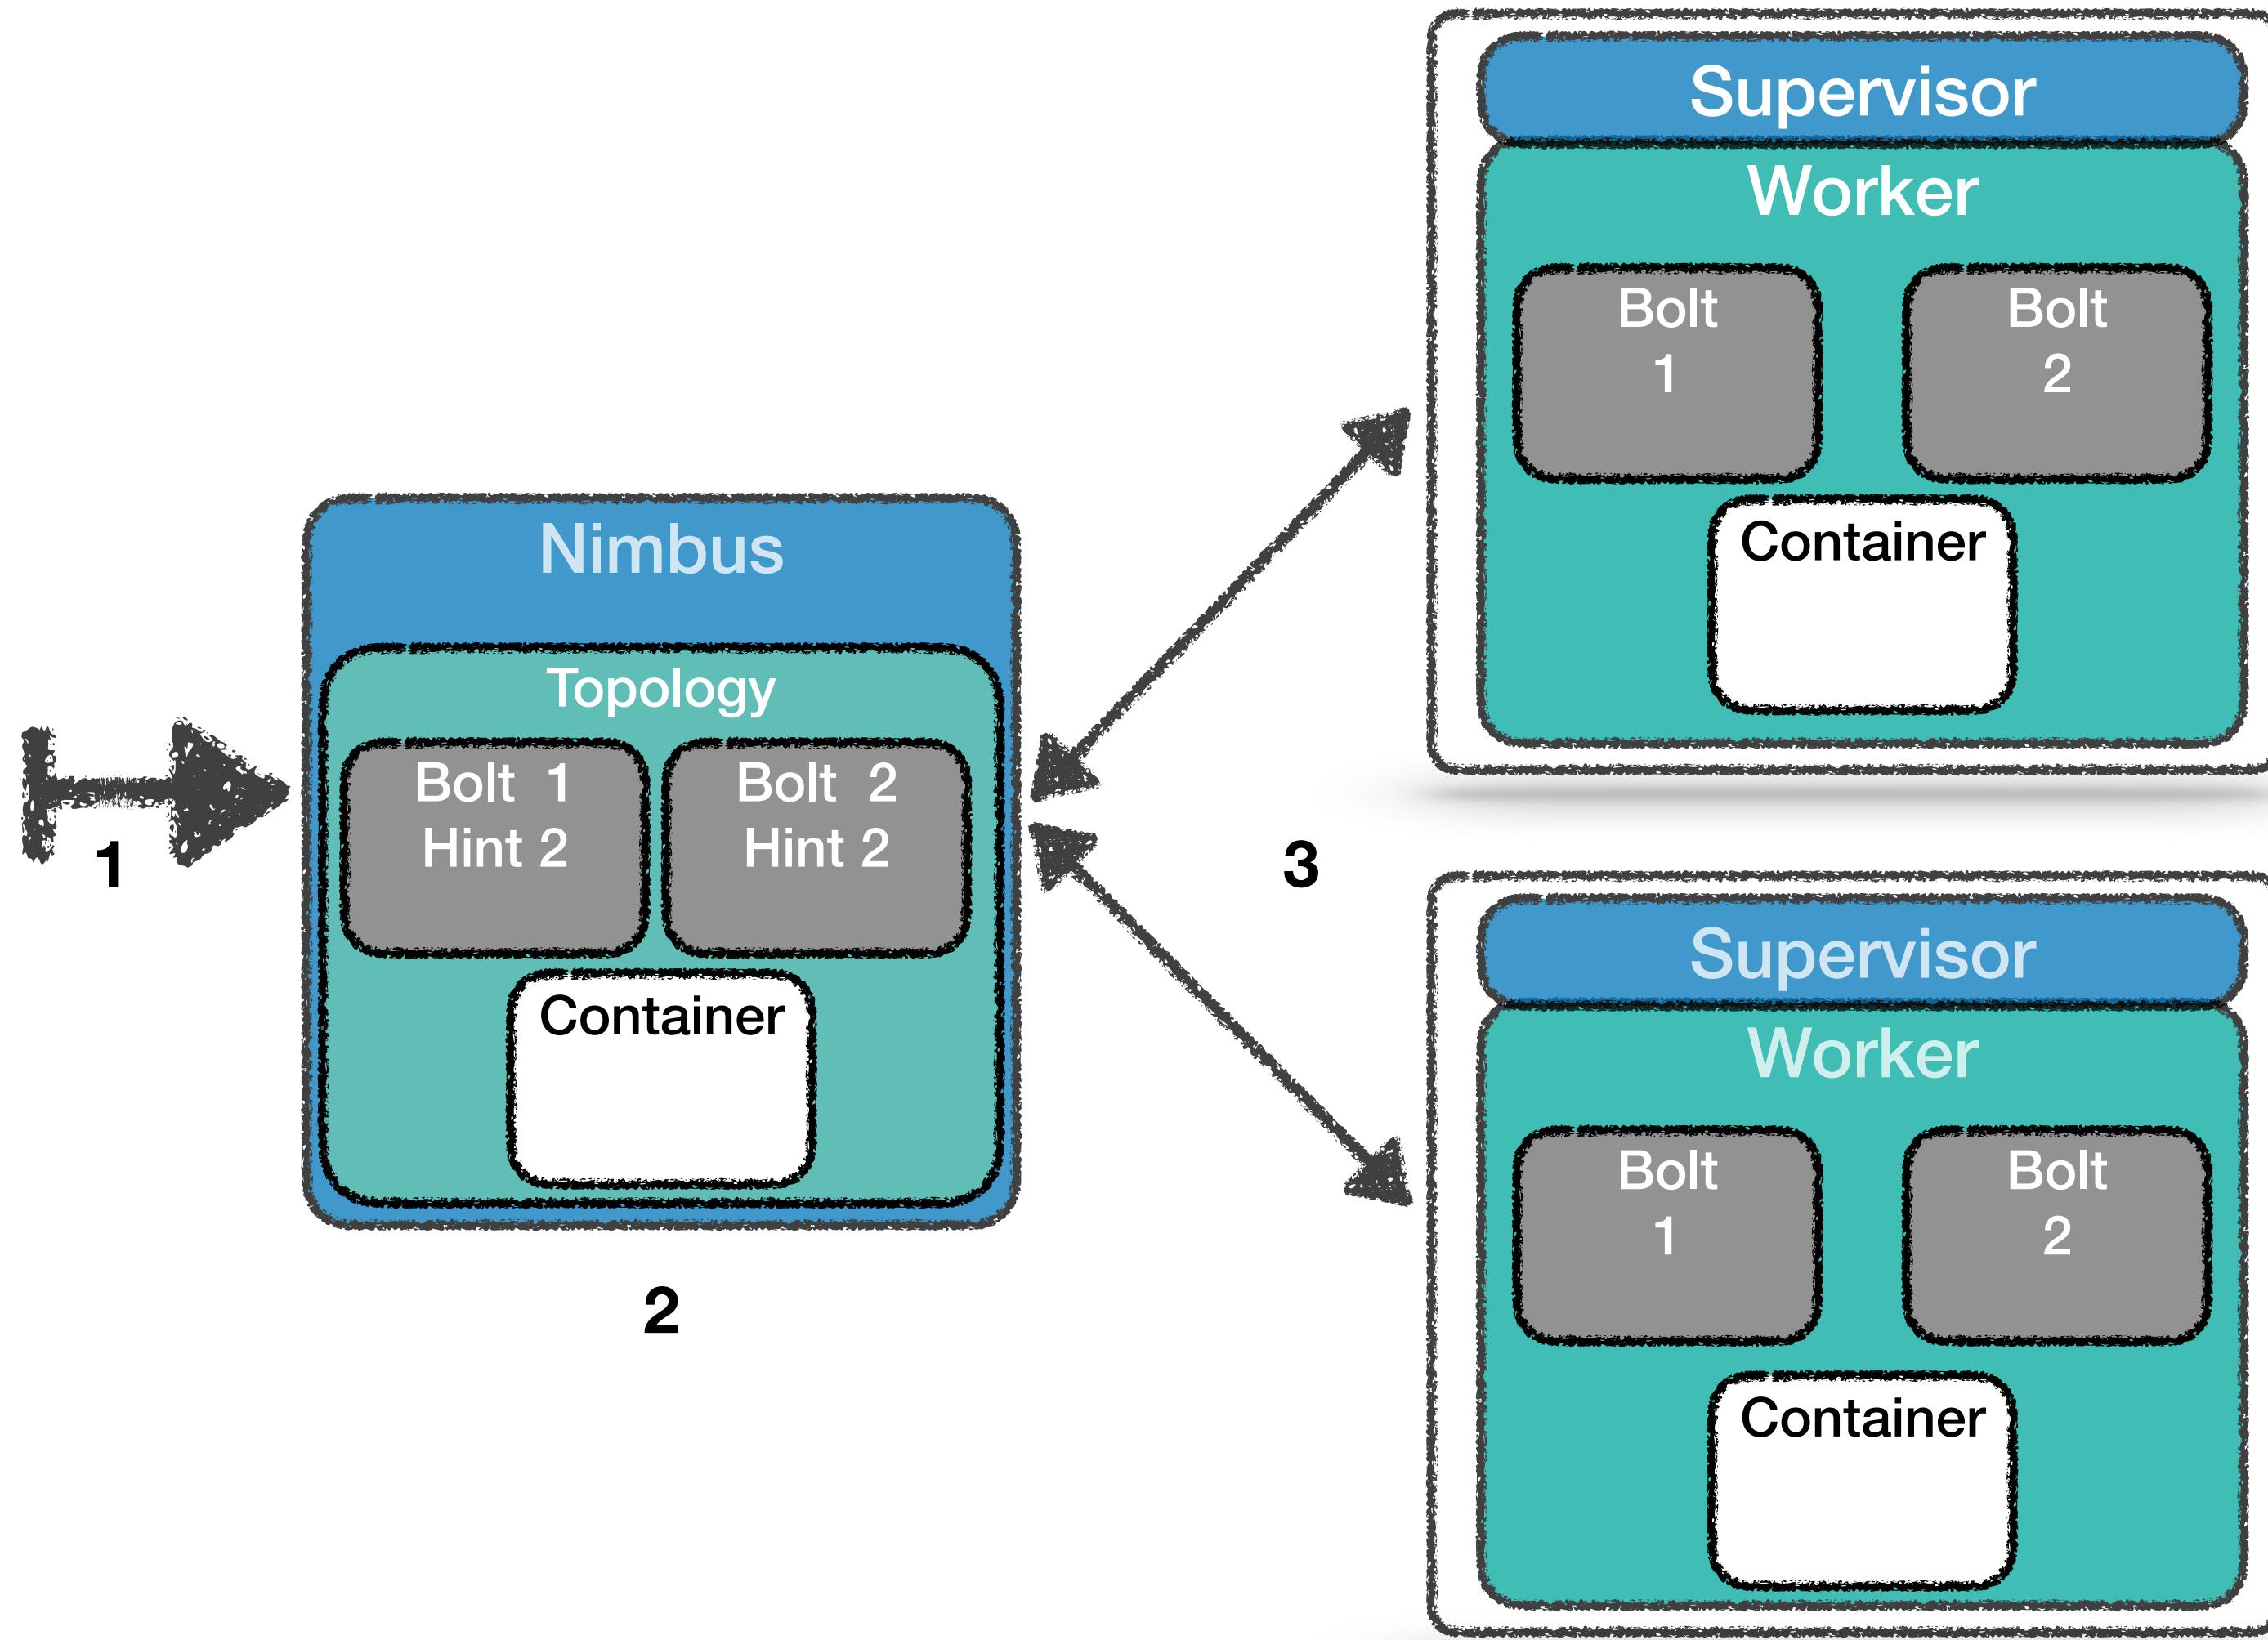
1. The topology's Jar and Config files are uploaded to Nimbus;
2. Nimbus builds the topology and serializes it;
3. The serialized topology is sent to the available workers;

# Dependency Management



- Higher resource consumption because of each bolt having a dedicated IOC container;
- Various container configurations, more difficult manage.

# Dependency Management



- One IOC container per worker shared by all bolt instances;
- It is necessary to synchronize the initialization/access to the container;
- This leads to a greater contention and increased topology startup times;
- Continue to depend on a external dependency (IOC container).

# Why Scala?



**Scala is a strongly statically typed language that:**

- Allows for a more concise and expressive syntax;
- Supports object-oriented programming (OOP) paradigm;
- Supports functional programming (FP) paradigm;
- Has a sophisticated type inference system;
- Runs on the Java Virtual Machine (JVM)

# Why Scala?

**Thin Cake Pattern** - A pure Scala approach to Modularity and DI, where the application is broken down into modules and orchestrated using Scala traits.

[/using-scala-trait-as-modules-or-the-thin-cake-pattern](#)

# Why Scala?

```
trait SimpleTopology extends SpoutModule with ProcessorModule with OutputModule {

    val SPOUT      : String = "SimpleSpout"
    val PROCESSOR : String = "SimpleBolt"
    val OUTPUT     : String = "SecondSimpleBolt"

    def name() = "SimpleTopology"

    def create(): StormTopology = {

        val builder = new TopologyBuilder()

        builder.setSpout(SPOUT, getSpout(), 1)

        builder.setBolt(PROCESSOR, getProcessor(), 1).shuffleGrouping(SPOUT)

        builder.setBolt(OUTPUT, getOutput(), 1).shuffleGrouping(PROCESSOR)

        builder.createTopology()
    }

    def config(): Config = new Config
}
```

# Why Scala?

```
trait ProcessorBolt extends BaseRichBolt {  
  
    var output    : OutputCollector = _  
    val processor: Processor[Instruction] = _  
  
    override def prepare(map: util.Map[_, _], topologyContext: context, output: OutputCollector): Unit = {  
        output = output  
    }  
  
    override def execute(tuple: Tuple): Unit = {  
  
        val instruction = tuple.getValueByField(INSTRUCTION.name).asInstanceOf[Instruction]  
        val results = processor.process(instruction)  
  
        results.map(result => output.emit(tuple, new Values(result)))  
        output.ack(tuple)  
    }  
  
    override def declareOutputFields(outputFieldsDeclarer: OutputFieldsDeclarer): Unit = {  
        outputFieldsDeclarer.declare(new Fields(INSTRUCTION.name))  
    }  
}
```

# Why Scala?

```
trait ProcessorModule {
    def getProcessor(): BaseRichBolt
}

trait SimpleProcessorModule extends ProcessorModule {
    override def getProcessor(): BaseRichBolt = new ProcessorBolt {
        lazy val processor: Processor[Instruction] = new RAMPProcessor
    }
}

trait MockedProcessorModule extends ProcessorModule {
    override def getProcessor(): BaseRichBolt = new ProcessorBolt {
        lazy val processor: Processor[Instruction] = new MockProcessor
    }
}
```

# Why Scala?

```
trait ProcessorModule {
  def getProcessor(): BaseRichBolt
}
trait StandardProcessorModule extends ProcessorModule {
  override def getProcessor(): BaseRichBolt = new ProcessorBolt {
    lazy val processor: Processor[Instruction] = new RAMPProcessor
  }
}

trait MockedProcessorModule extends ProcessorModule {
  override def getProcessor(): BaseRichBolt = new ProcessorBolt {
    lazy val processor: Processor[Instruction] = new MockProcessor
  }
}
```

Lazy Evaluation: the key for the serialization issue

# Why Scala?

```
object SimpleTopologyDeployer {  
  
    def main(args: Array[String]): Unit = {  
  
        val simpleTopology = new SimpleTopology with KafkaSpoutModule with SimpleProcessorModule with KafkaOutputModule  
  
        StormSubmitter.submitTopology(simpleTopology.name(), simpleTopology.config(), simpleTopology.create())  
    }  
}
```

```
object MockedTopologyDeployer {  
  
    def main(args: Array[String]): Unit = {  
  
        val simpleTopology = new SimpleTopology with MockedSpoutModule with MockedProcessorModule with MockedOutputModule  
  
        StormSubmitter.submitTopology(simpleTopology.name(), simpleTopology.config(), simpleTopology.create())  
    }  
}
```

# Mock Tests

```
"SimpleTopology" should "consume an instruction, process it and output the result" in {
    Given("a simple create instruction")
    val instructions = List(Instruction()
        .withIType(CREATE)
        .withHeaders(Map(MeterKey.ID.name -> "5769609"))
        .withEvent("mocked event"))

    val simpleTopology = new SimpleTopology
        with MockSpoutModule
        with MockProcessorModule
        with MockOutputModule

    When("the instruction is processed by the topology")
    val results = runTopology(Map(simpleTopology.SPOUT -> instructions), simpleTopology)

    Then("the result contains two valid tuples emitted by the Kafka Spout and Processor Bolt")
    val instructionsTuple = results.get(simpleTopology.SPOUT)
    val processorTuple    = results.get(simpleTopology.PROCESSOR)

    instructionsTuple.isDefined shouldBe true
    processorTuple.isDefined shouldBe true
}
```

# Mock Tests

ScalaTest: A simple  
Testing Library for Scala

```
"SimpleTopology" should "consume an instruction, process it and output the result" in {  
  
    Given("a simple create instruction")  
  
    val instructions = List(Instruction()  
        .withIType(CREATE)  
        .withHeaders(Map(MeterKey.ID.name -> "5769609"))  
        .withEvent("mocked event"))  
  
    val simpleTopology = new SimpleTopology  
        with MockSpoutModule  
        with MockProcessorModule  
        with MockOutputModule  
  
    When("the instruction is processed by the topology")  
  
    val results = runTopology(Map(simpleTopology.SPOUT -> instructions), simpleTopology)  
  
    Then("the result contains two valid tuples emitted by the Kafka Spout and Processor Bolt")  
  
    val instructionsTuple = results.get(simpleTopology.SPOUT)  
    val processorTuple    = results.get(simpleTopology.PROCESSOR)  
  
    instructionsTuple.isDefined shouldBe true  
    processorTuple.isDefined shouldBe true
```

# Mock Tests

ScalaTest: A simple  
Testing Library for Scala

```
"SimpleTopology" should "consume an instruction, process it and output the result" in {  
  
    Given("a simple create instruction")  
  
    val instructions = List(Instruction()  
        .withIType(CREATE)  
        .withHeaders(Map(MeterKey.ID.name -> "5769609"))  
        .withEvent("mocked event"))  
  
    val simpleTopology = new SimpleTopology  
        with MockSpoutModule  
        with MockProcessorModule  
        with MockOutputModule  
  
    When("the instruction is processed by the topology")  
  
    val results = runTopology(Map(simpleTopology.SPOUT -> instructions), simpleTopology)  
  
    Then("the result contains two valid tuples emitted by the Kafka Spout and Processor Bolt")  
  
    val instructionsTuple = results.get(simpleTopology.SPOUT)  
    val processorTuple    = results.get(simpleTopology.PROCESSOR)  
  
    instructionsTuple.isDefined shouldBe true  
    processorTuple.isDefined shouldBe true
```

It's using  
storm.Testing  
LocalCluster

# Mock Tests

ScalaTest: A simple  
Testing Library for Scala

```
"SimpleTopology" should "consume an instruction, process it and output the result" in {
```

```
    Given("a simple create instruction")
```

```
    val instructions = List(Instruction()  
        .withIType(CREATE)  
        .withHeaders(Map(MeterKey.ID.name -> "5769609"))  
        .withEvent("mocked event"))
```

```
    val simpleTopology = new SimpleTopology  
        with MockSpoutModule  
        with MockProcessorModule  
        with MockOutputModule
```

It's using  
storm.Testing  
LocalCluster

```
    When("the instruction is processed by the topology")
```

```
    val results = runTopology(Map(simpleTopology.SPOUT -> instructions))
```

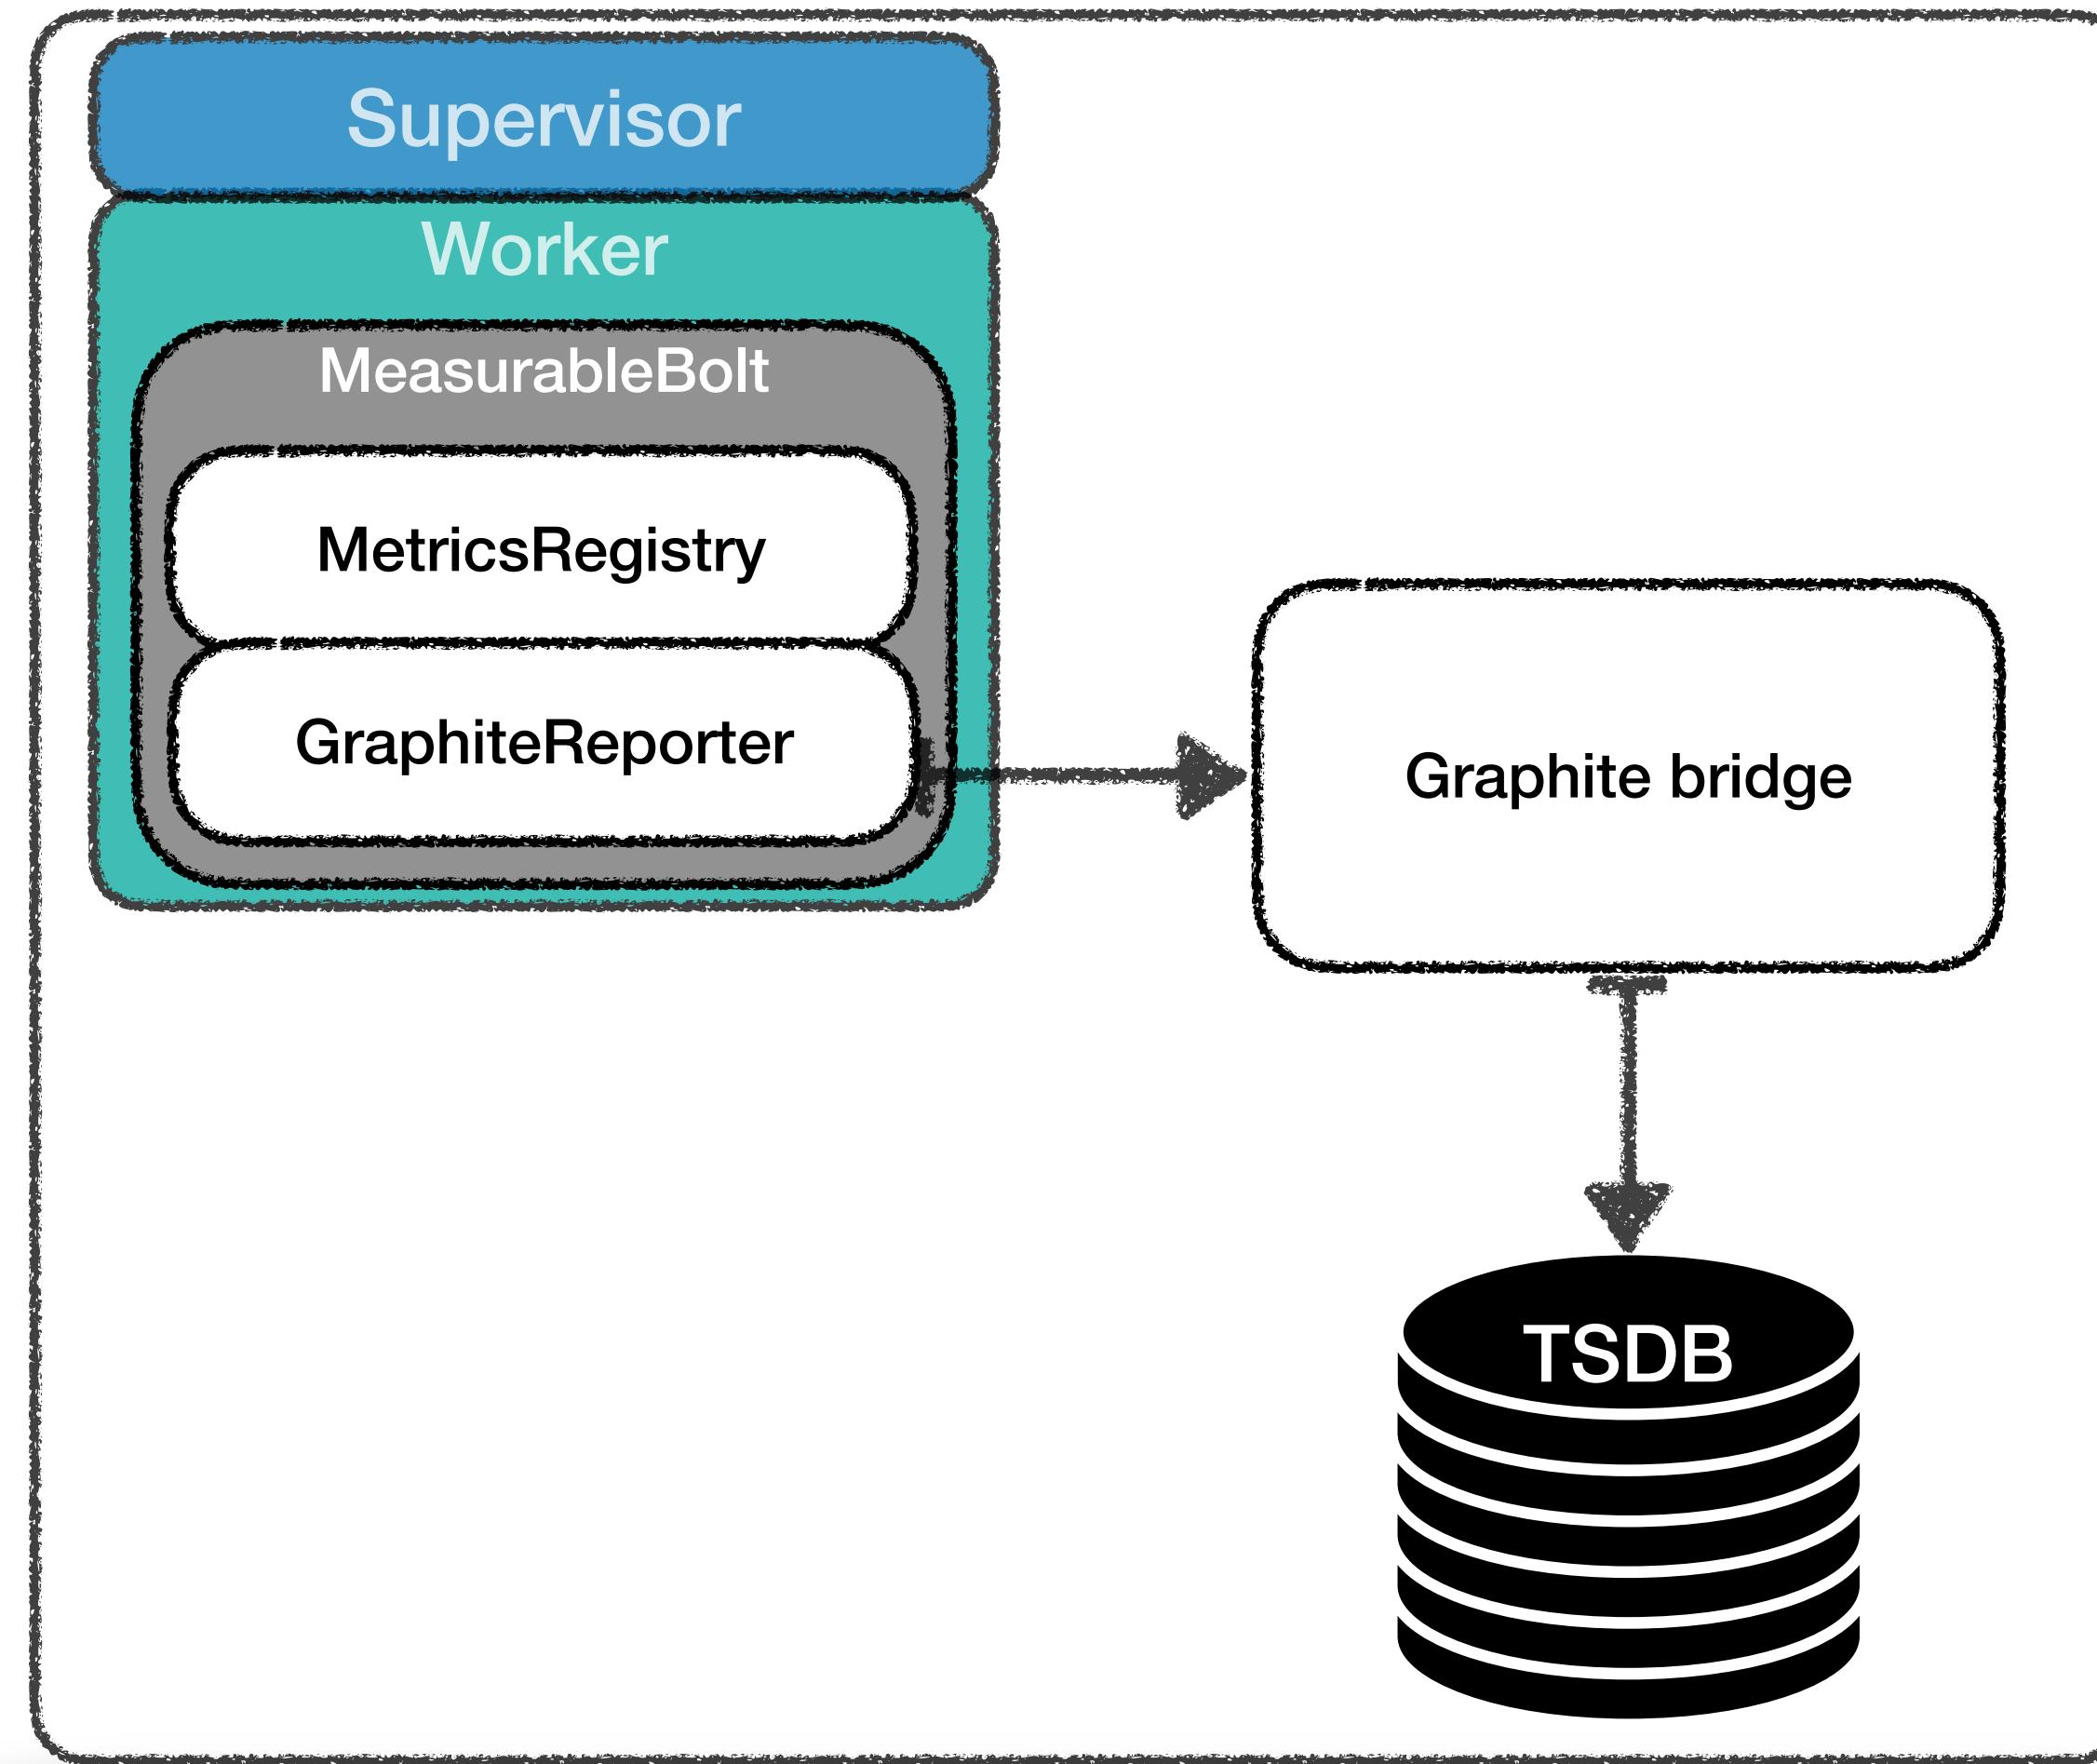
```
    Then("the result contains two valid tuples emitted by the topology")
```

```
    val instructionsTuple = results.get(simpleTopology.SPOUT)  
    val processorTuple = results.get(simpleTopology.PROCESSOR)
```

```
    instructionsTuple.isDefined shouldBe true  
    processorTuple.isDefined shouldBe true
```

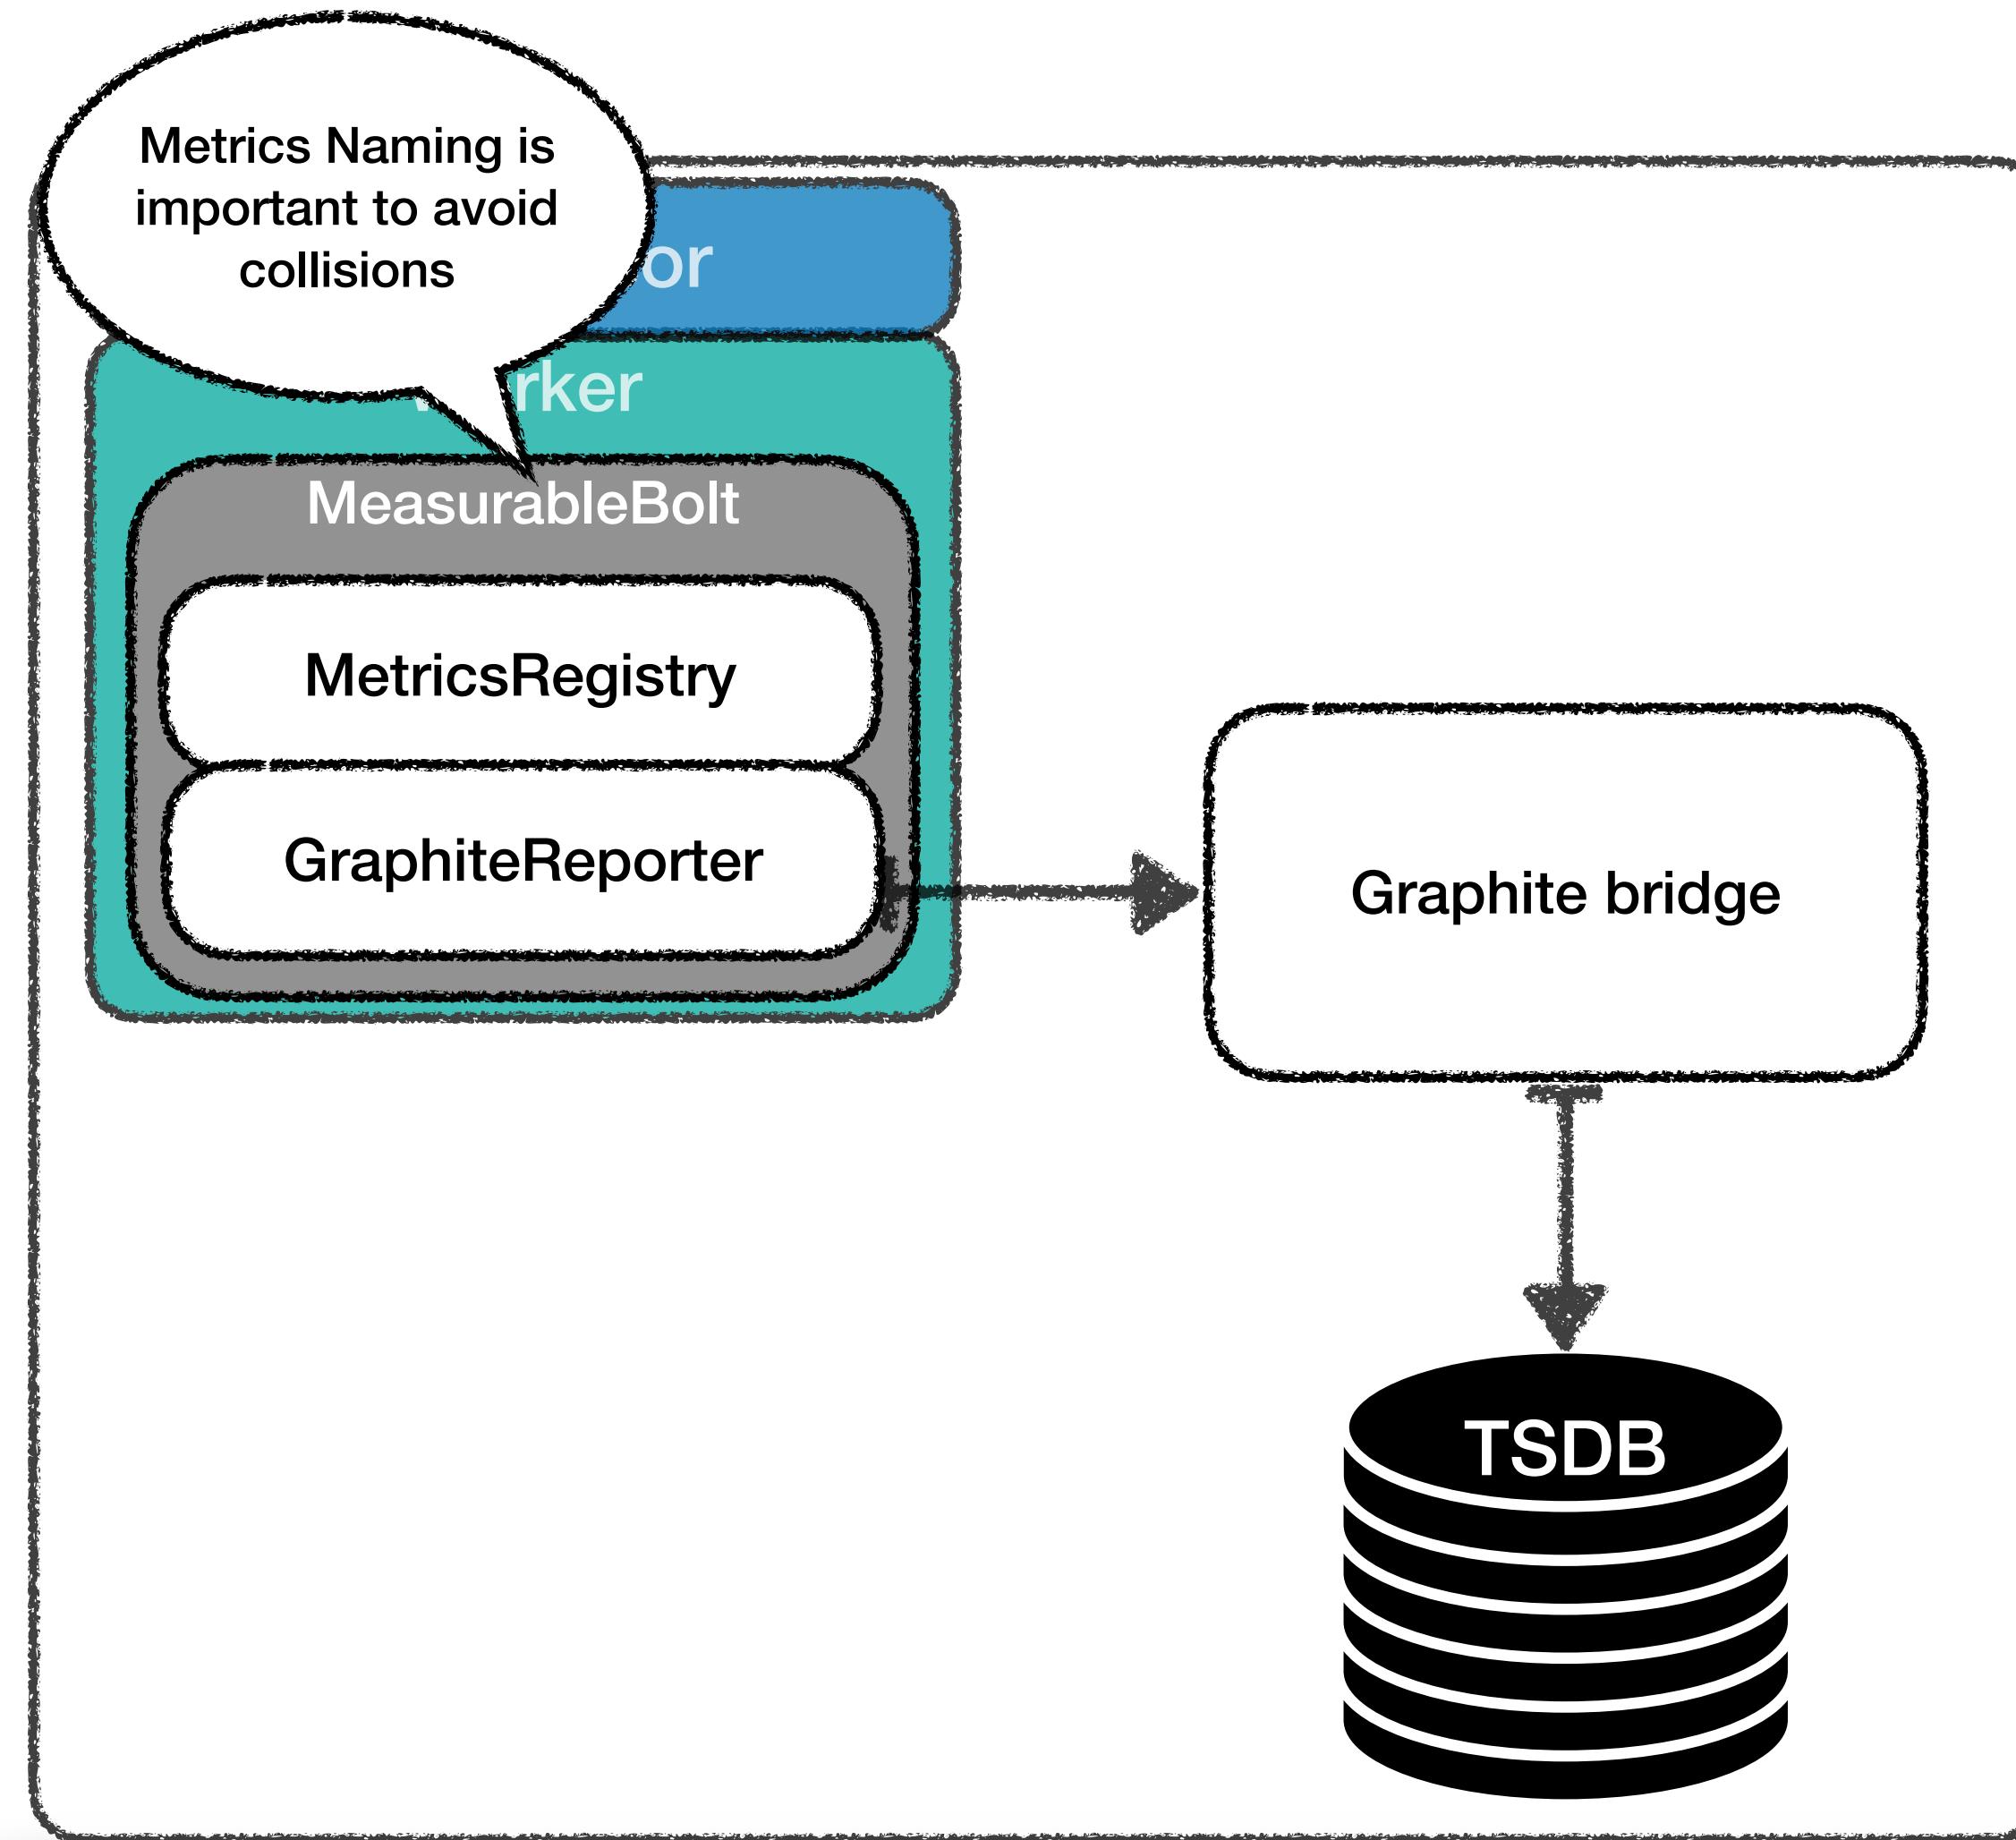
It only captures  
emitted tuples from  
Spouts/Bolts

# Metrics



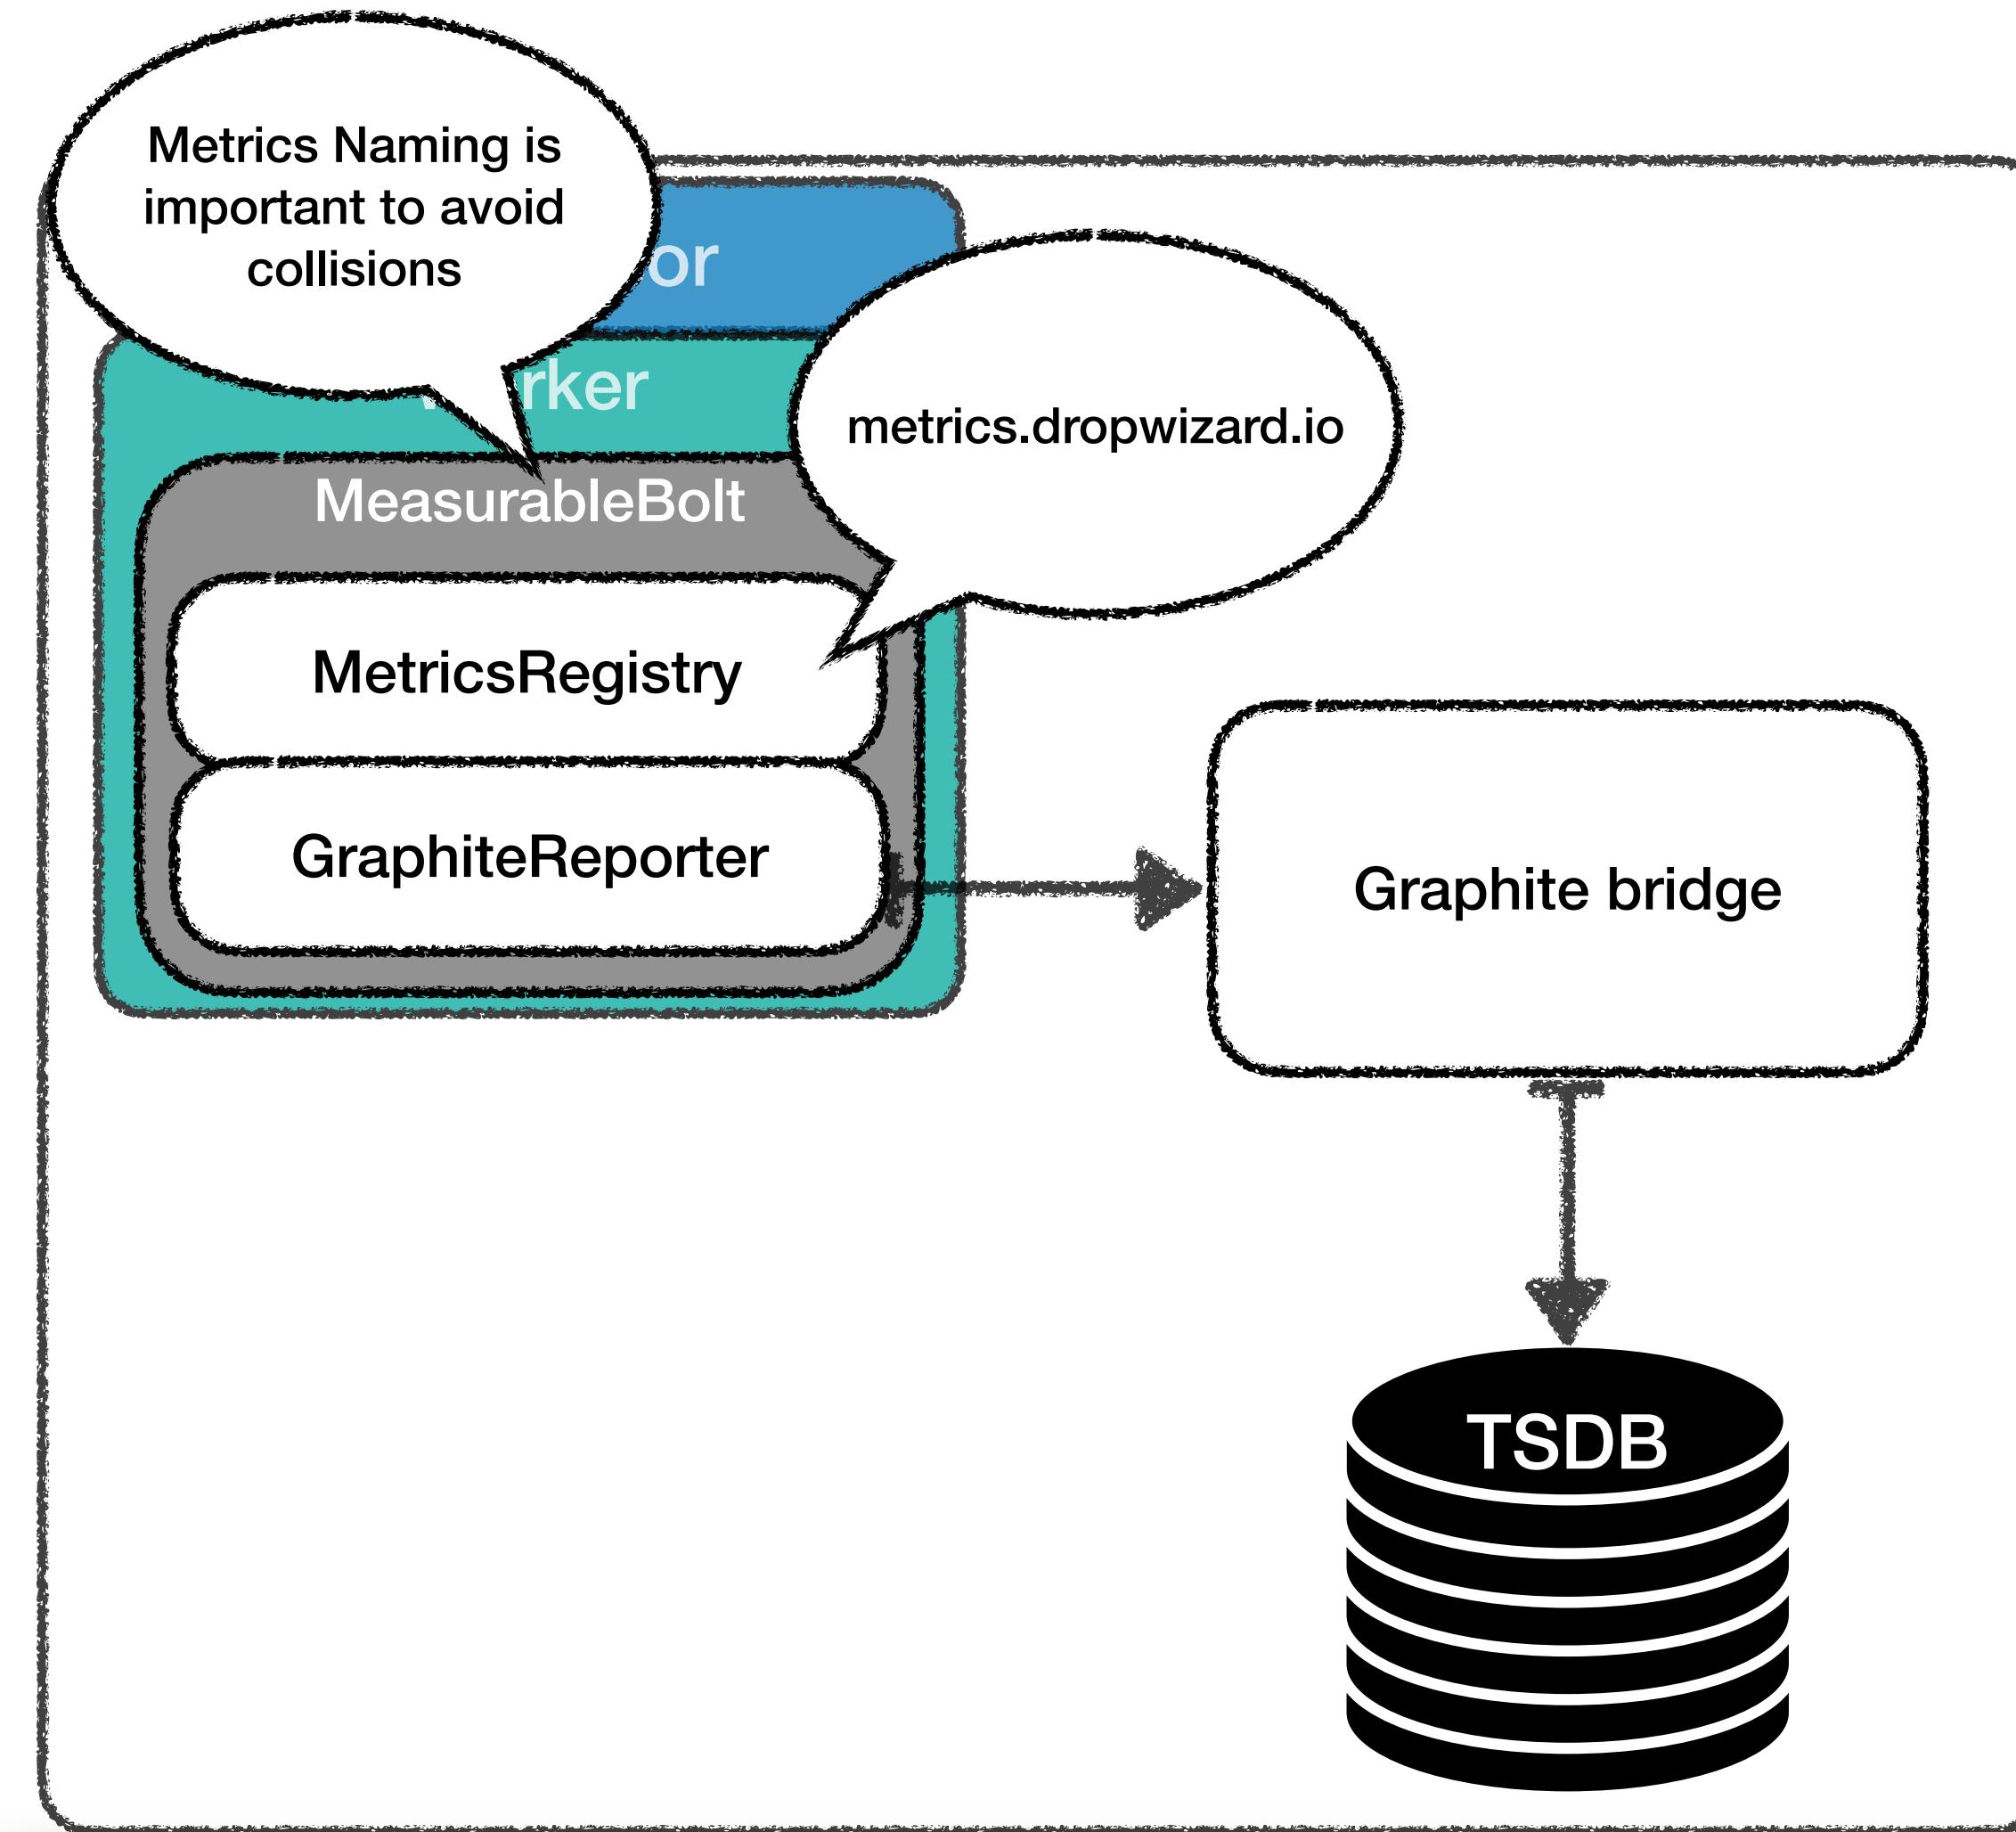
[/Flutter-Global/topology-management-lib](#)

# Metrics



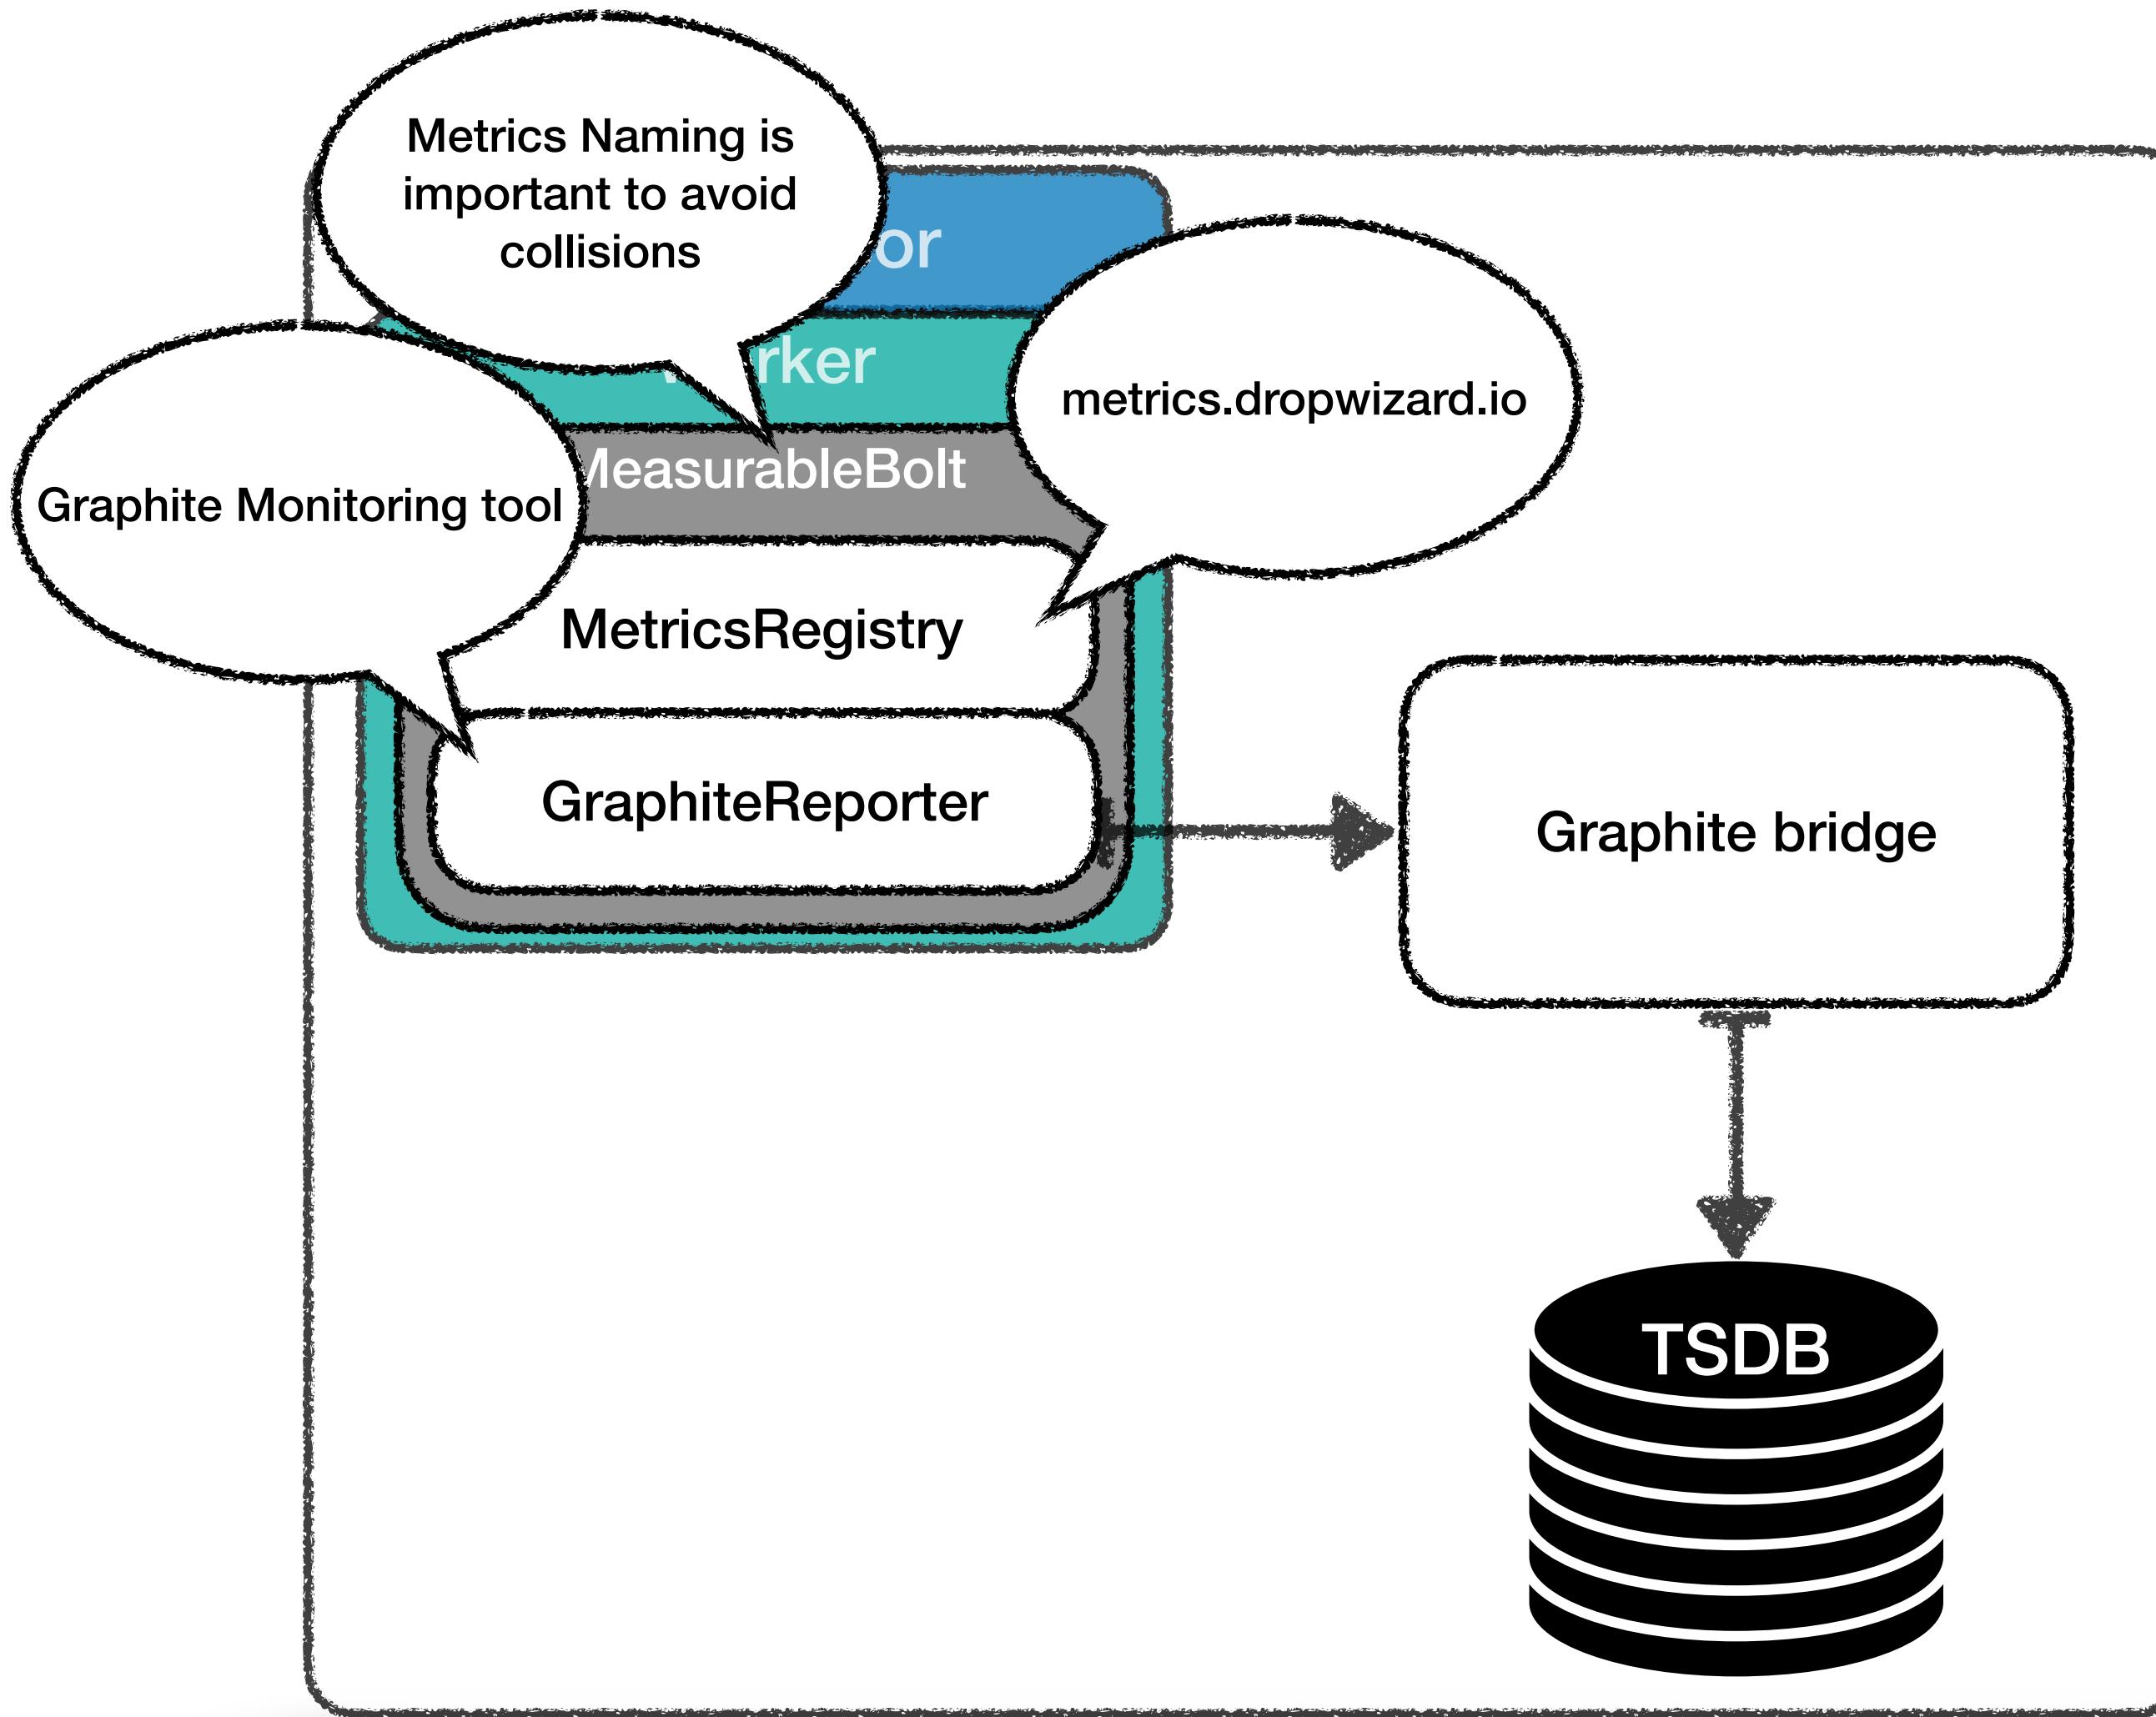
[/Flutter-Global/topology-management-lib](#)

# Metrics



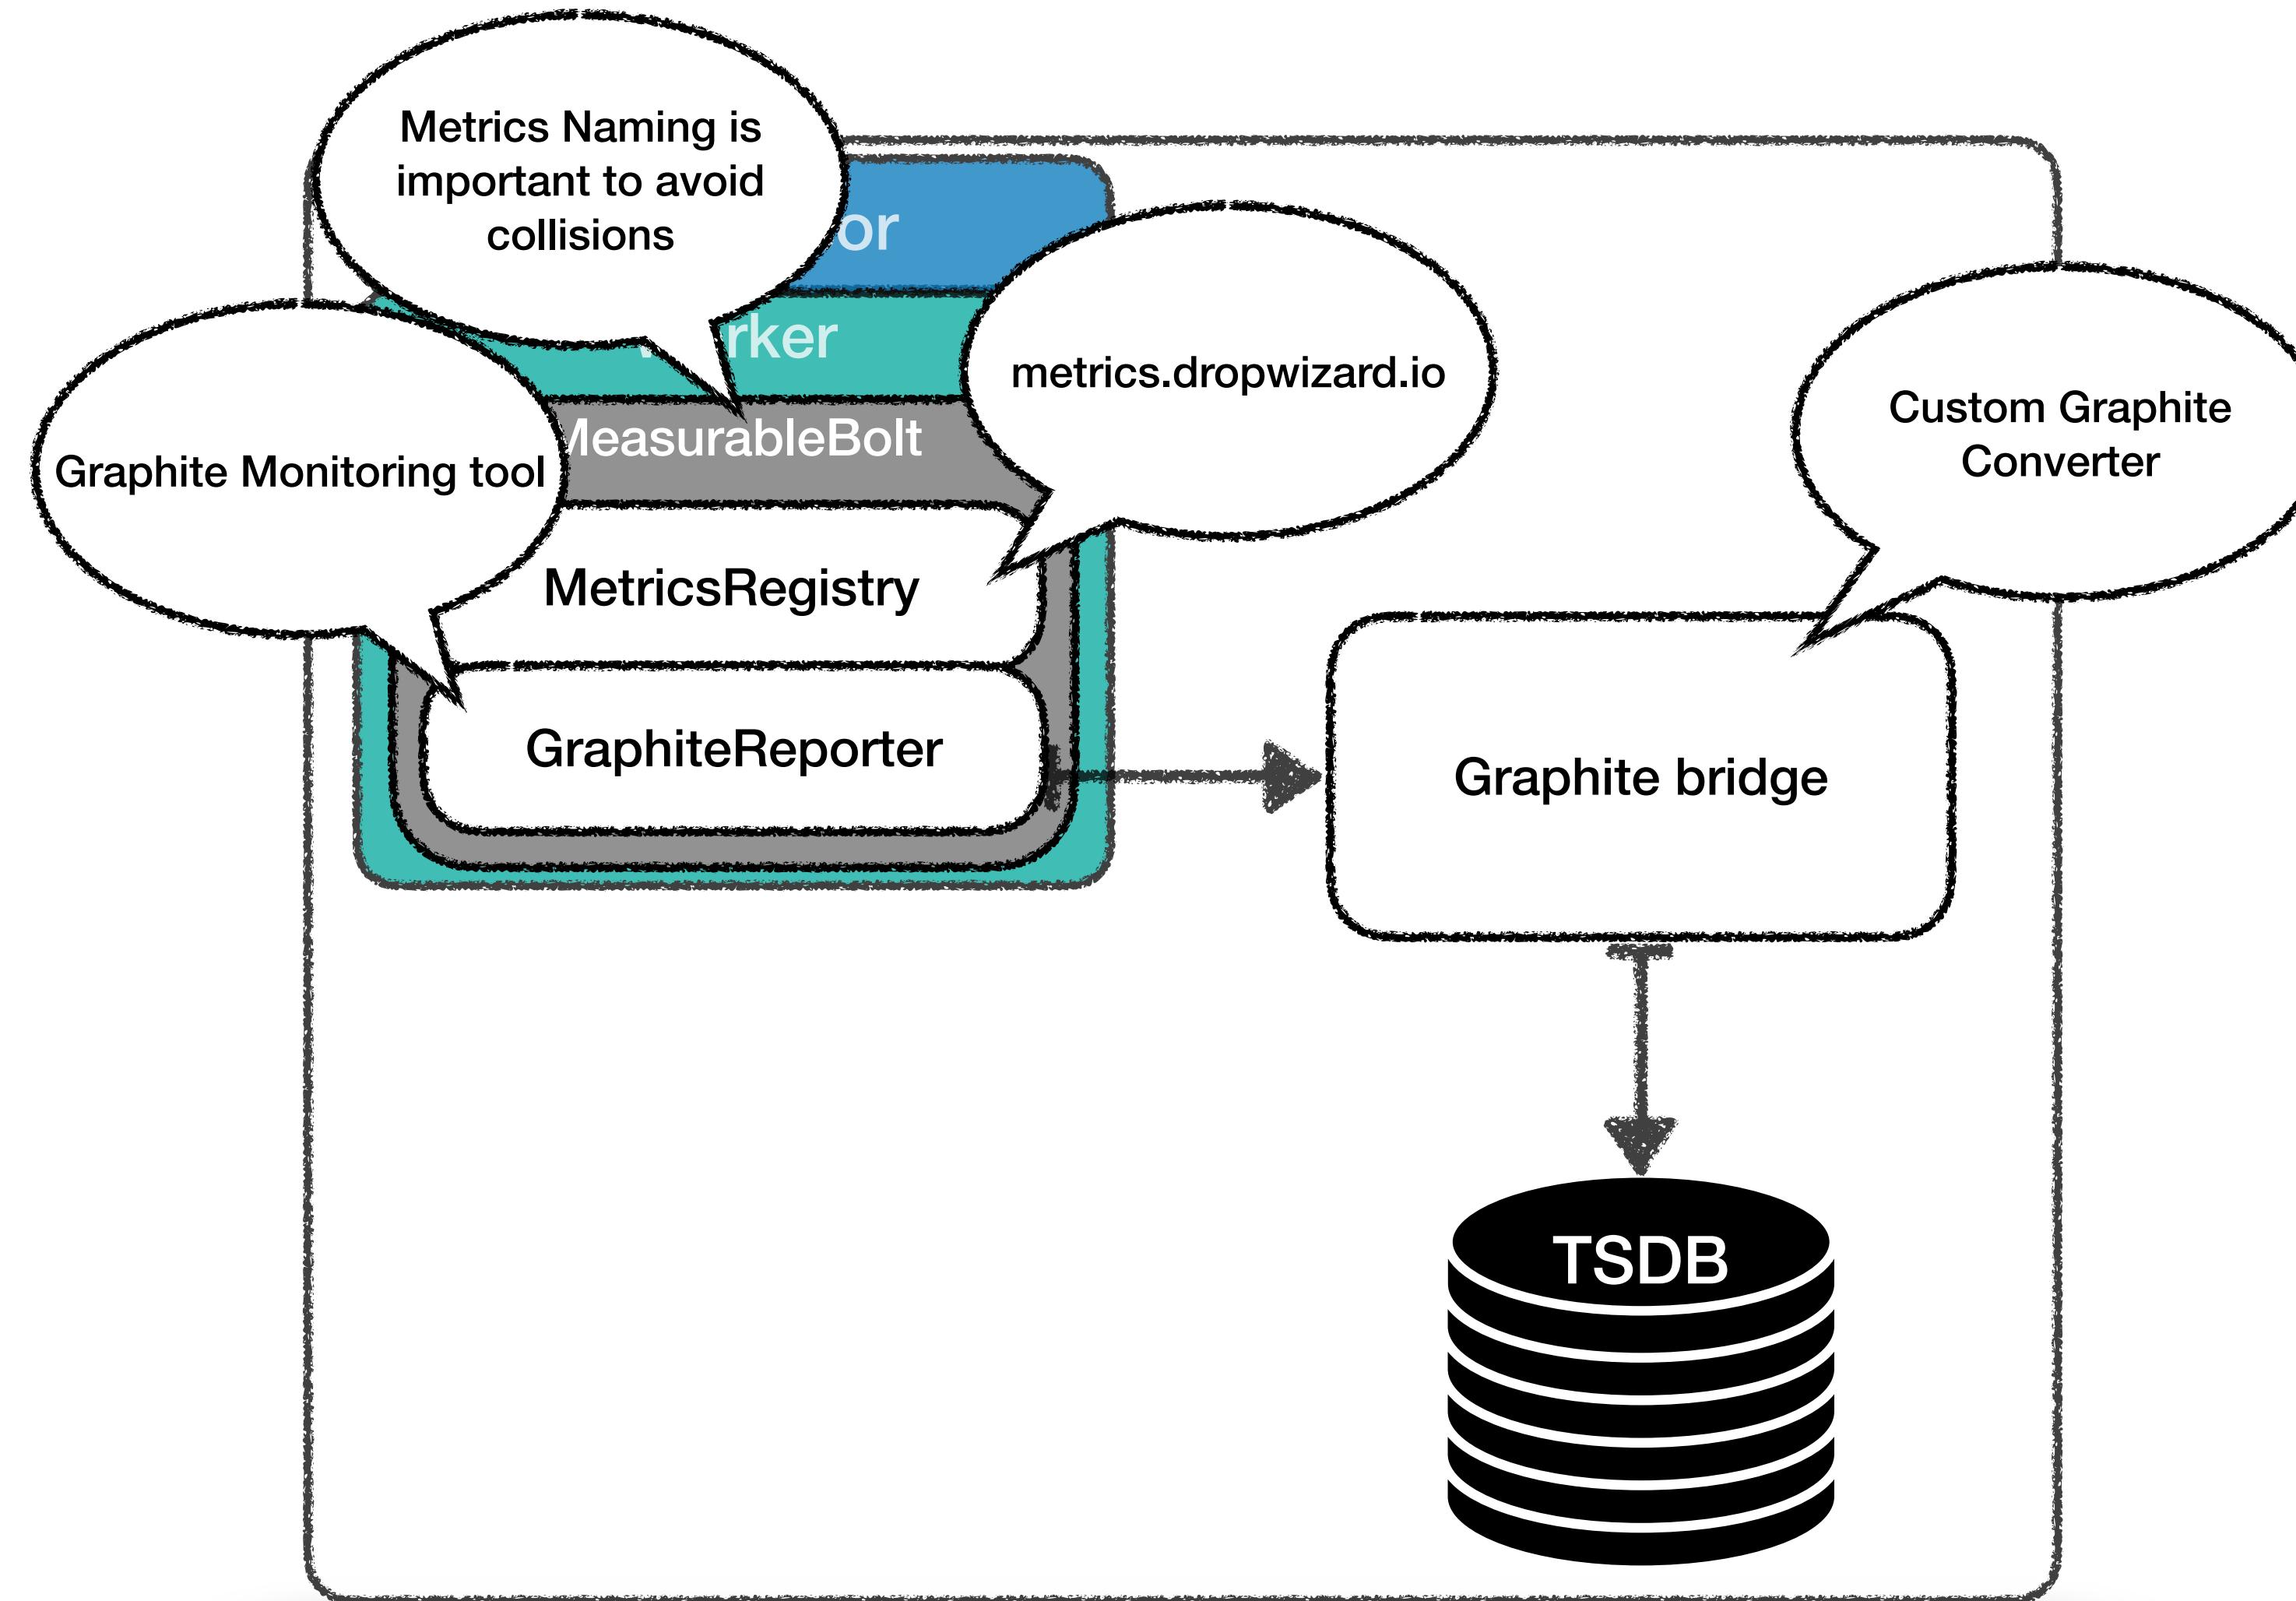
[/Flutter-Global/topology-management-lib](#)

# Metrics



[/Flutter-Global/topology-management-lib](#)

# Metrics



[/Flutter-Global/topology-management-lib](#)

# Metrics

```
trait ProcessorBolt extends MeasurableRichBolt {

    val processor: Processor[Instruction]

    override def apply(tuple: Tuple, output: OutputCollector): Unit = {
        val instruction = tuple.getValueByField(INSTRUCTION.name).asInstanceOf[Instruction]
        val results = processor.process(instruction)
        results.map(result => output.emit(tuple, new Values(result)))
        output.ack(tuple)
        getCounter("BoltCounter").inc()
    }

    override def declareOutputFields(outputFieldsDeclarer: OutputFieldsDeclarer): Unit = {
        outputFieldsDeclarer.declare(new Fields(INSTRUCTION.name))
    }
}
```

[/Flutter-Global/topology-management-lib](#)

# Metrics

Custom Bolt in  
management library

```
trait ProcessorBolt extends MeasurableRichBolt {  
  
    val processor: Processor[Instruction]  
  
    override def apply(tuple: Tuple, output: OutputCollector): Unit = {  
  
        val instruction = tuple.getValueByField(INSTRUCTION.name).asInstanceOf[Instruction]  
  
        val results = processor.process(instruction)  
  
        results.map(result => output.emit(tuple, new Values(result)))  
        output.ack(tuple)  
        getCounter("BoltCounter").inc()  
    }  
  
    override def declareOutputFields(outputFieldsDeclarer: OutputFieldsDeclarer): Unit = {  
        outputFieldsDeclarer.declare(new Fields(INSTRUCTION.name))  
    }  
}
```

[/Flutter-Global/topology-management-lib](#)

# Metrics

```
trait ProcessorBolt extends MeasurableRichBolt {  
  
    val processor: Processor[Instruction]  
  
    override def apply(tuple: Tuple, output: OutputCollector): Unit = {  
  
        val instruction = tuple.get(INSTRUCTION.name).asInstanceOf[Instruction]  
        val result = processor.process(instruction)  
        val results = List(result).map(tuple._1.copy(values = Values(result)))  
        output.emit(tuple, results)  
        output.ack(tuple)  
        getCounter("BoltCounter").inc()  
    }  
  
    override def declareOutputFields(outputFieldsDeclarer: OutputFieldsDeclarer): Unit = {  
        outputFieldsDeclarer.declare(new Fields(INSTRUCTION.name))  
    }  
}
```

Custom Bolt in management library

Also has Timers, Meters, Gauges & Histograms

[/Flutter-Global/topology-management-lib](#)

# Metrics

[DEMO Grafana Dashboards](#)

# Exercise

Checkout [STT Project](#)

Implement a Topology that consumes strings of sentences from a Kafka topic:

- Consume from the “test-input-topic” **stream** using a KafkaSpout (“localhost:31092”);
- Process each sentence:
  - Split them into **words**;
  - Have a **counter** for each **word** that is incremented at every occurrence of that word in any sentence.
  - For each word emit a **resulting** string containing the following format “word = <word>, count = <count>”
- Publish the **results** to the “test-output-topic” **stream**;

Deploy topology:

```
bin/storm jar /stt/stt-storm/target/stt-storm-1.0.5-SNAPSHOT.jar  
com.ppb.feeds.stt.storm.topology.deploy.STTDeployer
```

